Class Topics (클래스 홈페이지 참조)

- □ Part 1: Fundamental concepts and principles
- □ Part 2: 빠른 컴퓨터를 위한 ISA design
 - Topic 1 Computer performance and ISA design (Ch. 1)
 - Topic 2 RISC (MIPS) instruction set (Chapter 2)
 - Topic 3 Computer arithmetic and ALU (Chapter 3)
- □ Part 3: ISA 의 효율적인 구현 (pipelining, cache memory)
 - Topic 4 Processor design (Chapter 4)
 - Single cycle implementation
 - Pipelined implementation
 - Topic 5 Memory system design (Chapter 5)



COMPUTER ORGANIZATION AND DESIGN



The Hardware/Software Interface

Chapter 4

The Processor (Implementing ISA)

Part 2: Pipelining

Some of authors' slides are modified

Pipelined Laundry (복습)

- □ 3-stage pipelining problem
 - Infinite loads: speedup = 3 = number of stages

		•	ı			9	
☐ Need m	ore har	dware?	What i	f we ha	ve more	e hardw	vare?
						Tim	$\stackrel{\text{le}}{\longrightarrow}$
빨래	Washer	Drier	정리	Washer	Drier	정리]
빨래	Washer	Drier	정리				
빨래		Washer	Drier	정리			
빨래	·		Washer	Drier	정리		
		•		Į		ı	



Pipelined instruction Execution(복습)

- □ 3-stage pipeline
 - Overlapped execution, instruction-level parallelism

	In	struction	n_1	In	2	
	Fetch	Decode	Execute	Fetch	Decode	Execute
				1		Time
Instruction ₁	F	D	Е			
Instruction ₂		F	D	Е		
Instruction ₃			F	D	Е	
Instruction ₄				F	D	Е
1 1 1 1				C	hapter 4 — T	he Processor —

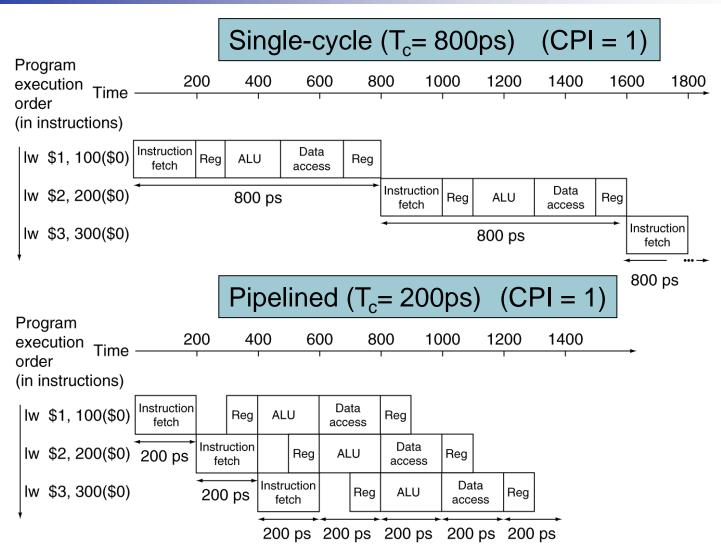
MIPS Pipeline (반복)

- ☐ Five stages, one step per stage
 - 1. IF: Instruction fetch from memory
 - 2. ID: Instruction decode & register read
 - 3. EX: Execute operation or calculate address
 - 4. MEM: Access memory operand
 - 5. WB: Write result back to register

Instr	Instr fetch	Register read	ALU op	Memory access	Register write	Total time
lw	200ps	100 ps	200ps	200ps	100 ps	800ps
SW	200ps	100 ps	200ps	200ps		700ps
R-format	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps

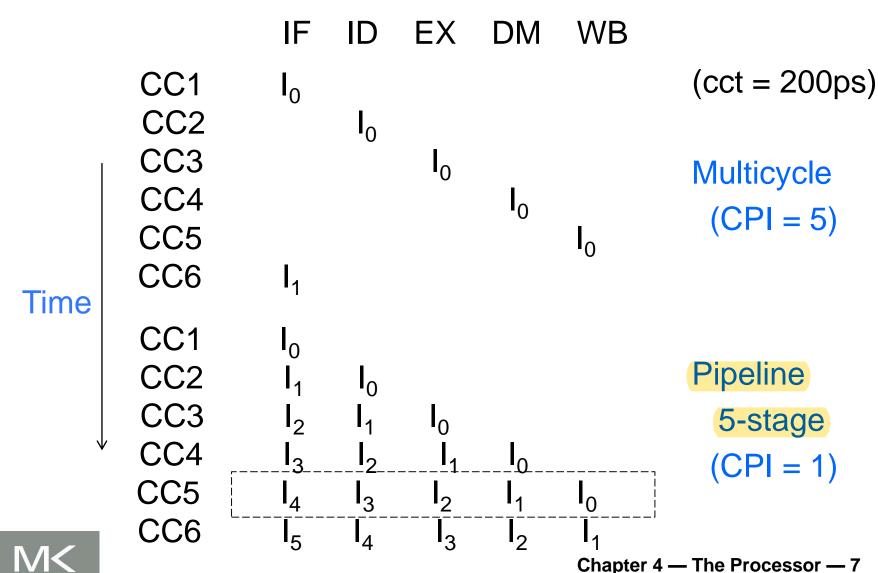


Pipeline Performance





Multicycle vs. Pipelining (반복)



Multi-Cycle Pipeline Diagram

□ Traditional form

Program execution order (in instructions)

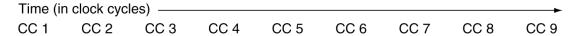
lw \$10, 20(\$1)	Instruction fetch	Instruction decode	Execution	Data access	Write back				
sub \$11, \$2, \$3		Instruction fetch	Instruction decode	Execution	Data access	Write back			
add \$12, \$3, \$4	,		Instruction fetch	Instruction decode	Execution	Data access	Write back		
lw \$13, 24(\$1)				Instruction fetch	Instruction decode	Execution	Data access	Write back	
add \$14, \$5, \$6					Instruction fetch	Instruction decode	Execution	Data access	Write back

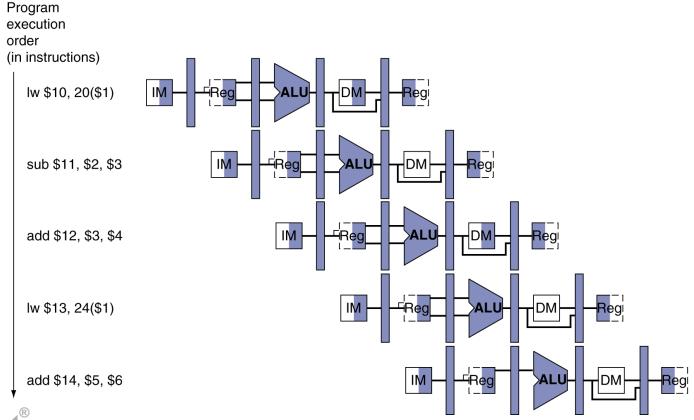


Pipeline Diagram

□ Form showing resource usage

Time

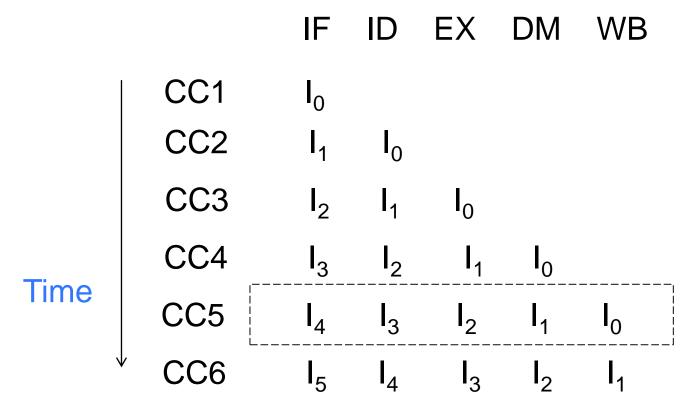






Pipeline Diagram

- ☐ Each stage run different instructions
 - Only one set of hardware



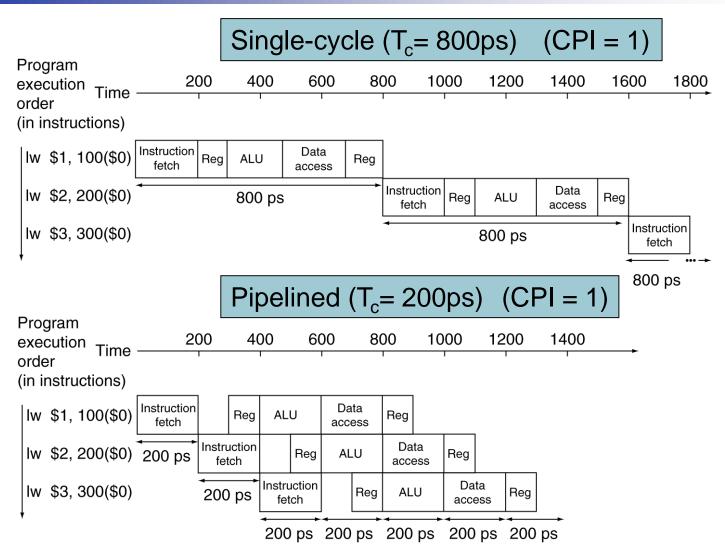


Pipeline Speedup

- ☐ If all stages are balanced
 - i.e., all take the same time
 - Ideal speedup = number of stages
- ☐ If not balanced, speedup is less
- Speedup due to increased throughput
 - Latency (time for each instruction) does not decrease



Pipeline Performance (반복)





Pipelining (미리보기)

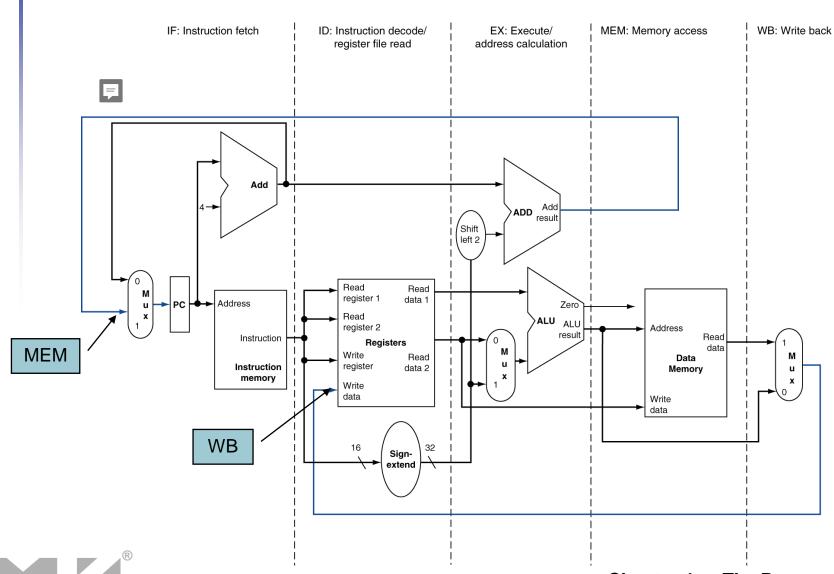
- What makes it easy?
 - All instructions are the same length
 - Just a few instruction formats
 - Memory operands appear only in loads and stores
- What makes it hard?
 - Structural hazards: suppose we had only one memory
 - Data hazards: instruction depends on previous one
 - Control hazards: need to worry about branch
- ☐ We'll build a simple pipeline and look at these issues



Pipelined Datapath and Control

(Implementation for Pipelining)

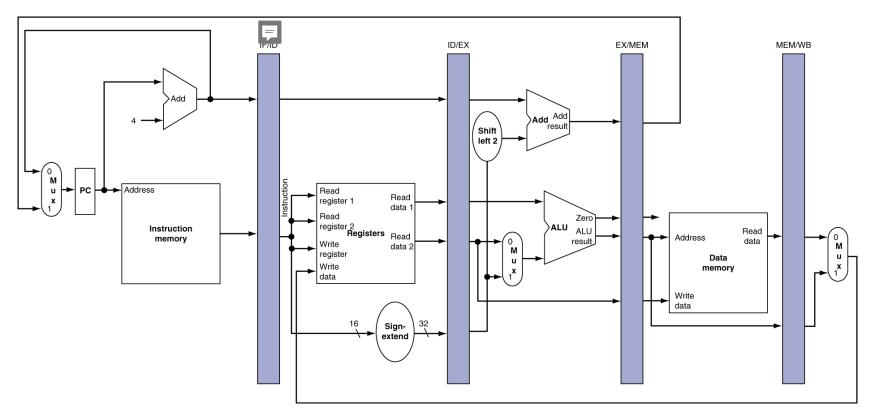
MIPS Pipelined Datapath





Pipeline registers

- □ Need registers between stages IF ↑ ID ↑ EX
 - To hold information produced in previous cycle



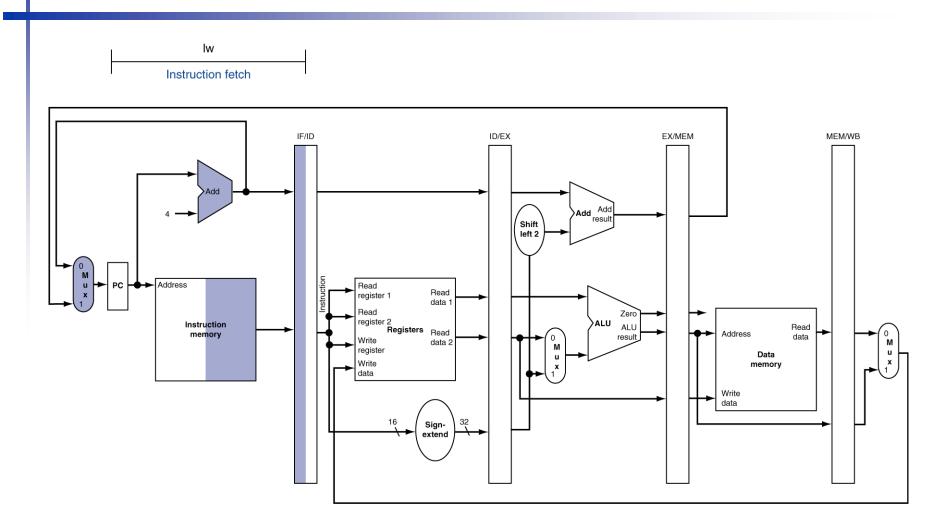


Pipeline Operation

- ☐ Cycle-by-cycle flow of instructions through the pipelined datapath
 - We'll look at diagrams for load & store

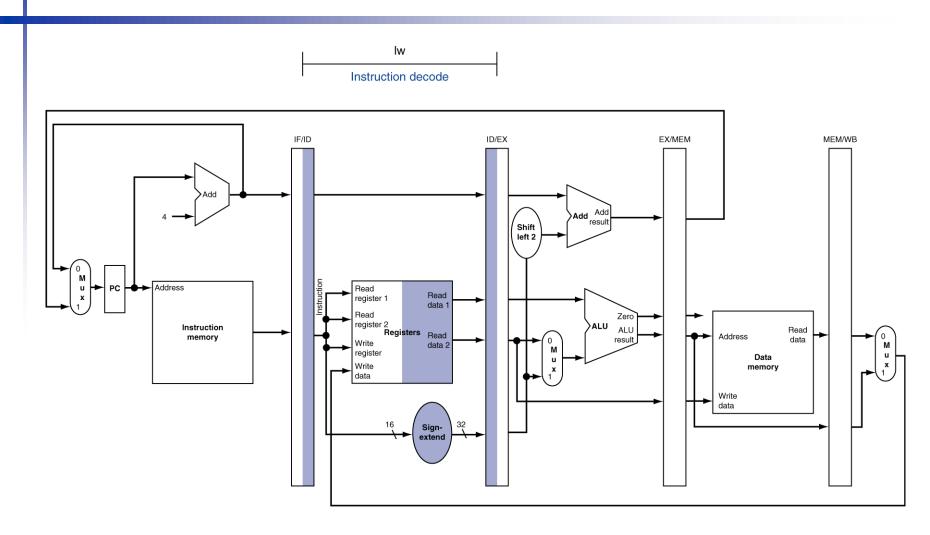


IF for Load, Store, ...



□ Information in IF/ID register: instruction, PC+4 (금통)

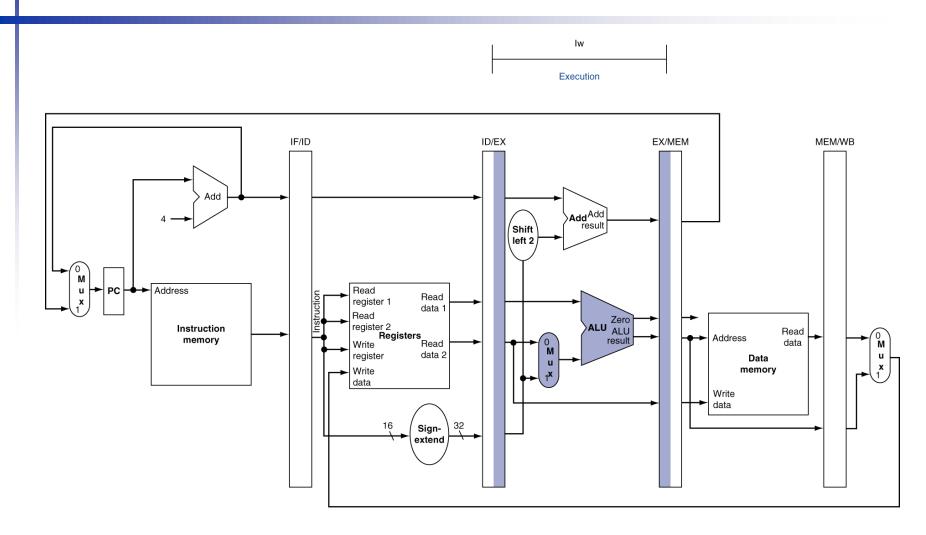
ID for Load, Store, ...



□ ID/EX: PC+4, registers, immediate, control signals (금통)

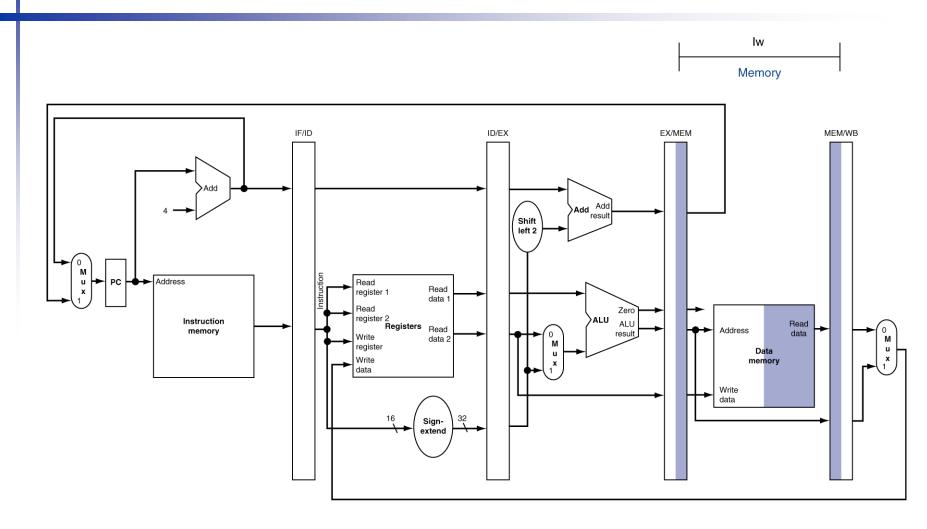


EX for Load



☐ EX/MEM register: memory address, control signals

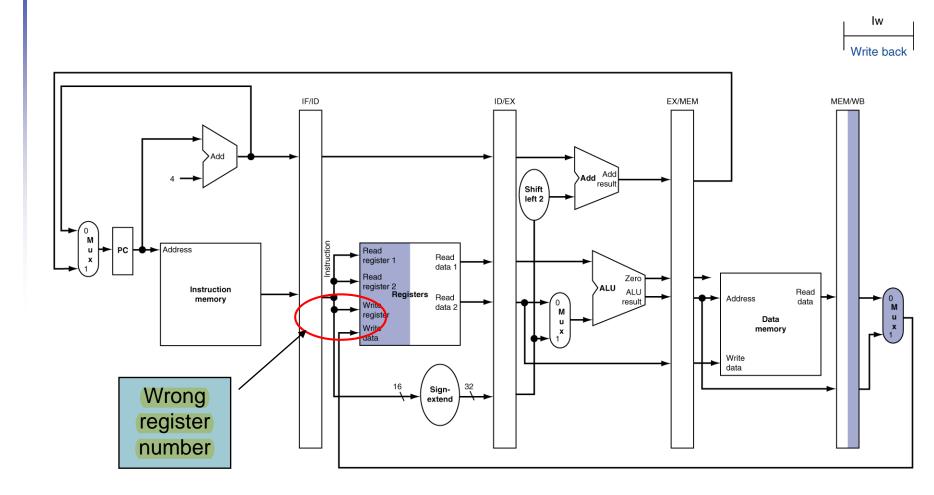
MEM for Load



☐ MEM/WB register: data read, control signals

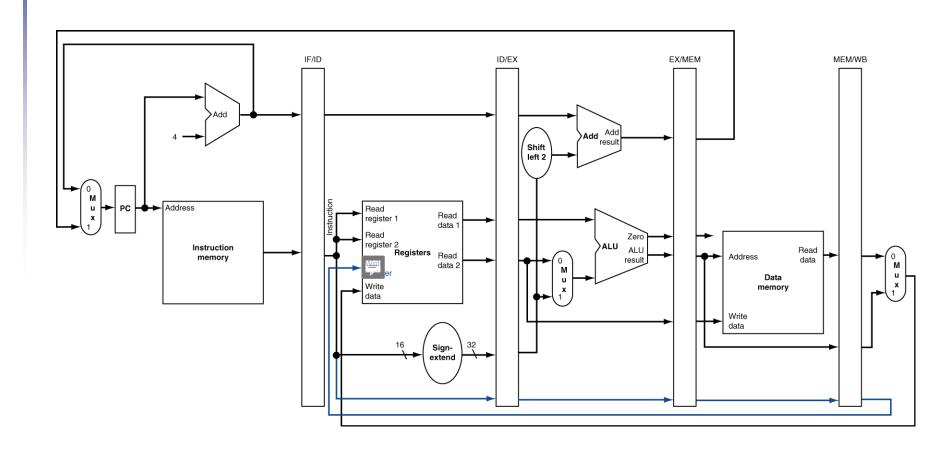


WB for Load



☐ Execution finished in WB; register write by WB

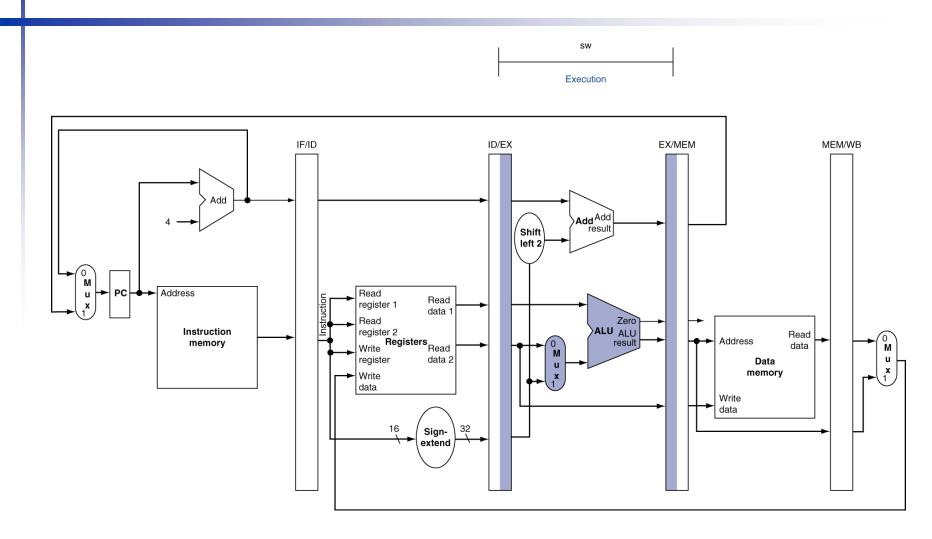
Corrected Datapath for Load



☐ Carry "write register no." information all the way through WB



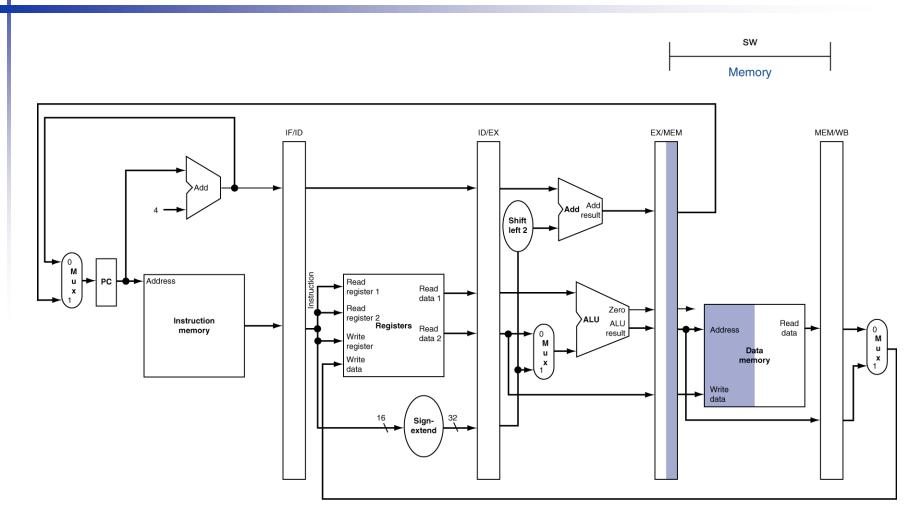
EX for Store



☐ EX/MEM: memory address, data to write, control signals



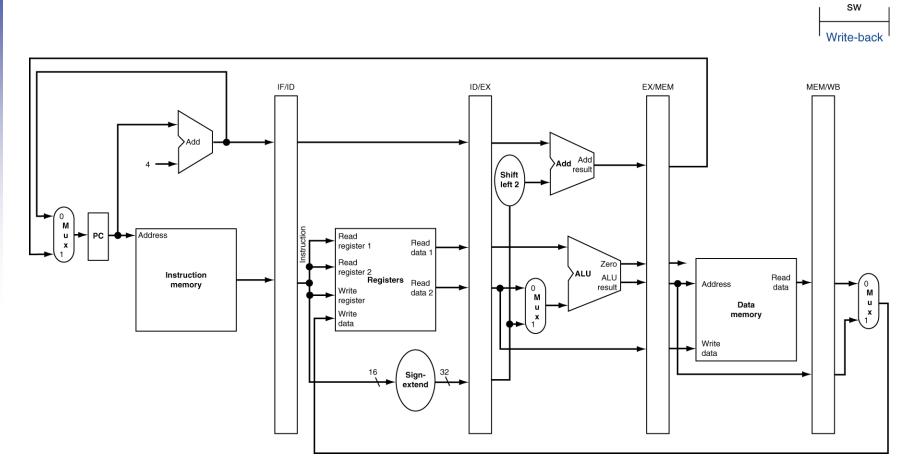
MEM for Store



☐ MEM/WB register: none (store is finished)



WB for Store



- Nothing to do; doesn't matter
 - Important: at every cycle, start (or finish) one instruction



ALU instructions and Branch

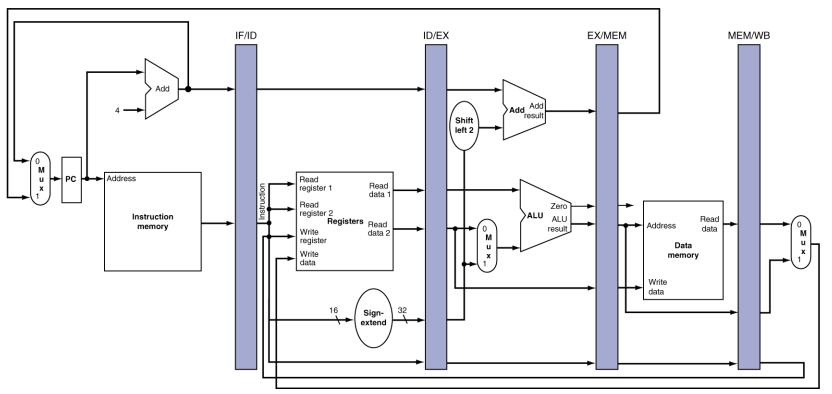
- ☐ ALU instructions ("add \$3, \$1, \$2")
 - EX/MEM: ALU result, write register no., control signals
 - NO operation in MEM stage
 - MEM/WB: ALU result, write register no., control signals
 - Execution finished in WB stage
- ☐ Branch instructions ("beq \$1, \$2, 8"
 - EX/MEM: branch target, zero or not, control signals
 - Execution finished in MEM stage
 - Update PC if branch is taken (do nothing if untaken)



Pipeline in Action

□ Each stage run a different instruction

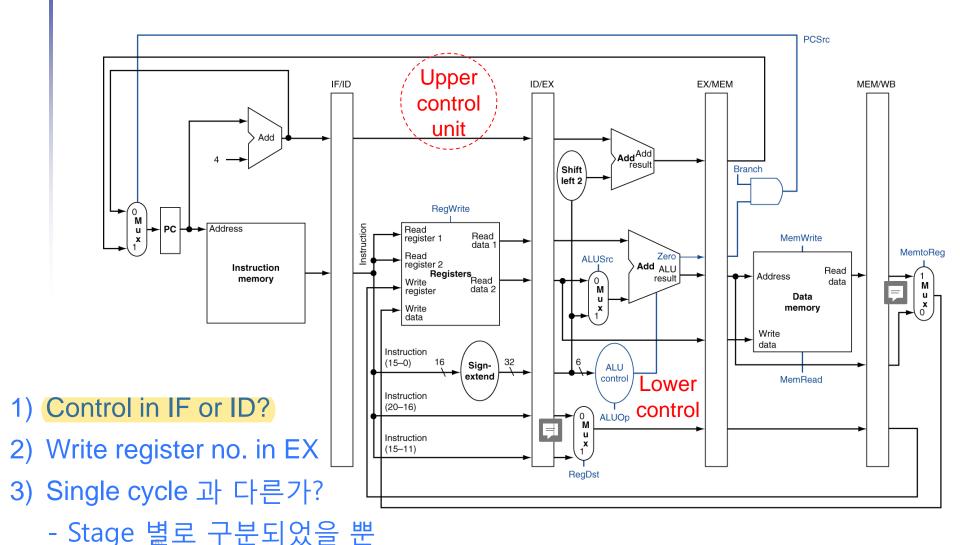
add \$14, \$5, \$6	lw \$13, 24 (\$1)	ac	dd \$12, \$3, \$4	sub \$11, \$2, \$3	lw \$10, 20(\$1)
Instruction fetch	Instruction decod	e	Execution	Memory	Write-back





Pipelined Control

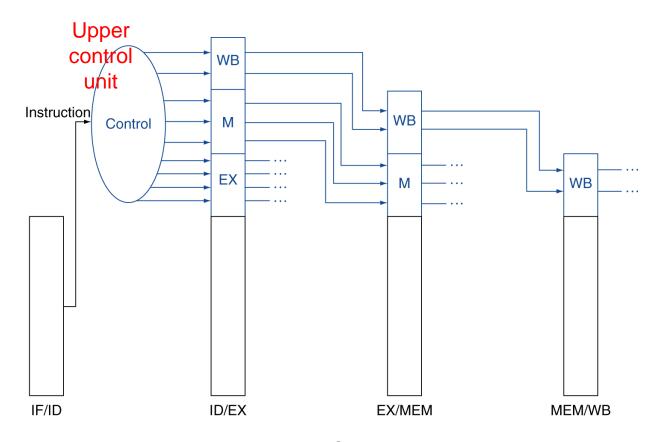
Pipelined Control (Simplified)



Pipelined Control

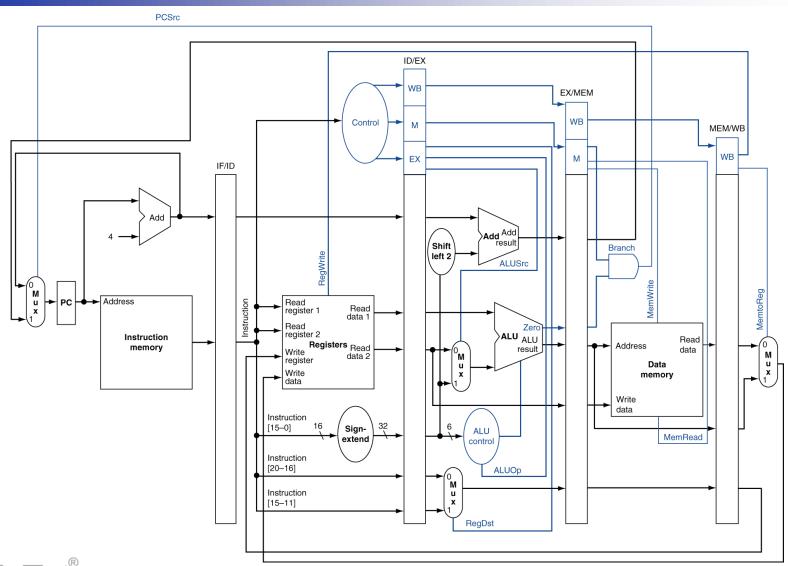
- ☐ Control signals derived from instruction
 - As in single-cycle implementation

Single cycle 과 동일한 control unit





Pipelined Control

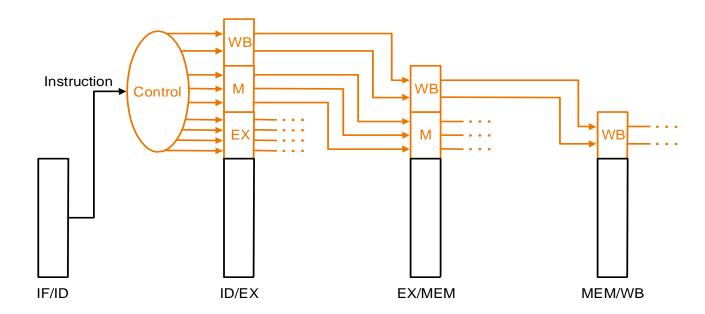




Pipeline Control

□ Same truth table we saw in single cycle (stage 별로 구분)

	Execution/Address Calculation stage control lines					y access ontrol lin	•	stage control lines		
Instruction	Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg	
R-format	1	1	0	0	0	0	0	1	0	
lw	0	0	0	1	0	1	0	1	1	
SW	Х	0	0	1	0	0	1	0	Х	
beq	Х	0	1	0	1	0	0	0	Х	



Pipelining and ISA Design (부연)

- MIPS ISA designed for pipelining
 - All instructions are 32-bits
 - Easier to fetch and decode in one cycle
 - + c.f. x86: 1- to 17-byte instructions
 - Few and regular instruction formats
 - Can decode and read registers in one step
 - Load/store addressing
 - Calculate address in 3rd, access memory in 4th stage
 - Alignment of memory operands
 - Memory access takes only one cycle





COMPUTER ORGANIZATION AND DESIGN



The Hardware/Software Interface

Chapter 4

Pipeline Hazards

(Problem-solving for powerful machines)

Hazards (미리보기)

- ☐ Situations that prevent starting the next instruction in the next cycle
 - 1) Structural hazards
 - A required resource is busy
 - 2) Data hazards
 - Need to wait for previous instruction to complete its data read/write
 - 3) Control (or branch) hazards
 - Deciding on control action depends on previous instruction



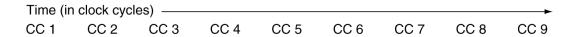
Structure Hazards

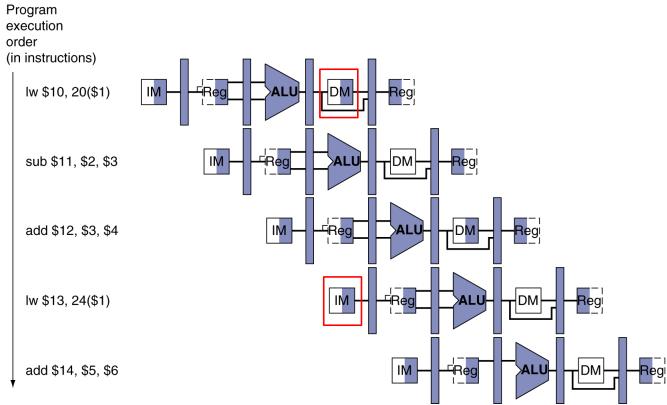
- ☐ Conflict for use of a resource
 - Different stages want to use same functional unit
- ☐ In MIPS pipeline with a single memory
 - Load/store requires data access
 - Instruction fetch would have to stall for that cycle
 - Would cause a pipeline "bubble"
- ☐ Hence, pipelined datapaths require separate instruction/data memories
 - Or separate instruction/data caches



Structural Hazards

□ Integrated IM and DM – impact on CPI? $(1 \rightarrow 1.2)$







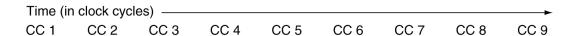
Structural Hazards

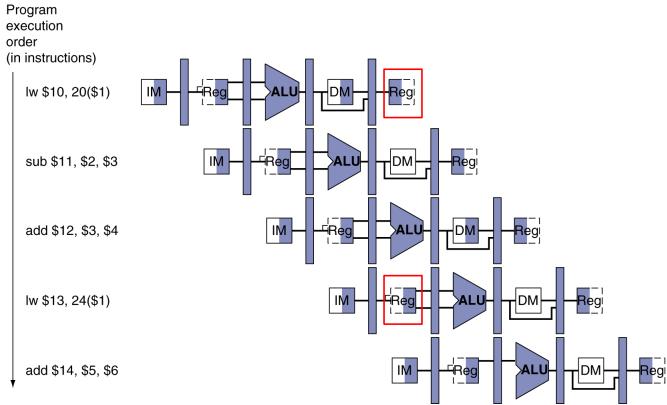
- ☐ Different stages want to use same functional unit
 - Can avoid with harwdare duplication
- But what about register access?
 - ID vs. write back
 - Divide a clock period into two halves
 - Register access: 100 ps, other operations: 200 ps
- Our datapath has no structural hazards
- Why the name "structural" hazards?



Structural Hazards

☐ Register access: divide a clock period into two halves



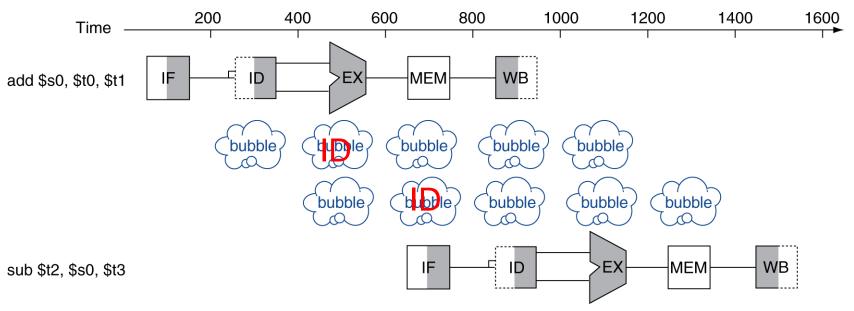




Data Hazards

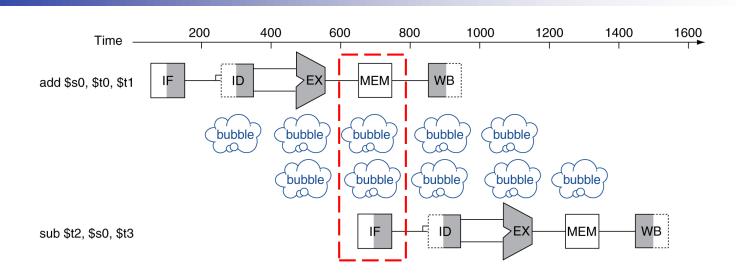
Data Hazards

□ An instruction depends on completion of data access by a previous instruction

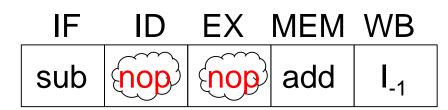




Pipeline Stall, Bubble



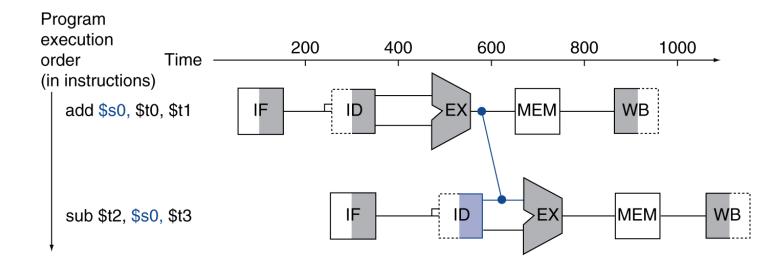
- □ Lose two cycles: increase in CPI
 - What is bubble: nop instruction (e.g., sll \$0, \$0, 0)
 - All it does is to increment PC by 4





Forwarding (aka Bypassing)

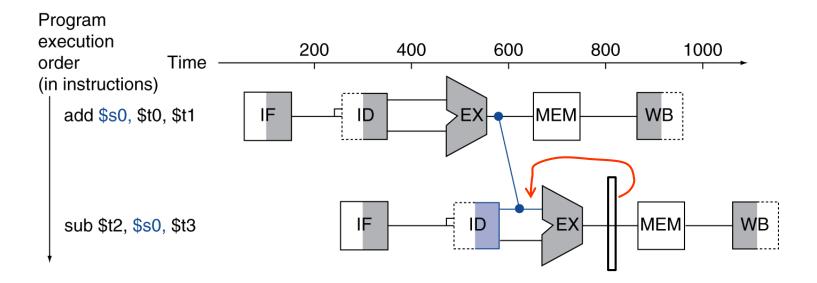
- Use result when it is computed
 - Don't wait for it to be stored in a register
 - Requires extra connections in the datapath





Forwarding (aka Bypassing)

☐ Add feedback datapath from EX/MEM pipeline register





Data Hazards and Forwarding (1)

Data Hazards in ALU Instructions

☐ Consider this sequence:

```
sub $2, $1,$3
and $12,$2,$5
or $13,$6,$2
add $14,$2,$2
sw $15,100($2)
```

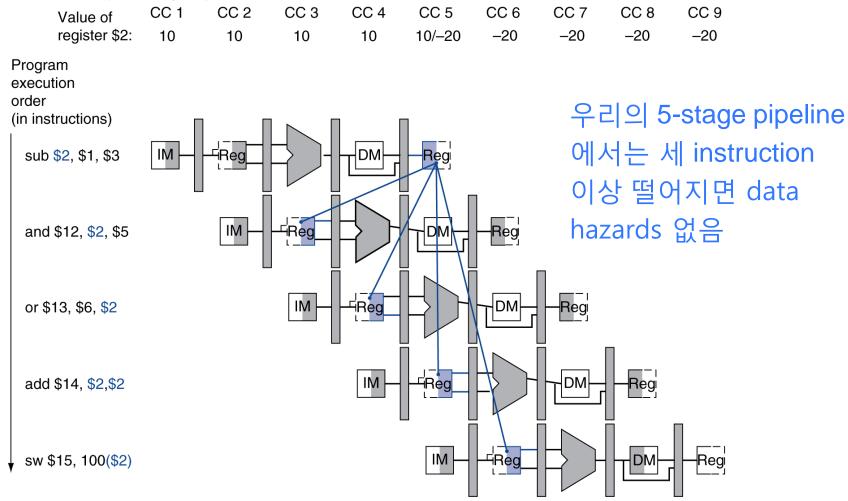
Data dependence vs. parallelism

- □ RAW (read after write) data dependence (프로그램 특성)
- We can resolve hazards with forwarding
 - How do we detect when to forward?



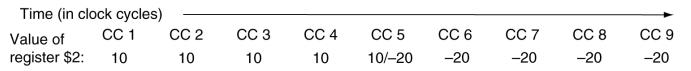
RAW Data Dependence

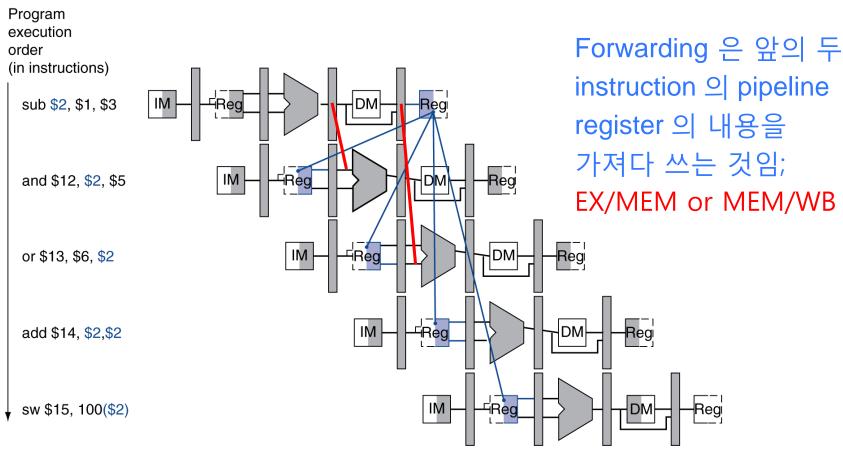
Time (in clock cycles)





Dependence and Forwarding

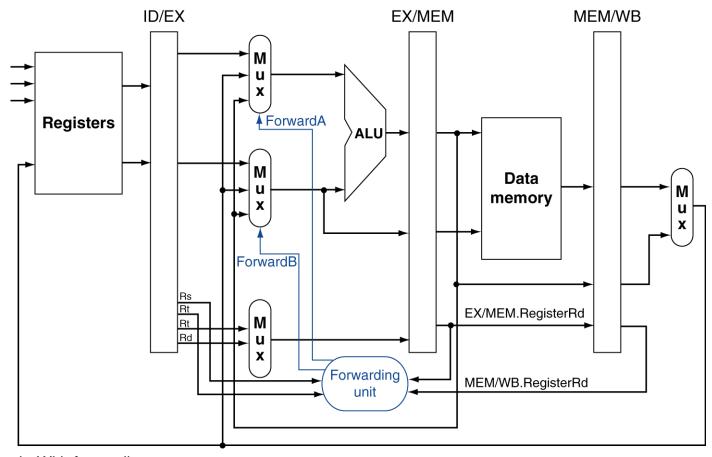






Datapath for Forwarding

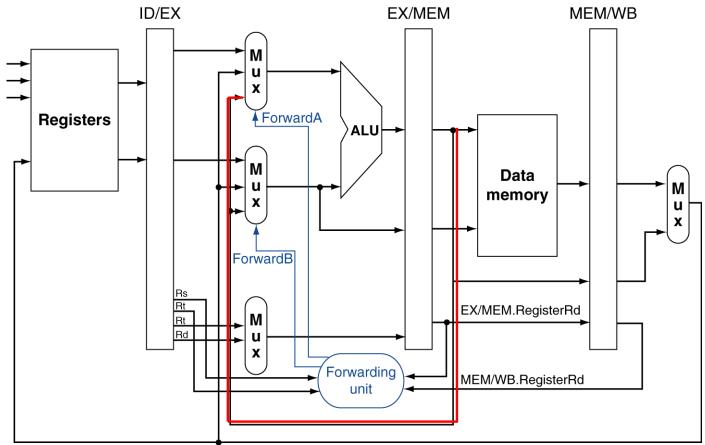
More datapath and control

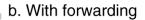




Datapath for Forwarding

☐ Case study: forwarding from EX/MEM to ALU upper input







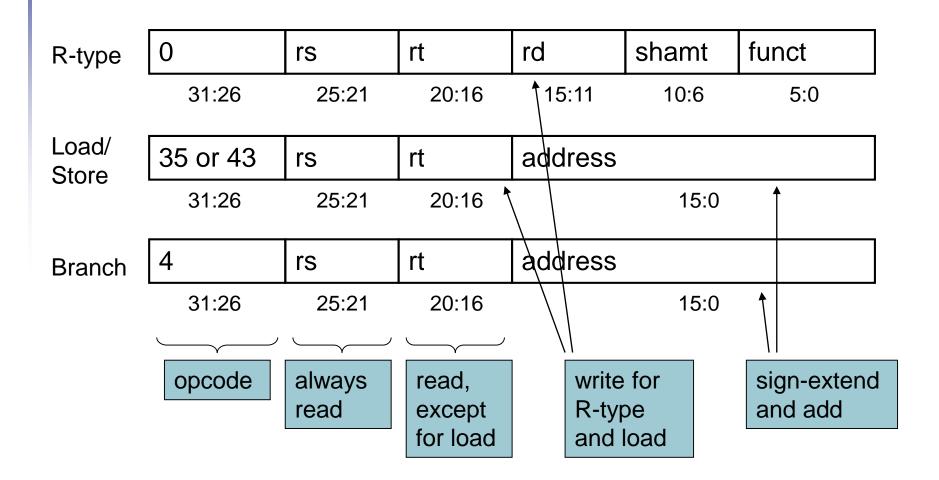
Detecting EX Hazard (condition 1)

- □ Register number for ALU upper operand
 - ID/EX.RegisterRs
- □ 바로 앞선 instruction 이 결과를 쓰는 register
 - EX/MEM.RegisterRd
- Data hazards when
 - EX/MEM.RegisterRd = ID/EX.RegisterRs

Fwd from EX/MEM pipeline reg



MIPS Instruction Format (반복)



Detecting EX Hazard (condition 2)

- But only if forwarding instruction will write to a register!
 - EX/MEM.RegWrite = 1
 - ALU or lw instructions
- Some instructions don't forward data to others
 - They don't write data to registers

(or they don't compute new register values)

beq \$t1, \$t2, 4

sw \$t1, 8(\$gp)



Detecting EX Hazard (condition 3)

- ☐ And only if Rd for forwarding instruction is not \$zero
 - EX/MEM.RegisterRd ≠ 0
- □ 컴파일러는 필요시 destination 이 \$zero 인 instruction 만듬 (예: sll \$0, \$0, 0)
 - 이런 instruction 으로 부터는 forwarding 받으면 안 됨



Forwarding Conditions

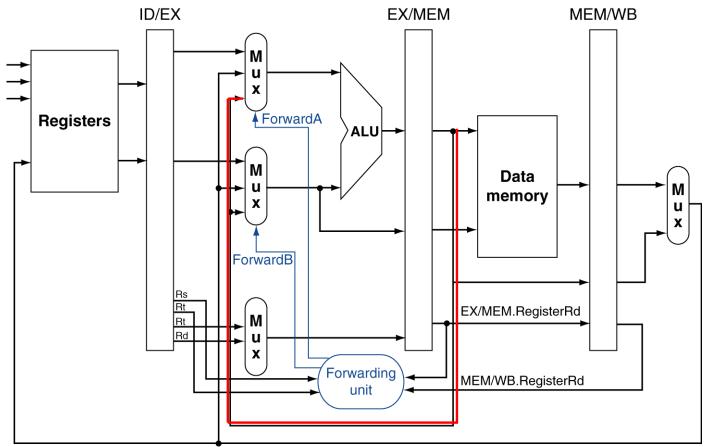
- ☐ Forwarding from EX/MEM to ALU upper input
 - if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)
 and (EX/MEM.RegisterRd = ID/EX.RegisterRs))

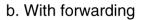
ForwardA = 10



Datapath for Forwarding (반복)

□ Case study: forwarding from EX/MEM to ALU upper input







Data Hazards and Forwarding (2)

Forwarding Conditions

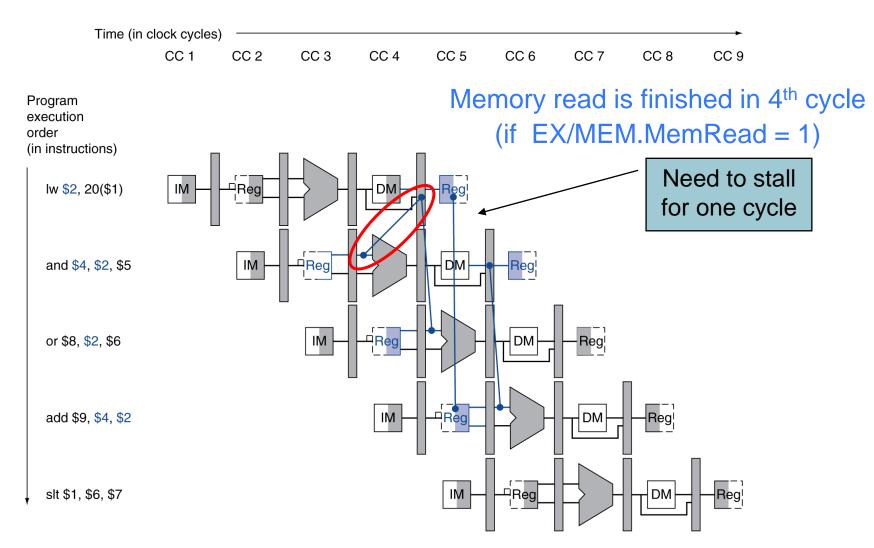
- □ Forwarding from EX/MEM to ALU upper input (반복)
 - if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)
 and (EX/MEM.RegisterRd = ID/EX.RegisterRs))

ForwardA = 10

- What is missing?
 - If "load" instruction in MEM stage, cannot forward
 - if EX/MEM.MemRead = 1



Load-Use Data Hazard (미리보기)





Control for Forwarding

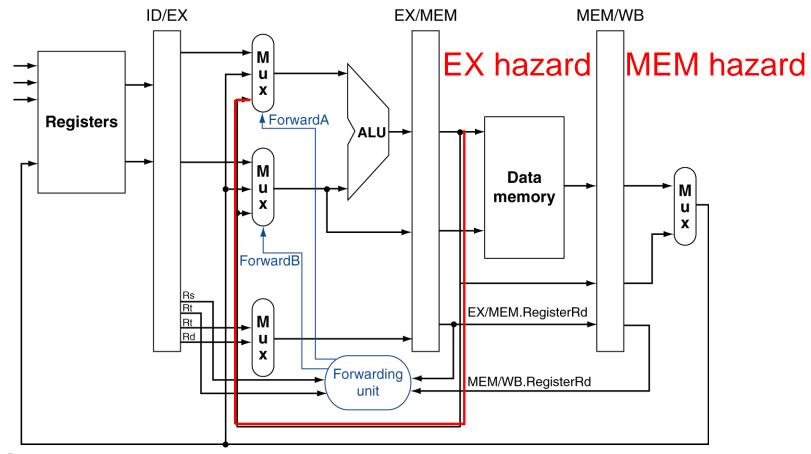
- EX hazard (if EX/MEM.MemRead = 0)
- We did it!
- if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0) and (EX/MEM.RegisterRd = ID/EX.RegisterRs))
 ForwardA = 10
- ALU •
 Iower input
- if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)
 input and (EX/MEM.RegisterRd = ID/EX.RegisterRt))
 - ForwardB = 10
 - ☐ MEM hazard
 - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0) and (MEM/WB.RegisterRd = ID/EX.RegisterRs))
 ForwardA = 01
 - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0) and (MEM/WB.RegisterRd = ID/EX.RegisterRt))

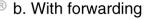




Forwarding Paths (반복)

■ Designed forwarding from EX/MEM to ALU upper input





Thinking about Forwarding

- Source of forwarding: EX/MEM or MEM/WB
 - Instructions that produce new register values
 - Producer: **ALU** instructions, **Iw**
 - Non-producer: sw, beq
- □ Destination of forwarding: ALU inputs, write data to DM
 - Where data is consumed
 - Any instruction can receive forwarding
 - Forwarding to DM write data exercise!



Double Data Hazard

☐ Consider the sequence:

```
add $1,$1,$2
add $1,$1,$3
add $1,$1,$4
```

- Both hazards occur
 - Want to use the most recent
- □ Revise MEM hazard condition
 - Only fwd if EX hazard condition isn't true



Revised Forwarding Condition(참고)

☐ MEM hazard

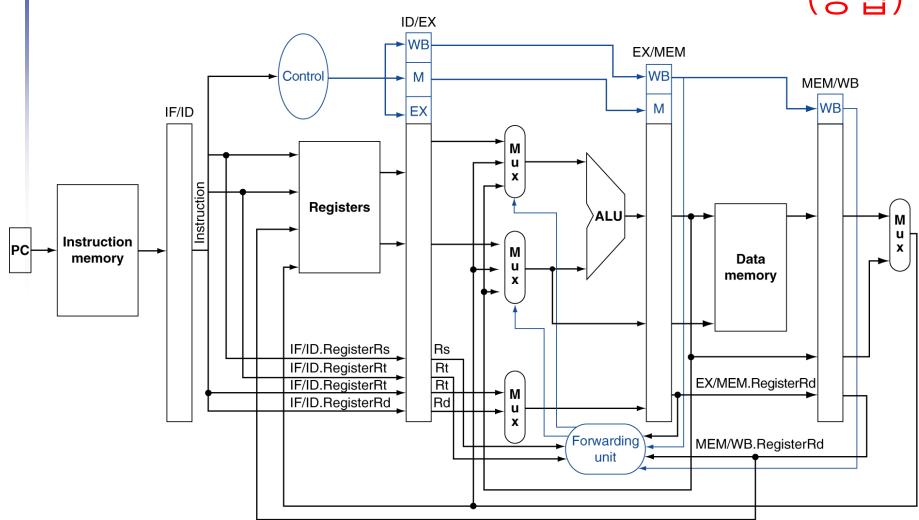
if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)

if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)



Datapath and Control with Forwarding

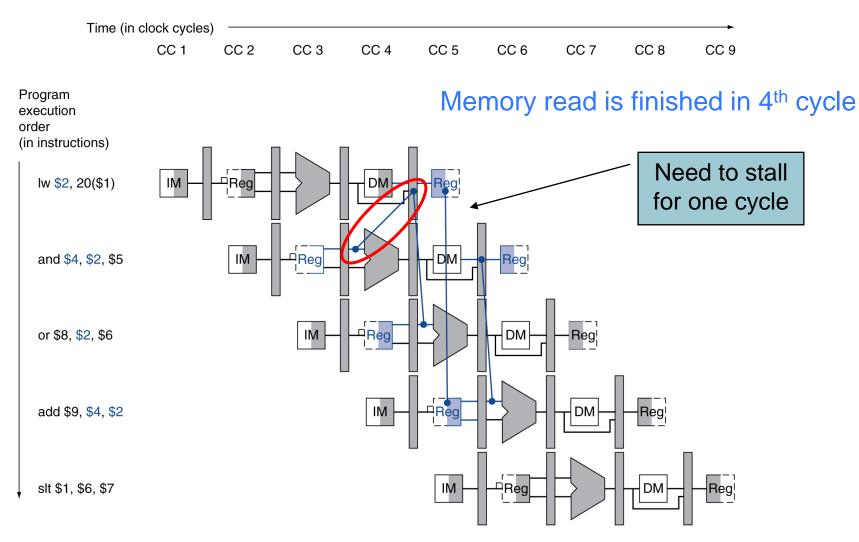
(종합)





Load-Use Data Hazard

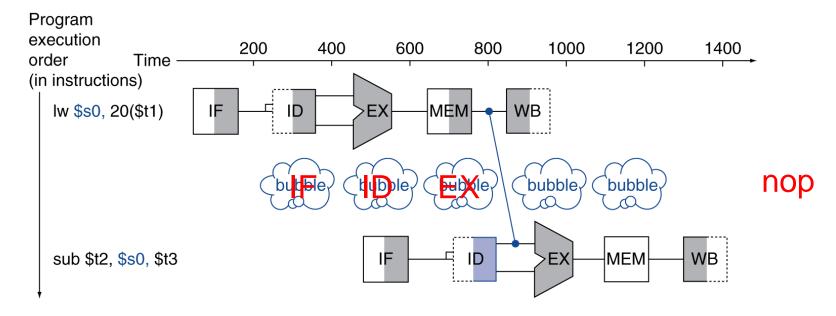
Load-Use Data Hazard (반복)





Load-Use Data Hazard

- Can't always avoid stalls by forwarding
- Only data hazard that require stall in our 5-stage design
 - CPI: $1 \to 1.05$



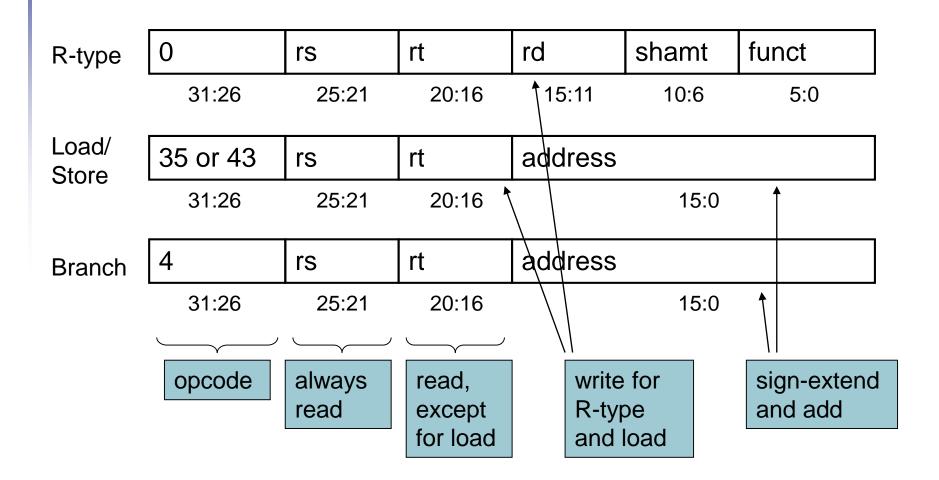


Load-Use Hazard Detection

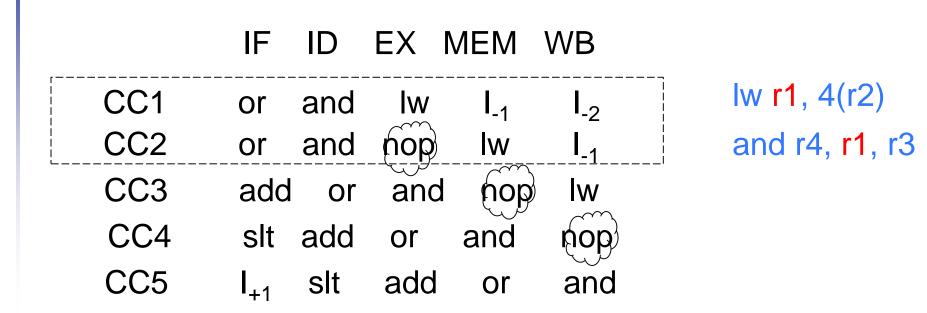
- ☐ Check when using instruction is decoded in ID stage
- □ ALU operand register numbers in ID stage are given by
 - IF/ID.RegisterRs, IF/ID.RegisterRt
- □ Load-use hazard when
 - ID/EX.MemRead and ((ID/EX.RegisterRt = IF/ID.RegisterRs) or (ID/EX.RegisterRt = IF/ID.RegisterRt))
- ☐ If detected, stall (i.e., insert **nop** instruction or bubble)



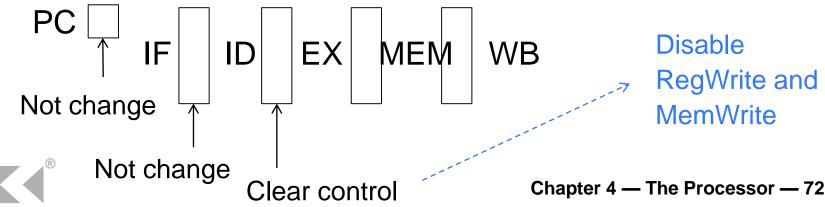
MIPS Instruction Format (반복)



How to Stall the Pipeline



To insert **nop** (or bubble), at the end of clock, do:





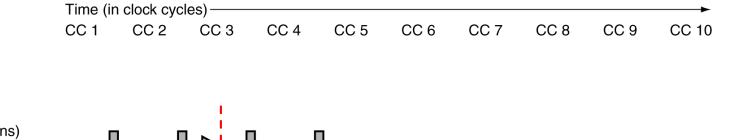


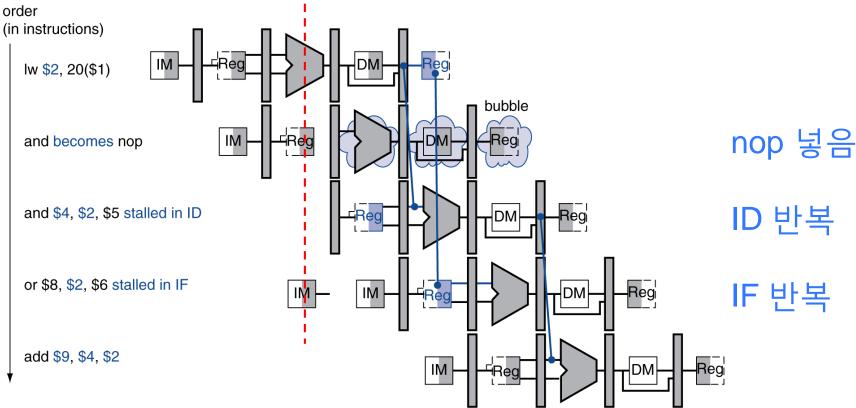
How to Stall the Pipeline (부연)

- ☐ Force control values in ID/EX register to 0
 - Disable "RegWrite" and "MemoryWrite"
 - EX executes "nop" (no-operation)
- ☐ Prevent update of PC and IF/ID register
 - ID stage: same instruction is decoded again
 - IF stage: same instruction is fetched again
- □ 1-cycle stall allows MEM to read data for \(\)\rm \(\)
 - Can subsequently forward to EX stage



Stall/Bubble in the Pipeline

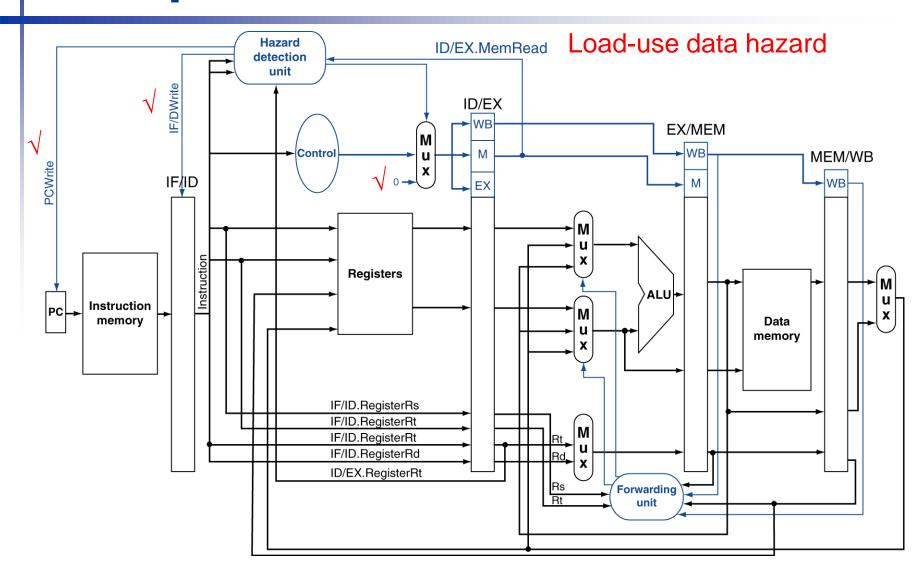






Program execution

Datapath with Hazard Detection







Data Hazards and Compilers

Software Solution

- Stalls reduce performance
 - But are required to get correct results
- Compiler can arrange code to avoid hazards and stalls



Software Solution

- □ Let compiler guarantee no hazards
 - Where do we insert **nop**'s?

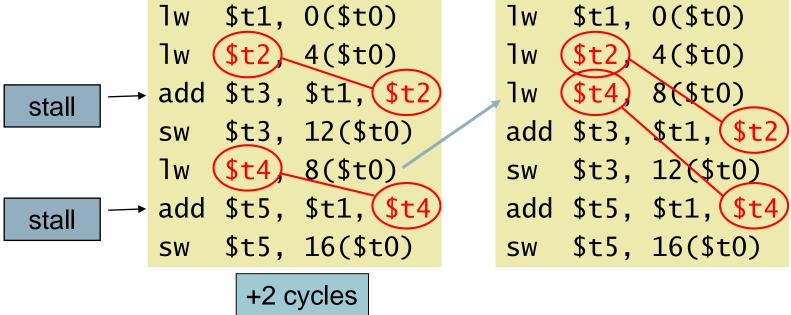
```
sub $2, $1, $3
and $12, $2, $5
or $13, $6, $2
add $14, $2, $2
sw $15, 100($2)
```

- ☐ Problem: this really slows us down!
 - Use forwarding!



Code Scheduling to Avoid Stalls

- Reorder code to avoid load-use data hazard
- \Box C code for A = B + E; C = B + F;
 - Assume forwarding





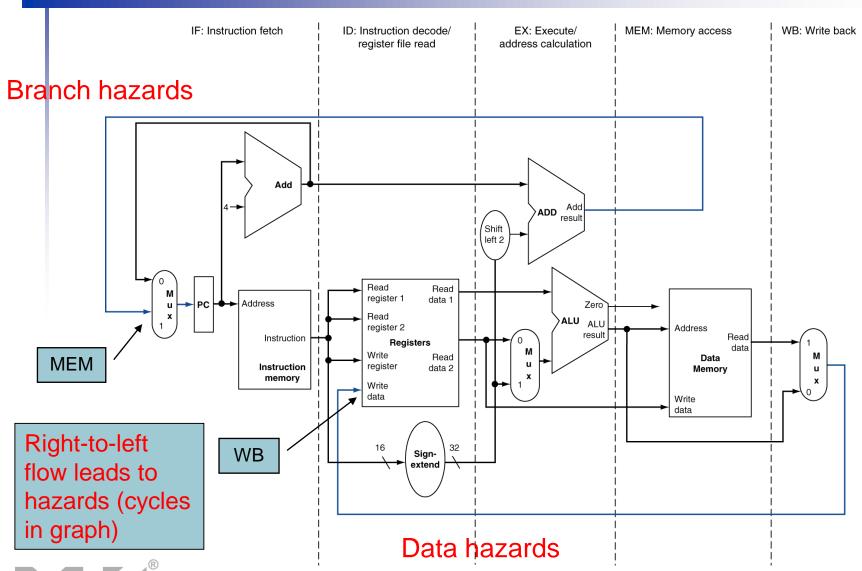
Software Solution

- ☐ Compiler can arrange code to avoid load-use stalls
 - Requires knowledge of the pipeline structure
 - What if the implementation changes?
 - Layer violation and compiler dependence
- □ Hardware can also do it without difficulty
 - Runtime (dynamic) out-of-order execution and inorder completion (Topic 4-3 참고)
 - No compiler dependence

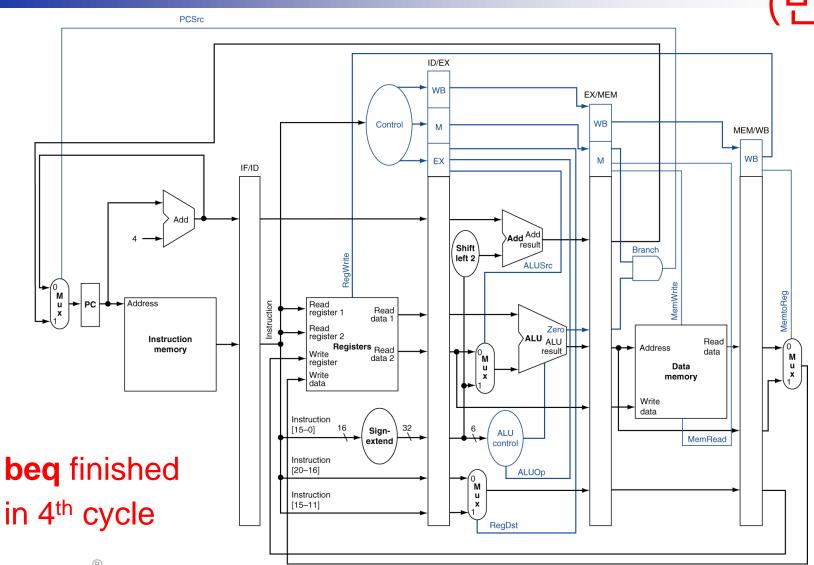


Control (or Branch) Hazards

MIPS Pipelined Datapath (반복)



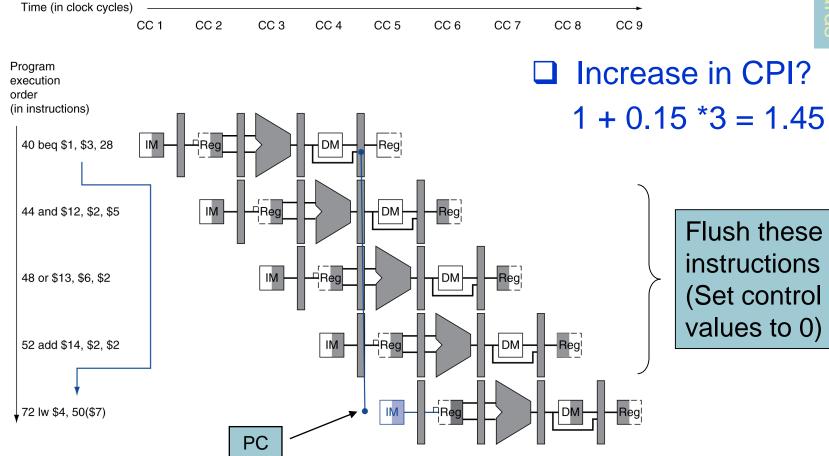
Pipelined Datapath and Control





Branch Hazards

If branch outcome determined in MEM



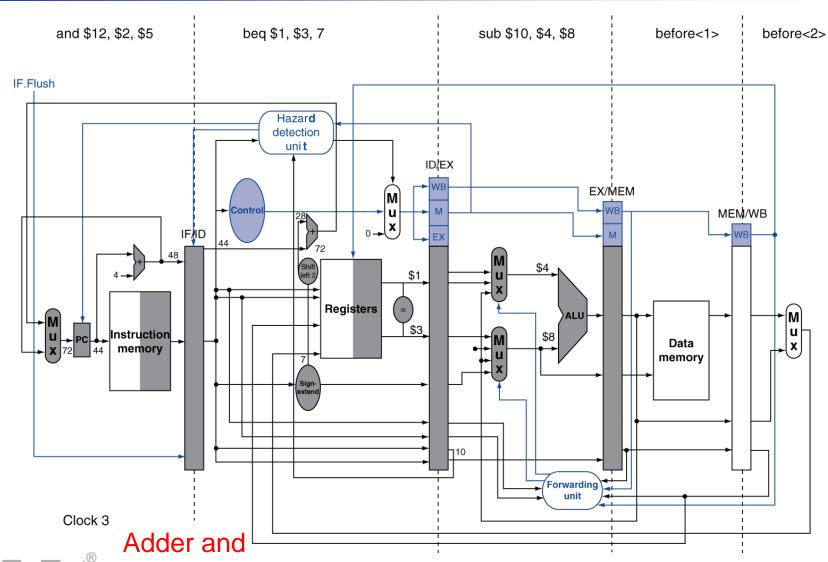


Control Hazards

- Branch determines flow of control
 - Fetching next instruction depends on branch outcome
 - Pipeline can't always fetch correct instruction
- ☐ Branch delay: 3 cycles
 - Change datapath to finish beq/bne in ID stage
 - Branch delay: 1 cycle



Moving Branch to ID Stage





comparator

Reducing "Branch Delay"

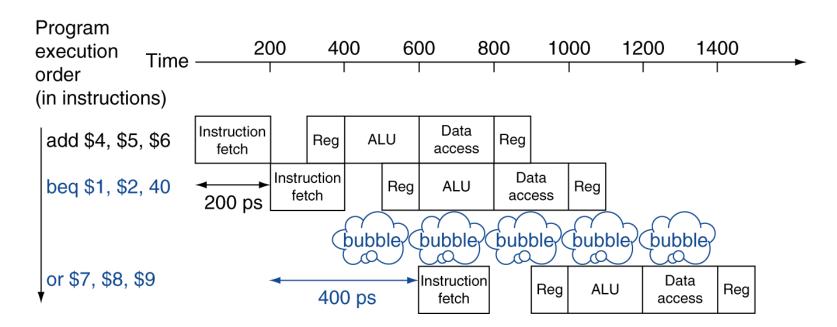
- Move hardware to determine outcome to ID stage
 - Target address adder
 - Register comparator
 - Equality check (beq/bne): bit by bit comparison
 - † Faster than subtract (**blt** → **slt** and **beq**)
- ISA design with pipelining in mind



Stall on Branch

■ Wait until branch outcome determined before fetching next instruction

☐ Increase in CPI?1 + 0.15 * 1 = 1.15



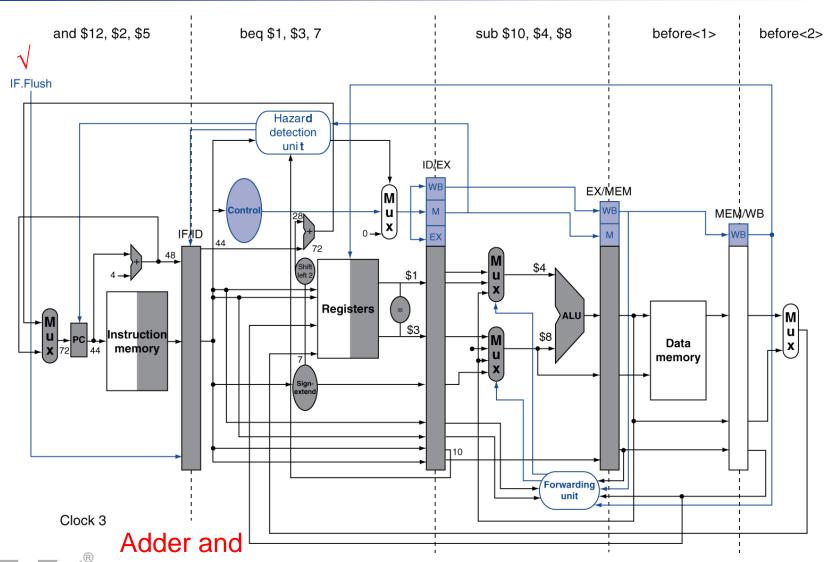


Moving Branch to ID Stage

- Evaluate branch decision in ID
 - Use IF.Flush control signal to insert NOP
 - sII \$0, \$0, 0 (all 0s; opcode & function field are 0s)
- ☐ Side effect: more forwarding and data hazards
 - Additional forwarding to register comparator
 - From EX/MEM, MEM/WB registers
 - Additional stalls
- Think about ARM conditional instruction (skip)



Moving Branch to ID Stage (반복)





comparator

Additional Data Hazards

- ☐ If a comparison register is a destination of immediately preceding ALU instruction
 - Need 1 stall cycle



Additional Data Hazards

- ☐ If a comparison register is a destination of immediately preceding load instruction
 - Need 2 stall cycles

그러나 control hazard 를 줄이는 이득이 더 큼



Branch Prediction

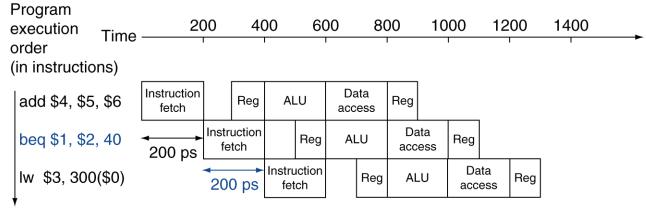
Branch Prediction

- Predict outcome of branch
 - Only stall if prediction is wrong
- ☐ In MIPS pipeline
 - Can predict branches not taken (predict not taken)
 - Fetch instruction after branch, with no delay
 - Prediction accuracy: 40%
 - Cannot use "predict taken" in our design



MIPS with Predict Not Taken

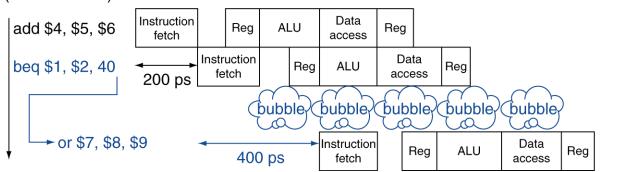
Prediction correct



☐ Increase in CPI?

1 + 0.15 * 0.6 * 1 = 1.09

Prediction incorrect





More–Realistic Branch prediction

- ❖ Advanced pipelining 에 대해 미리 공부 (Topic 4-3 참고)
 - 상세한 방법보다는 개념만 이해하면 됨

- Longer pipelines can't readily determine branch outcome early
 - Stall penalty becomes unacceptable
- □ Branch prediction is important in deep pipeline



More-Realistic Branch Prediction

- ☐ Static branch prediction (before run)
 - Based on typical branch behavior
 - Example: loop and if-statement branches
 - Predict backward branches taken
 - Predict forward branches not taken
- Dynamic branch prediction (runtime)
 - Hardware measures actual branch behavior
 - e.g., record recent history of each branch
 - Assume future behavior will continue the trend
 - If wrong, stall while re-fetching, and update history



Dynamic Branch Prediction

- ☐ Branch prediction buffer (aka branch history table)
 - Indexed by recent branch instruction addresses
 - Stores outcome (taken/not taken)
- □ Branch target buffer
 - Cache of target addresses
 - If hit and instruction is branch predicted taken, can fetch target immediately



Dynamic Branch Prediction

- ☐ To execute a branch (speculative execution)
 - Check table, expect the same outcome
 - Start fetching from fall-through or target

Addrogoo of

If wrong, flush pipeline and flip prediction

	branch instructions	or unta	Branch target
Hardware data structure			

- ☐ Hardware overhead vs. performance
 - Prediction accuracy over 90% in desktop processors

Dranch taken

1-Bit Predictor: Shortcoming

□ Inner loop branches mispredicted twice!

```
outer: ...

inner: ...

beq ..., ..., inner

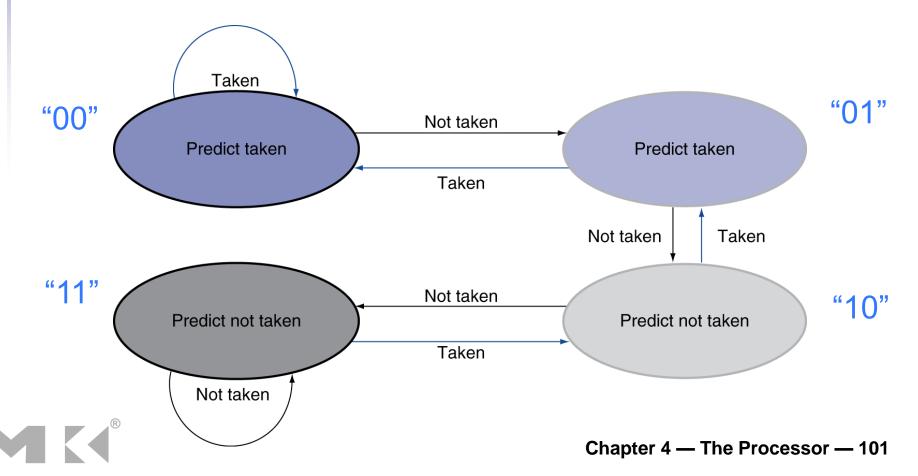
beq ..., outer
```

- Mispredict as taken on last iteration of inner loop
- Then mispredict as not taken on first iteration of inner loop next time around



2-Bit Predictor

- ☐ Only change prediction on two successive mispredictions
 - Inner loop branches mispredicted only once!



Pipeline Summary

The BIG Picture

- □ Pipelining improves performance by increasing instruction throughput
 - Executes multiple instructions in parallel
 - Each instruction has the same latency
- Subject to hazards
 - Structure, data, control
- ☐ Instruction set design affects complexity of pipeline implementation (RISC vs. CISC)



Homework #12 (see Class Homepage)

- 1) Write a report summarizing the materials discussed in Topic 4-2 (이번 주 수업 내용)
- ** 문장으로 써도 좋고 파워포인트 형태의 개조식 정리도 좋음
- 2) Solve Chapter 4 exercises 4.8, 4.12, 4.13.2, 4.13.4

- □ Due: see Blackboard
 - · Submit electronically to Blackboard

Section 4.10

Introduction to Advanced Pipelining

(From 5-Stage MIPS to Powerful MIPS for PC)