

Section 4.10

Introduction to Advanced Pipelining

(From 5-Stage MIPS to Powerful MIPS for PC)

(Optional)

Section 4.10

Instruction-Level Parallelism (ILP)

- Inside your PC

(Problem solving to develop
powerful machines)

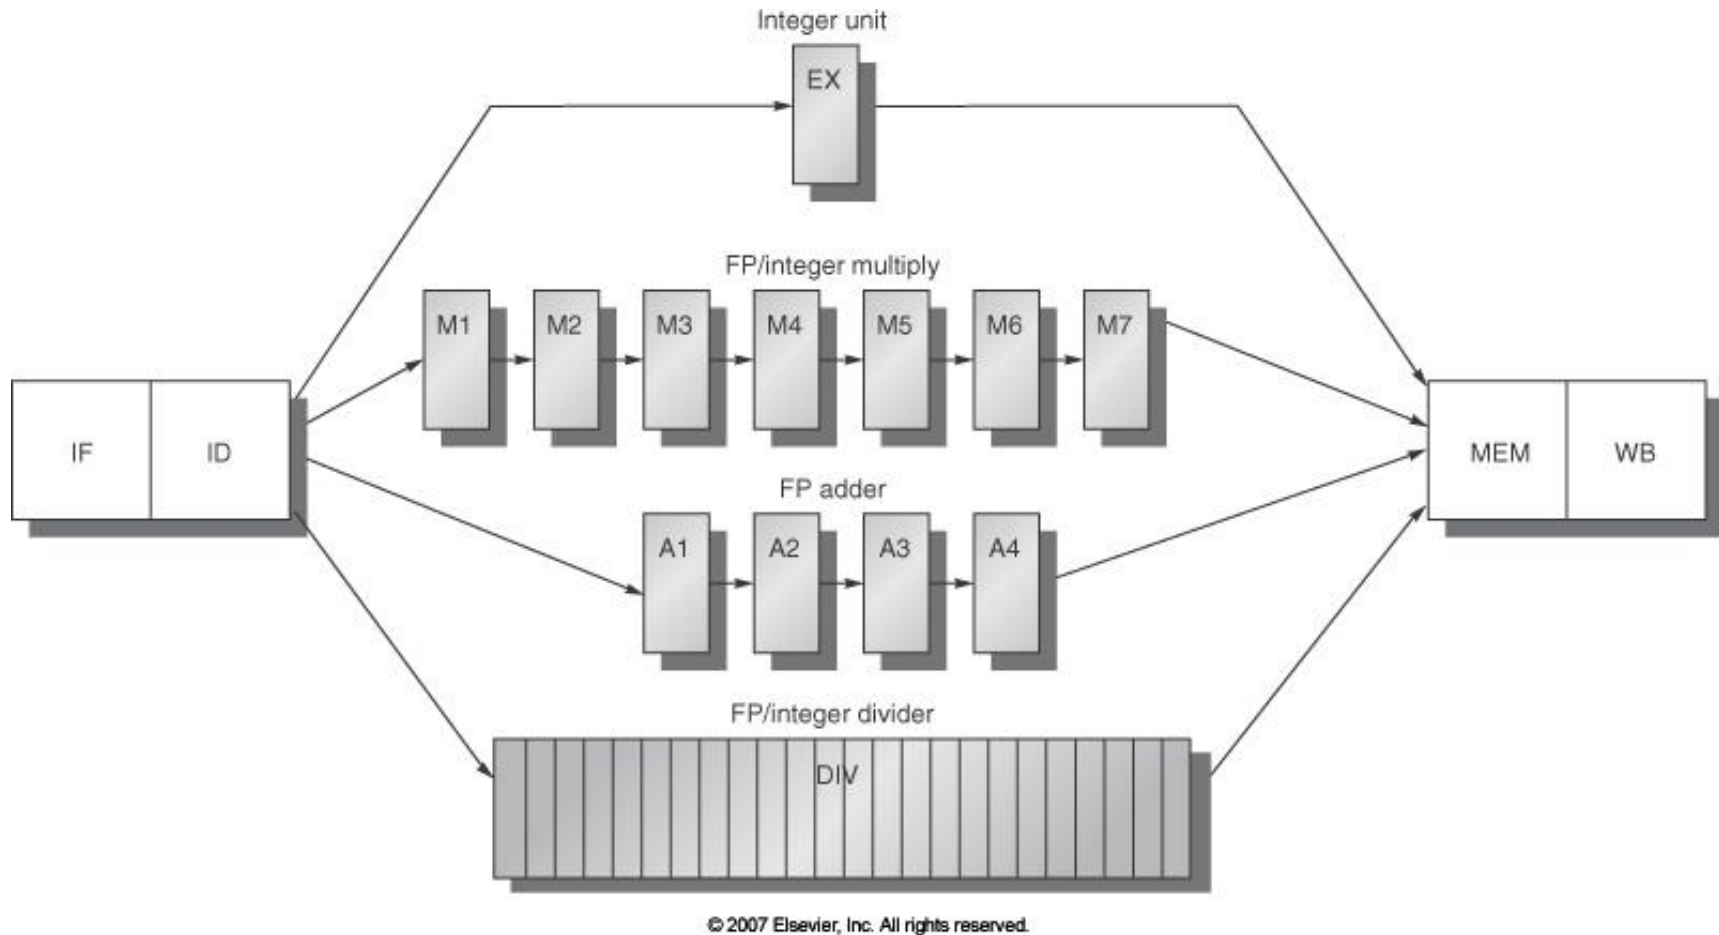
Some of authors' slides are modified

Key Concepts (다음의 용어를 이해)

- ☐ Floating-point pipelines
- ☐ Deep pipelines
- ☐ Multiple issue
 - Static (c.f., VLIW)
 - Dynamic (c.f., superscaler)
- ☐ Loop unrolling
- ☐ Speculation

❖ Computer Architecture: A Quantitative Approach,
Hennessy and Patterson

Floating-Point Pipelines

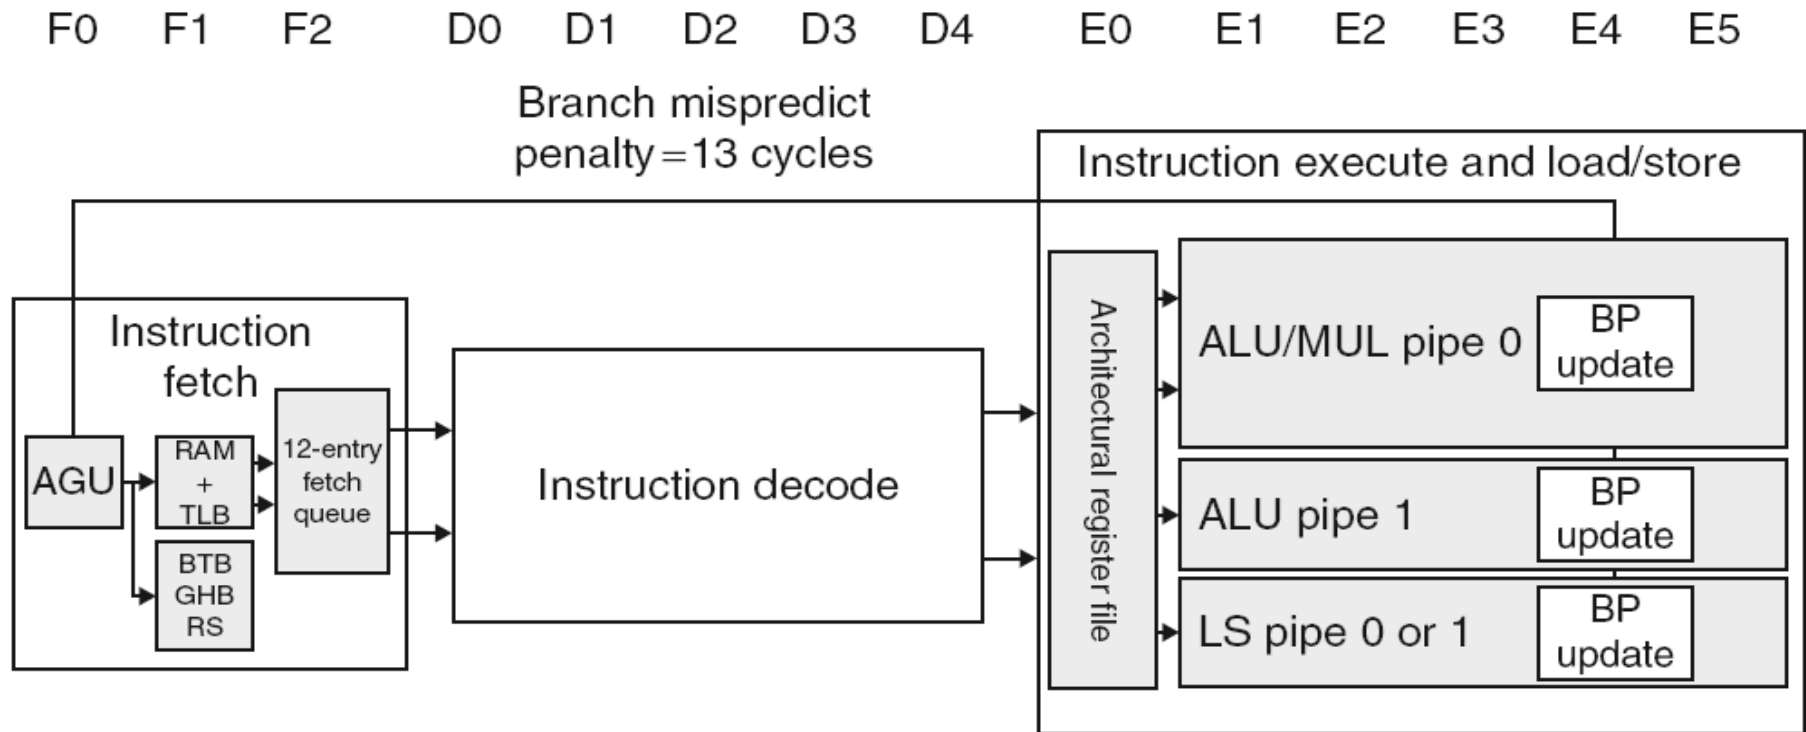


- ❑ Sometimes, better not to pipeline
 - Return on investment (e.g., divider)

Instruction-Level Parallelism (ILP)

- ❑ Pipelining: execute multiple instructions in parallel
- ❑ To increase ILP
 - Deeper pipeline (CPI = 1)
 - Less work per stage \Rightarrow shorter clock cycle
 - Multiple issue
 - Replicate pipeline stages \Rightarrow multiple pipelines
 - Start multiple instructions per clock cycle
 - † CPI = 0.25 with 4GHz 4-way multiple-issue
 - But dependencies reduce this in practice

ARM Cortex-A8 Pipeline (참고)



Key Concepts (다음의 용어를 이해)

- ❑ Floating-point pipelines
- ❑ Deep pipelines
- ❑ Multiple issue
 - Static (c.f., VLIW)
 - Dynamic (c.f., superscaler)
- ❑ Loop unrolling
- ❑ Speculation

Static Multiple Issue

- ❑ **Compiler** groups instructions that can be **issued** on a single cycle (“issue packets”)
- ❑ **Compiler** must remove some/all hazards
 - Reorder instructions; pad with “nop” if necessary

Instruction	Instruction	nop	nop
instruction	Instruction	Instruction	nop

⋮

- ❑ Think of **an issue packet** as a very long instruction
 - Specifies multiple concurrent operations
- ⇒ Very Long Instruction Word (VLIW)

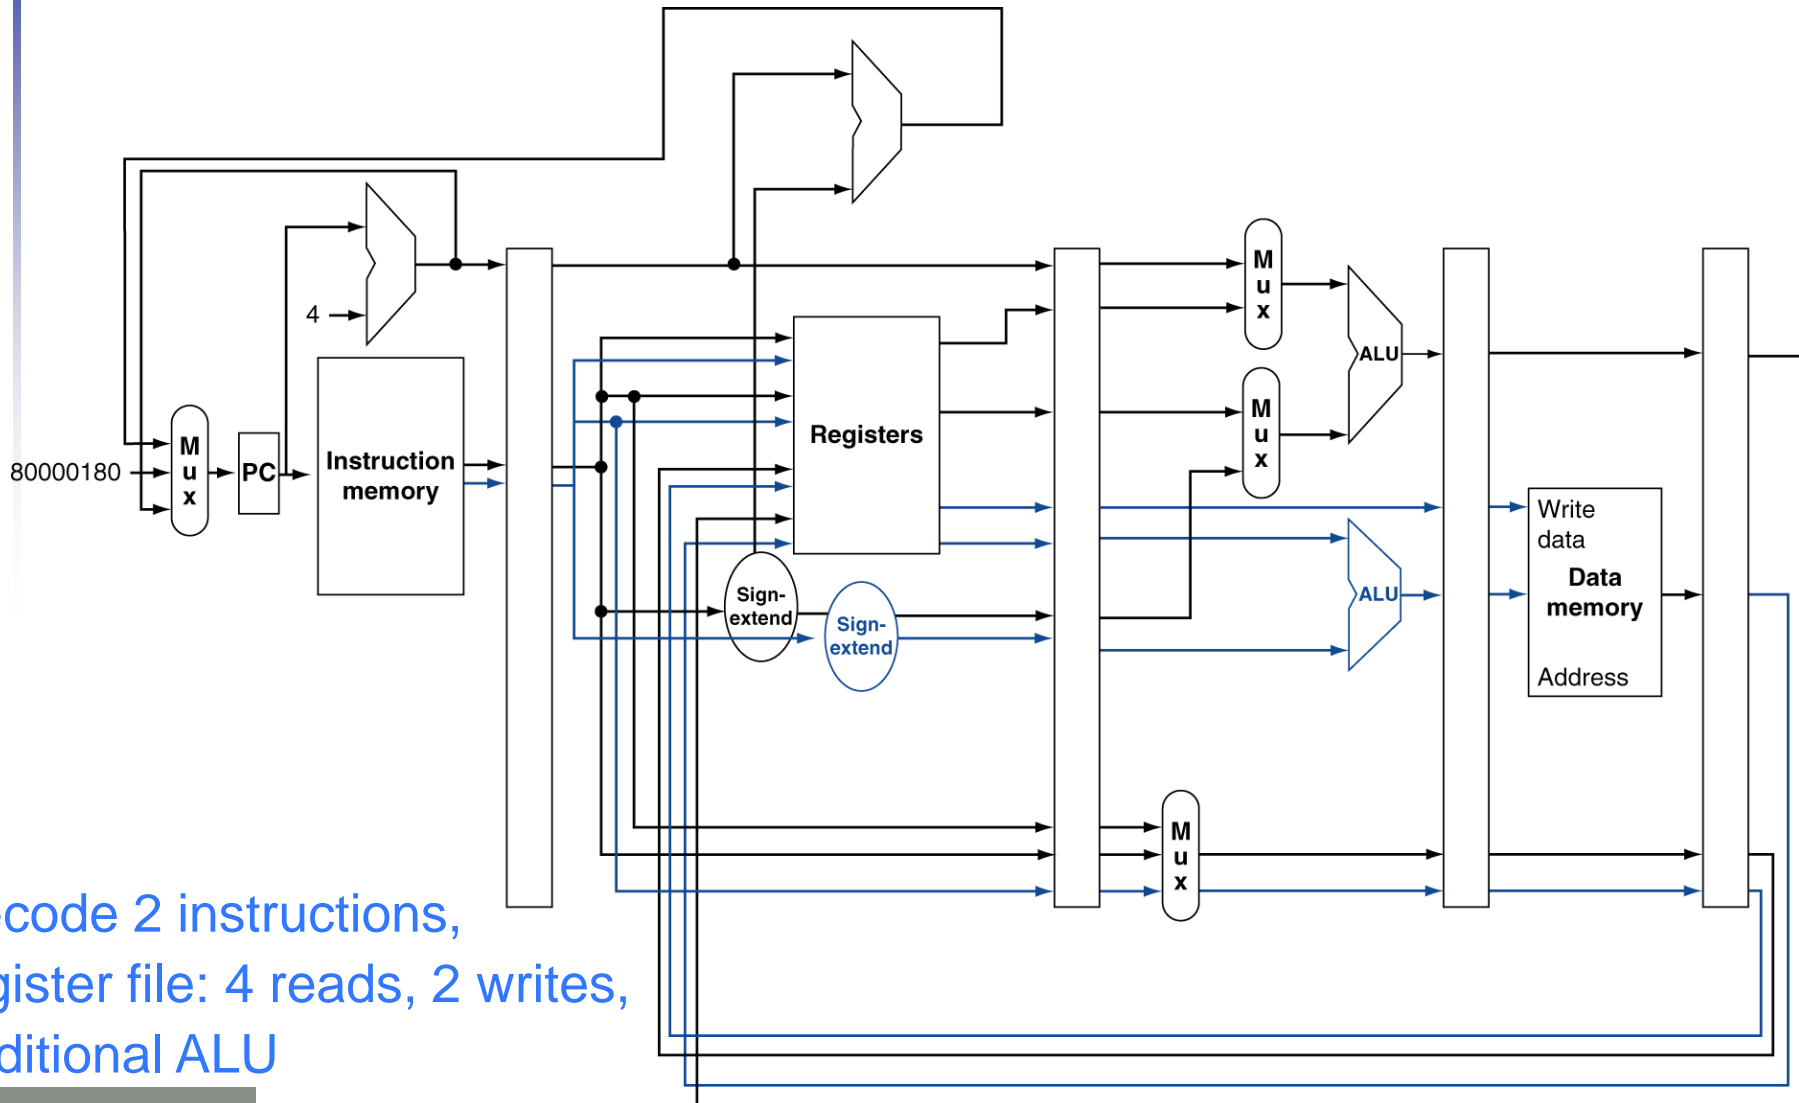
MIPS with Static Dual Issue

❑ Two-issue packets

- One ALU/branch instruction + one load/store instruction
- 64-bit aligned (pad unused instruction with “nop”)

Address	Instruction type	Pipeline Stages						
n	ALU/branch	IF	ID	EX	MEM	WB		
n + 4	Load/store	IF	ID	EX	MEM	WB		
n + 8	ALU/branch		IF	ID	EX	MEM	WB	
n + 12	Load/store		IF	ID	EX	MEM	WB	
n + 16	ALU/branch			IF	ID	EX	MEM	WB
n + 20	Load/store			IF	ID	EX	MEM	WB

MIPS with Static Dual Issue



Decode 2 instructions,
register file: 4 reads, 2 writes,
additional ALU

Hazards in the Dual-Issue MIPS

(참고)

- ❑ More instructions executing in parallel
- ❑ Load-use hazard
 - One cycle use latency, but now two instructions
- ❑ EX data hazard
 - Can't use ALU result in load/store in same packet

addu **\$t0**, \$t0, \$s2

sw **\$t0**, 4(\$s1)

- Split into two packets, effectively a stall

- ❑ More aggressive scheduling required

Scheduling Example

Loop: $a[i] = a[i] + k$

```
Loop: lw    $t0, 0($s1)      # $t0=array element
      addu  $t0, $t0, $s2     # add scalar in $s2
      sw    $t0, 0($s1)      # store result
      addi  $s1, $s1, -4      # decrement pointer
      bne   $s1, $zero, Loop # branch $s1!=0
```

	ALU/branch	Load/store	cycle
Loop:	nop	lw \$t0, 0(\$s1)	1
	addi \$s1, \$s1, -4	nop	2
	addu \$t0, \$t0, \$s2	nop	3
	bne \$s1, \$zero, Loop	sw \$t0, 4(\$s1)	4

□ $IPC = 5/4 = 1.25$ (c.f. peak $IPC = 2$)

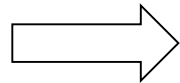
Loop Unrolling

- ❑ Replicate loop body to expose more parallelism
 - Reduces loop-control overhead
 - By compilers or programmers
- ❑ Use different registers per replication
 - Called “register renaming” (dependence 발생 방지)
 - Iteration 1: use \$t0
 - Iteration 2: use \$t1
 - Iteration 3: use \$t2
 - Iteration 4: use \$t3

Loop Unrolling

- ❑ Strip-mining or loop sectioning
 - Fragmenting a large loop into sections or strips

```
for (i = 0; i < 10000; i++)  
    a[i] += k;
```



```
for (i = 0; i < 10000; i += 4)  
{  
    a[i] += k;  
    a[i+1] += k;  
    a[i+2] += k;  
    a[i+3] += k;  
}  
// remaining iterations if any
```

Loop Unrolling Example

	ALU/branch	Load/store	cycle
Loop:	addi \$s1 , \$s1, -16	lw \$t0 , 0(\$s1)	1
	nop	lw \$t1 , 12(\$s1)	2
	addu \$t0 , \$t0 , \$s2	lw \$t2 , 8(\$s1)	3
	addu \$t1 , \$t1 , \$s2	lw \$t3 , 4(\$s1)	4
	addu \$t2 , \$t2 , \$s2	sw \$t0 , 16(\$s1)	5
	addu \$t3 , \$t4 , \$s2	sw \$t1 , 12(\$s1)	6
	nop	sw \$t2 , 8(\$s1)	7
	bne \$s1 , \$zero, Loop	sw \$t3 , 4(\$s1)	8

□ $IPC = 14/8 = 1.75$

- Closer to 2, but at cost of registers and code size

Key Concepts (다음의 용어를 이해)

- ❑ Floating-point pipelines
- ❑ Deep pipelines
- ❑ Multiple issue
 - Static (c.f., VLIW)
 - Dynamic (c.f., superscaler)
- ❑ Loop unrolling
- ❑ Speculation

Dynamic Multiple Issue

- ❑ “Superscalar” processors (why the name?)
- ❑ CPU decides whether to issue 0, 1, 2, ... instructions each cycle (at runtime)
 - Avoiding structural and data hazards
- ❑ Avoids the need for compiler scheduling
 - Though it may still help
 - Code semantics ensured by the CPU

Dynamic Pipeline Scheduling

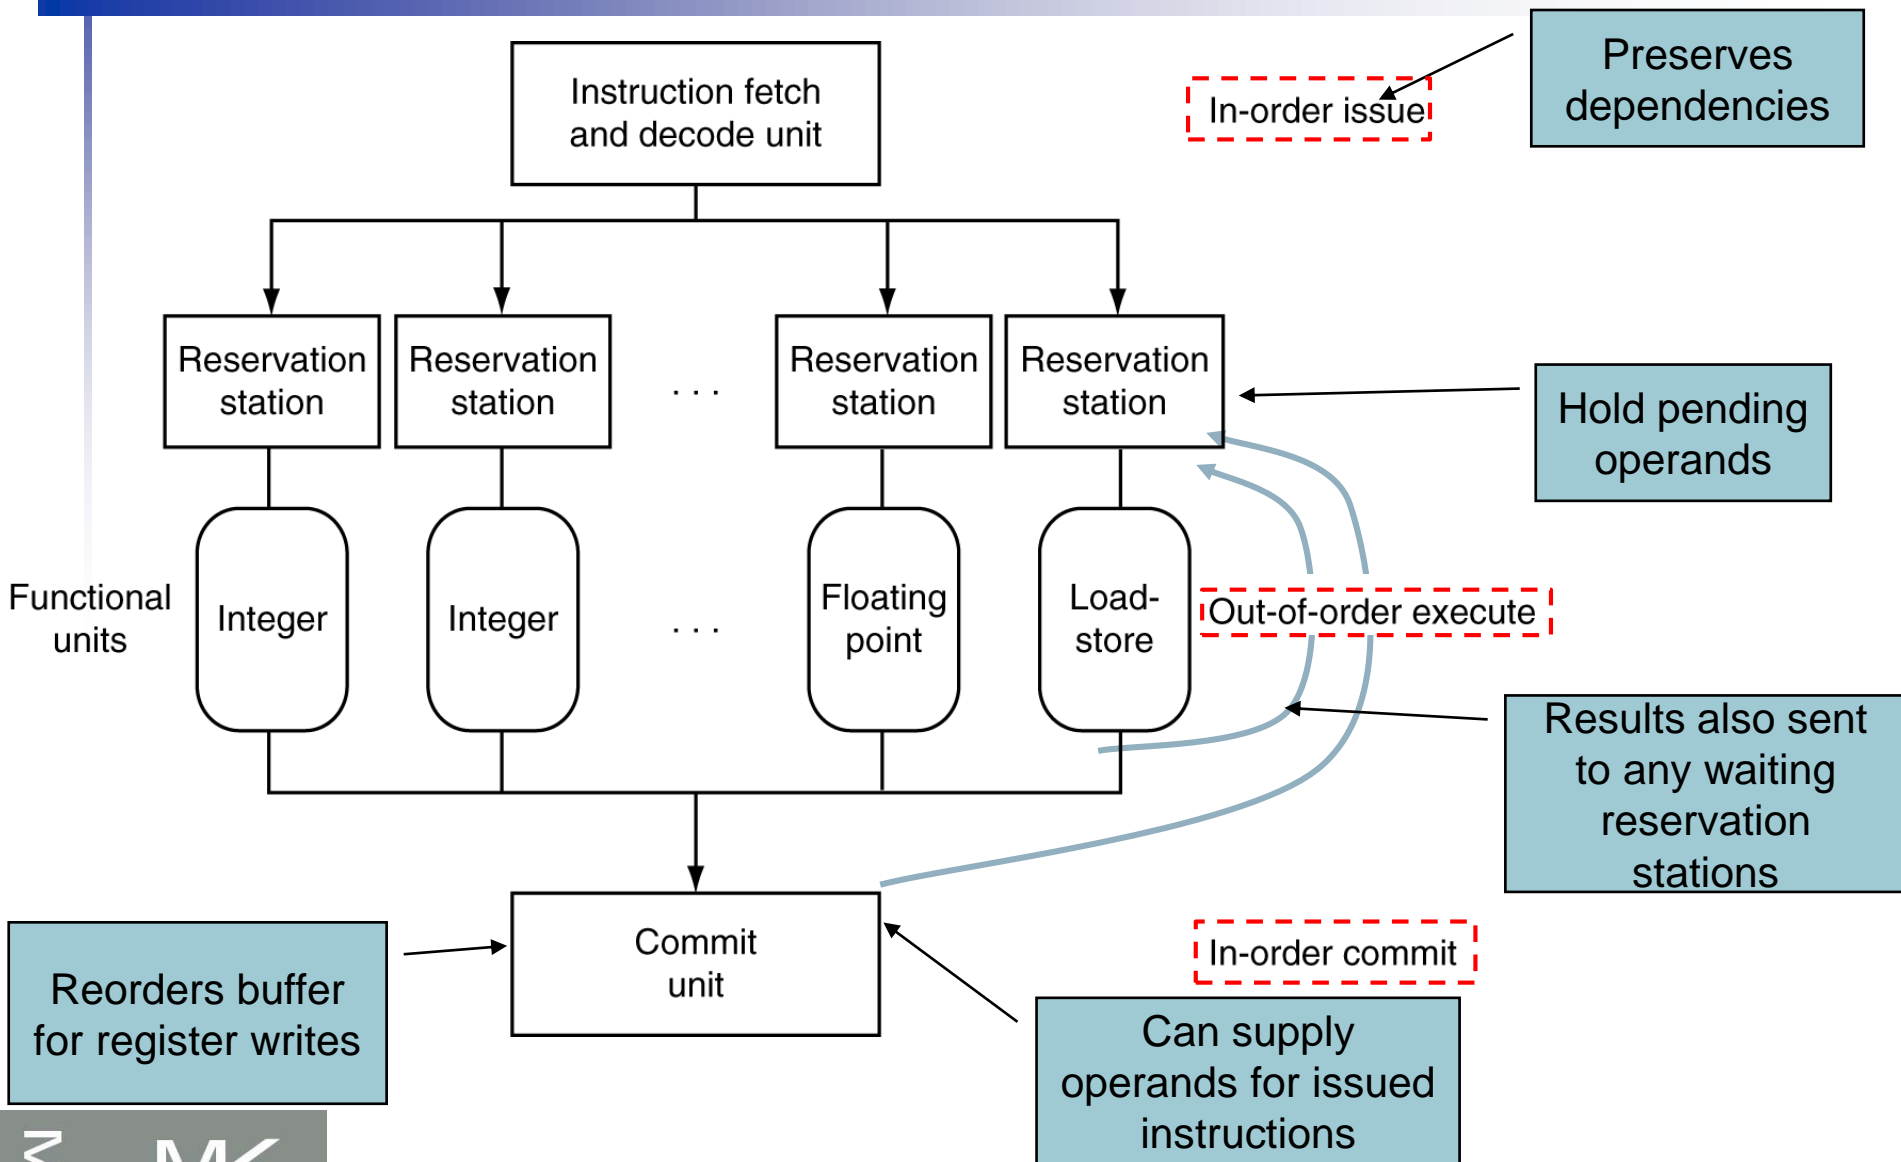
- ❑ CPU executes instructions out of order to avoid stalls
 - But commit result to registers in order

- ❑ Example

```
lw      $t0, 20($s2)
addu    $t1, $t0, $t2
sub      $s4, $s4, $t3
slti    $t5, $s4, 20
```

- Can start **sub** while **addu** is waiting for **lw**
 - But **sub** doesn't commit until **lw** and **addu** commit

Dynamic Pipeline Scheduling



Register Renaming (참고)

- ❑ Reservation stations and reorder buffer effectively provide register renaming
- ❑ On instruction issue to reservation station
 - If operand available in register file or reorder buffer
 - Copied to reservation station
 - If operand is not yet available
 - Will be provided to reservation station by a FU
 - Name of the functional unit is tracked

Why Do Dynamic Scheduling?

- ❑ Why not just let the compiler schedule code?
 - Not all stalls are predictable
 - e.g., cache misses
 - Can't always schedule around branches
 - Branch outcome is dynamically determined
 - Different implementations of an ISA have different latencies and hazards
 - No compiler dependence

Does Multiple Issue Work?

- ❑ Yes, but not as much as we'd like
 - Programs have real dependencies that limit ILP
 - Some dependencies are hard to eliminate
 - e.g., pointer aliasing
 - Some parallelism is hard to expose
 - Limited window size during instruction issue
 - Memory delays and limited bandwidth
 - Hard to keep pipelines full
 - Speculation can help if done well

Power Efficiency

- ❑ Complexity of dynamic scheduling and speculations requires power
- ❑ Multiple simpler cores may be better

Microprocessor	Year	Clock Rate	Pipeline Stages	Issue width	Out-of-order/ Speculation	Cores	Power
i486	1989	25MHz	5	1	No	1	5W
Pentium	1993	66MHz	5	2	No	1	10W
Pentium Pro	1997	200MHz	10	3	Yes	1	29W
P4 Willamette	2001	2000MHz	22	3	Yes	1	75W
P4 Prescott	2004	3600MHz	31	3	Yes	1	103W
Core	2006	2930MHz	14	4	Yes	2	75W
UltraSparc III	2003	1950MHz	14	4	No	1	90W
UltraSparc T1	2005	1200MHz	6	1	No	8	70W

Key Concepts (다음의 용어를 이해)

- ❑ Floating-point pipelines
- ❑ Deep pipelines
- ❑ Multiple issue
 - Static (c.f., VLIW)
 - Dynamic (c.f., superscaler)
- ❑ Loop unrolling
- ❑ Speculation

Speculation

- ❑ “Guess” what to do with an instruction
 - Start operation as soon as possible
 - Check whether guess was right
 - If so, complete the operation
 - If not, roll-back and do the right thing
- ❑ Common to static and dynamic multiple issue
- ❑ **Speculate on branch outcome**

Roll back if path taken is different
- ❑ Speculate on load
 - Roll back if location is updated

Class Topics (클래스 홈페이지 참조)

- ❑ Part 1: Fundamental concepts and principles
- ❑ Part 2: 빠른 컴퓨터를 위한 ISA design
 - Topic 1 Computer performance and ISA design (Ch. 1)
 - Topic 2 RISC (MIPS) instruction set (Chapter 2)
 - Topic 3 Computer arithmetic and ALU (Chapter 3)
- ❑ Part 3: ISA 의 효율적인 구현 (pipelining, cache memory)
 - Topic 4 Processor design (Chapter 4)
 - Topic 5 Memory system design (Chapter 5)
 - Cache memory
 - Cache memory and virtual memory

Section 4.9

Exceptions (skip)

Exceptions and Interrupts

- ❑ “Unexpected” events requiring change in flow of control
 - Different ISAs use the terms differently
- ❑ Exception
 - Arises within the CPU
 - e.g., undefined opcode, overflow, syscall, ...
- ❑ Interrupt
 - From an external I/O controller
- ❑ Dealing with them without sacrificing performance is hard

Section 4.11

**Real Stuff: The ARM
Cortex-A8 and Intel Core i7
Pipelines**

(skip)

ARM Cortex A8 and Intel i7

Processor	ARM A8	Intel Core i7 920
Market	Personal Mobile Device	Server, cloud
Thermal design power	2 Watts	130 Watts
Clock rate	1 GHz	2.66 GHz
Cores/Chip	1	4
Floating point?	No	Yes
Multiple issue?	Dynamic	Dynamic
Peak instructions/clock cycle	2	4
Pipeline stages	14	14
Pipeline schedule	Static in-order	Dynamic out-of-order with speculation
Branch prediction	2-level	2-level
1 st level caches/core	32 KiB I, 32 KiB D	32 KiB I, 32 KiB D
2 nd level caches/core	128-1024 KiB	256 KiB
3 rd level caches (shared)	-	2- 8 MB

Section 4.12

Instruction-Level Parallelism and Matrix Multiply

(skip)