

# Big Picture

## □ Part 3: implementation of ISA

- High-level organization, not circuits design
- Ch. 4: processor
  - Given ISA, what is a good implementation?
  - Datapath and control, pipelining
- Ch. 5: memory system design
  - 1) Memory systems: physical and virtual
  - 2) Memory hierarchy
  - 3) Cache memory: structure, operation and performance
  - 4) Cache and virtual memory

# **Chapter 5**

## **Large and Fast: Exploiting Memory Hierarchy**

### **Part 2**

Some of authors' slides are modified

# Virtual Memory

(OS topic; 가법계)

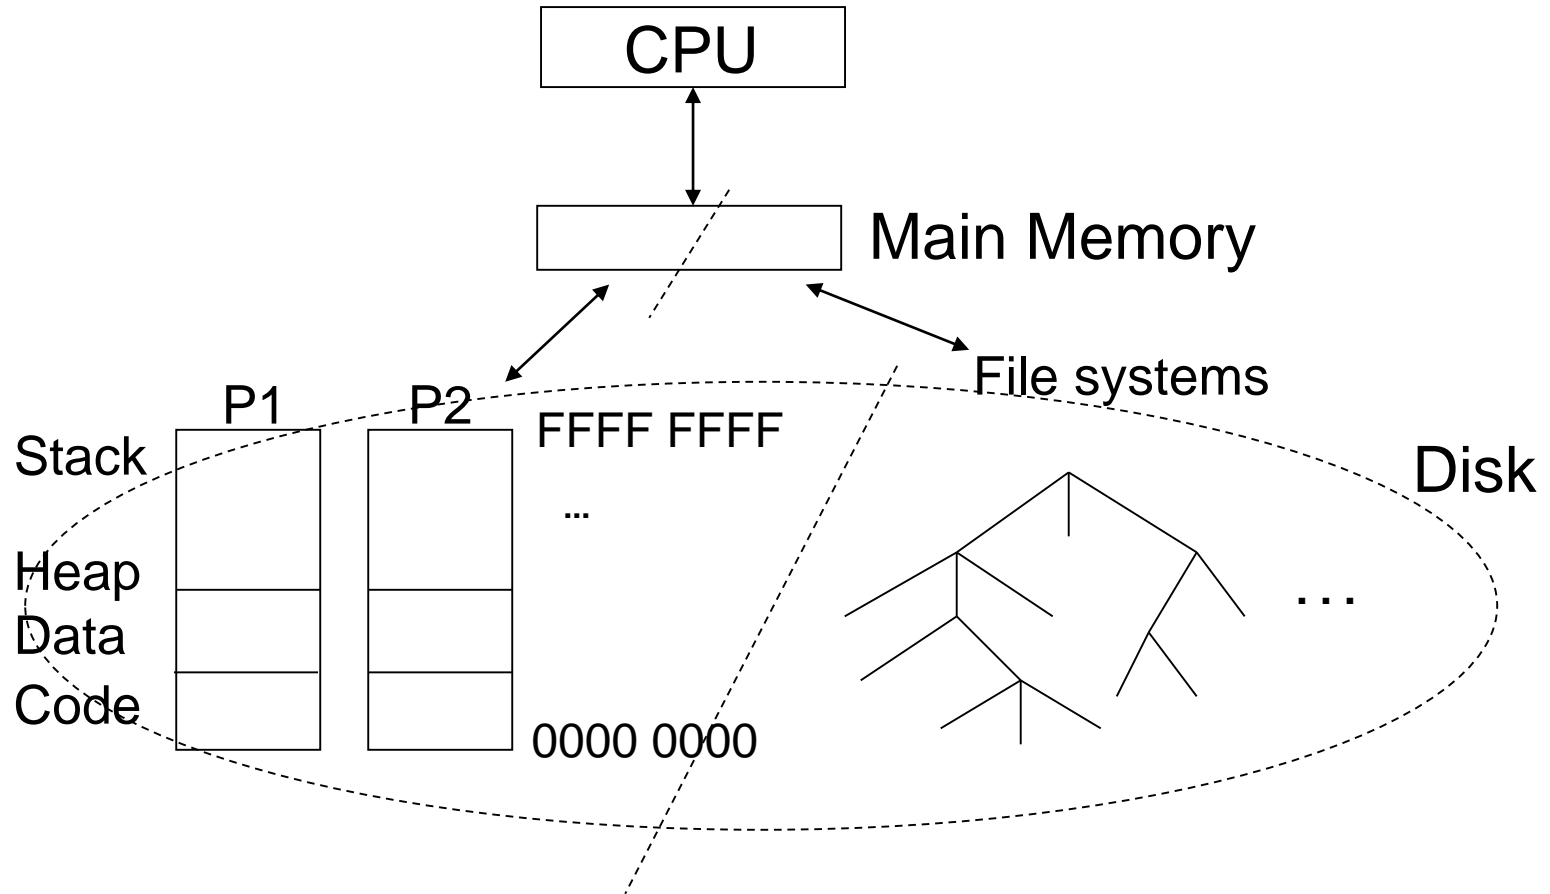
# Memory Management (반복)

- ❑ Cache memory management (Architecture topic)
  - Cache part of main memory
  - Implemented by hardware: fast, simple
    - Part of processor (on-chip cache)
    - Hardware accelerator: SW not know about it
- ❑ Virtual memory management (OS topic)
  - Use main memory as cache for disk
  - Implemented by software
    - Disk access is already slow (10 ms)
- ❑ Same principles (caching, locality, management) for both
  - Usage: independent of each other

# Memory Management

- ❑ Forget about cache memory for now
  - It is hardware accelerator, a part of CPU
  - With cache, CPU feels that main memory is faster
- ❑ Virtual memory management
  - Use main memory as a "cache" for disk
  - Focus only on main memory and disk
- ❑ Separation of concern

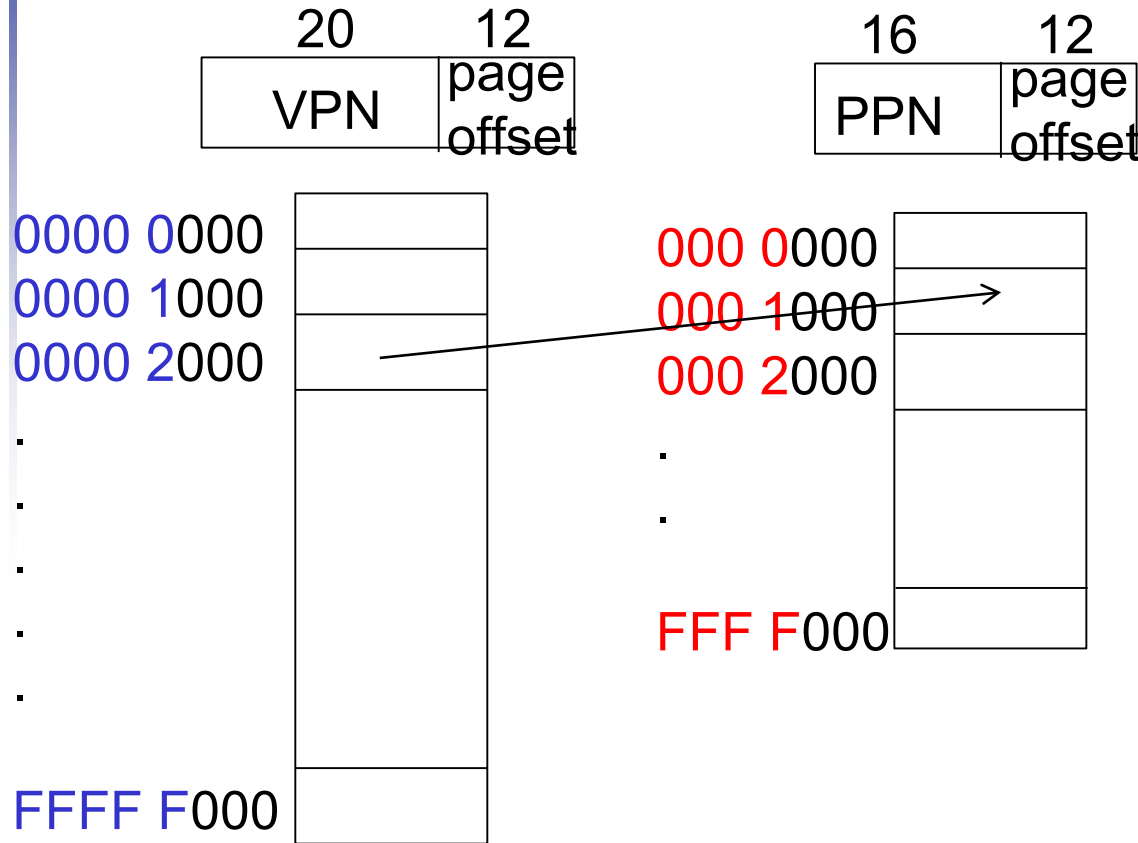
# Ignore Cache Memory



# Virtual Memory: Placement

VPN: virtual page no.

PPN: physical page no.



V	PPN	Disk location
0000 0		
0000 1		
0000 2	1 0001	
.		
.		
.		
.		
.		
FFFF F		

Page size:  $2^{12} = 4\text{KB}$

Virtual space:  $2^{32} = 4\text{GB}$

Main memory:  $2^{28} = 256\text{MB}$

Page Table  
(per-process)

# Virtual Memory

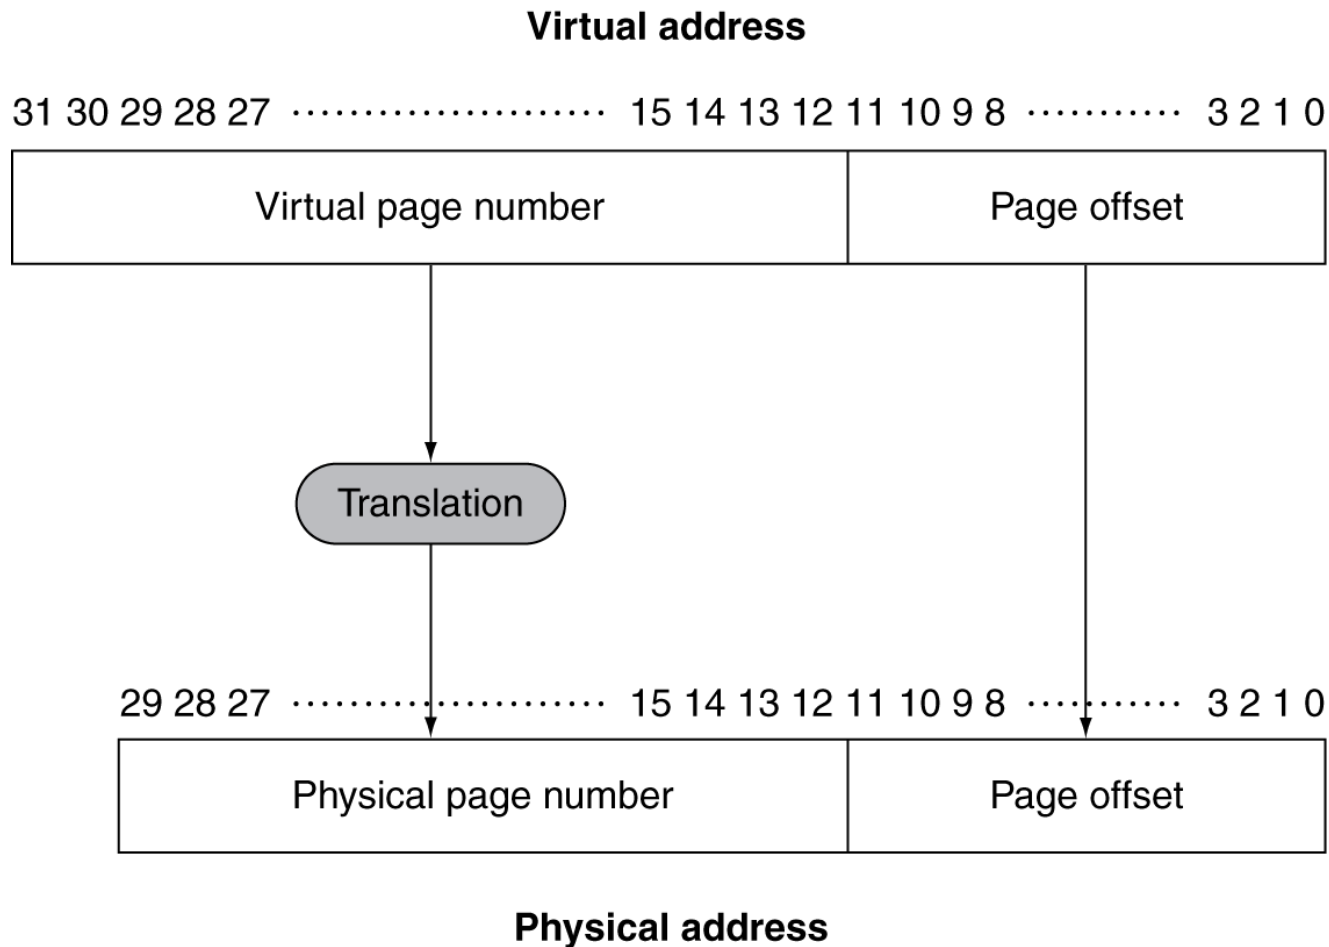
- ❑ What kind of mapping (placement) do we use?
  - Table data structure in software
- ❑ Use main memory as a “cache” for disk
- ❑ Multiple user programs share main memory
  - OS manages per-process page table
  - Protected from other programs
- ❑ CPU and OS translate virtual addresses to physical addresses



# Virtual Memory

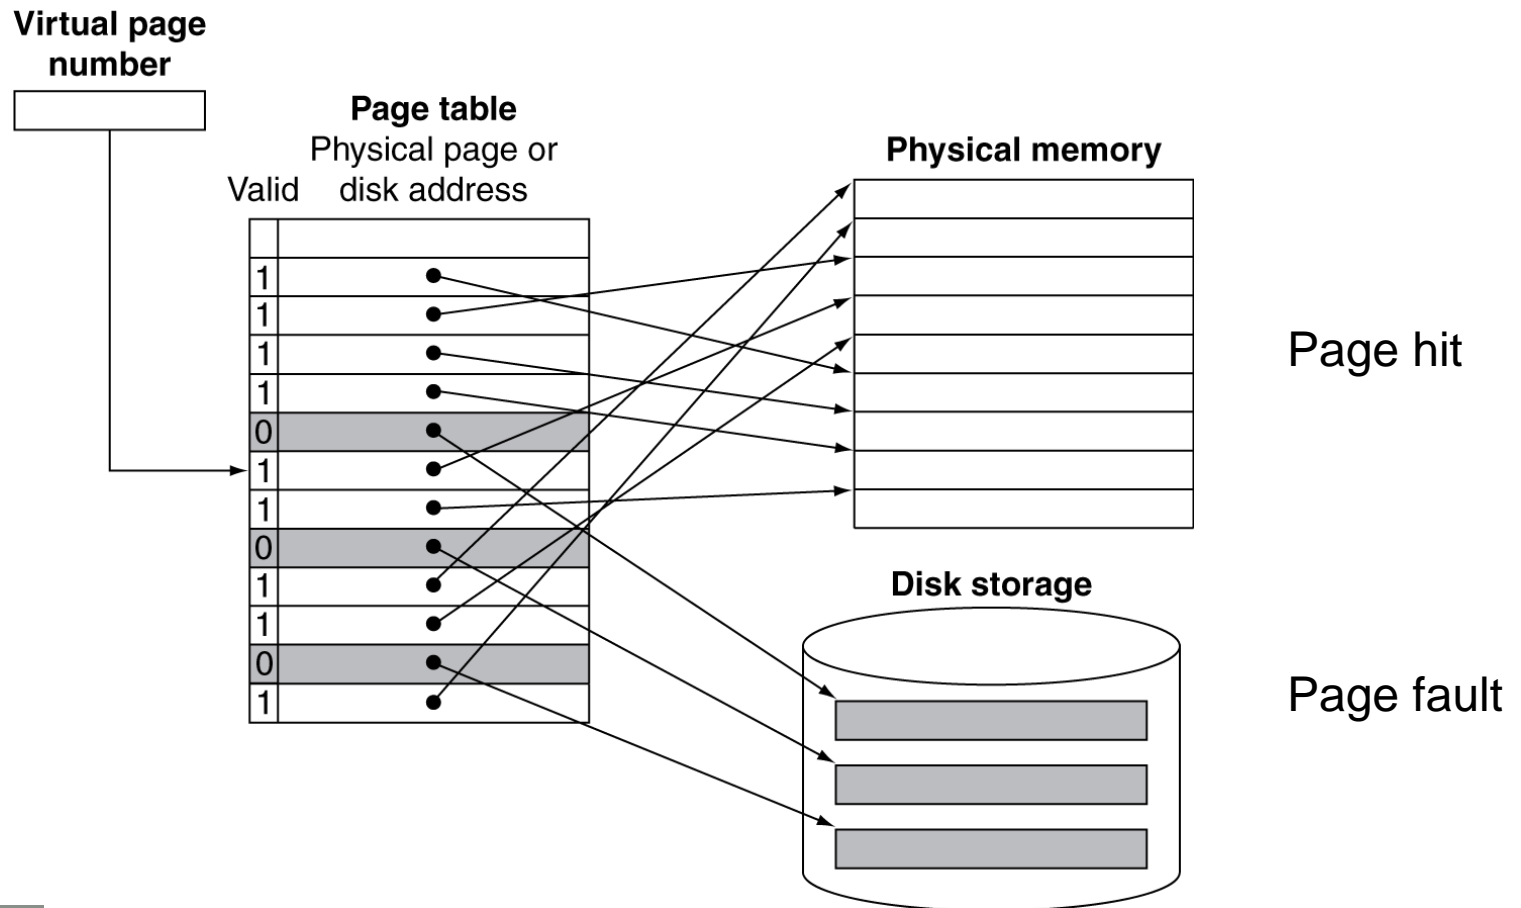
- ❑ When creating a process, OS:
  - Create space in disk or flash memory for all pages of process (swap space)
    - Create per-process page table
    - Record the location of each virtual page on disk
  - Also, track which process and which virtual page use each physical page
  - Page replacement (with approximated LRU)
- ❑ Process (state of a program)
  - Page table, program counter, registers

# Address Translation (V-to-P)



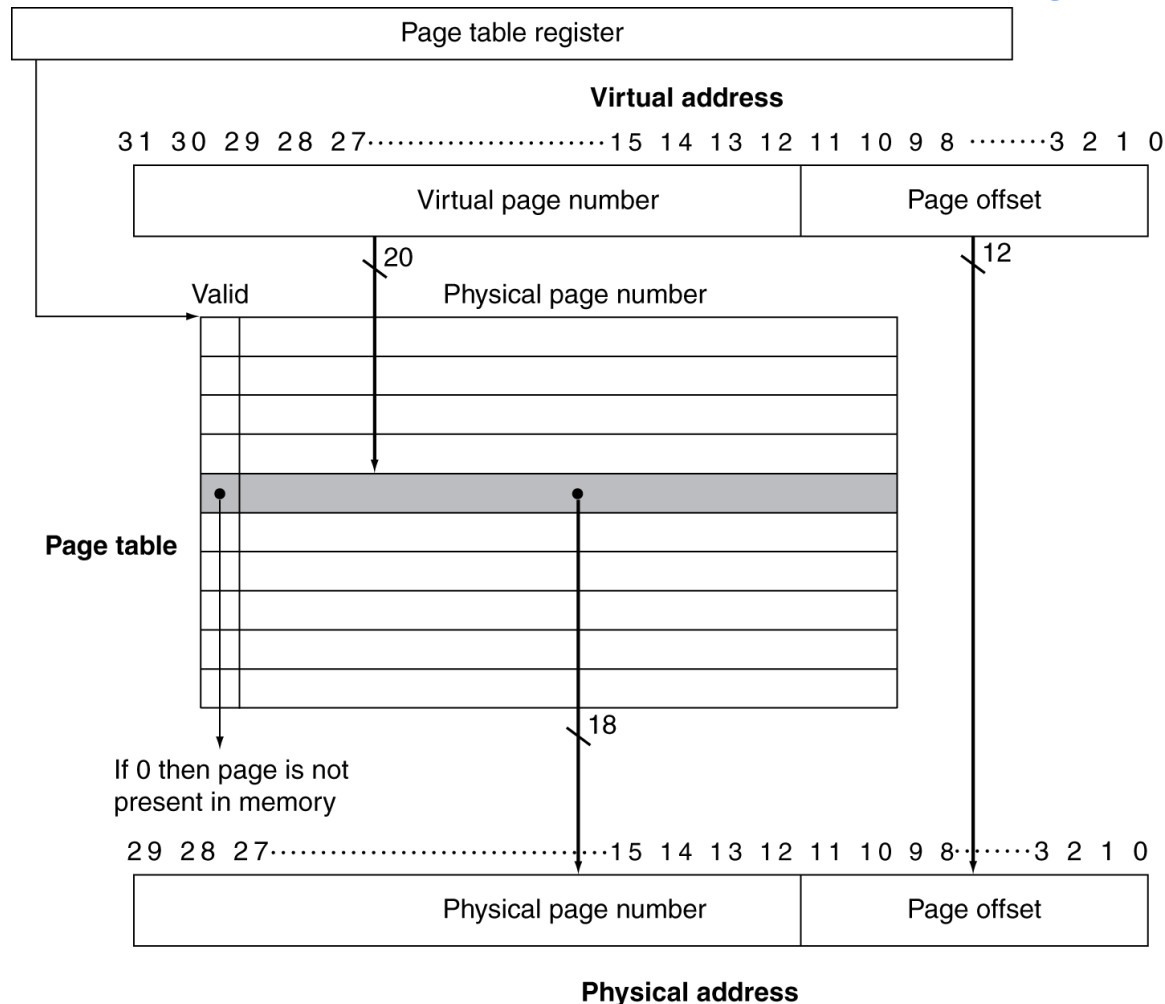
# Mapping Pages to Storage

## ❑ Per-process page tables



# Translation Using a Page Table

Process 별 page table 시작 주소



# Page Tables (부연)

- ❑ Stores placement information
  - Array of page table entries, indexed by virtual page number
  - Page table register in CPU points to page table in physical memory
- ❑ If page is present in memory
  - PTE (page table entry) stores the physical page no.
  - Plus other status bits (referenced, dirty, ...)
- ❑ If page is not present (page fault)
  - PTE can refer to location in swap space on disk

# Page Fault Penalty

- ❑ On page fault, the page must be fetched from disk
  - Takes millions of clock cycles
  - Handled by OS code
    - May switch process (in order of microseconds)
- ❑ Try to minimize page fault rate
  - Pages should be large enough
    - From 4KB up to 64KB
  - Fully associative placement
  - Smart replacement algorithms

# Four Issues in Virtual Memory

- ❑ Q1: placement
- ❑ Q2: identification
- ❑ Q3: write strategy
  - Use write-back (dirty bit in PTE)
- ❑ Q4: replacement policy
  - Some form of LRU (always approximated)

# Writes and Replacement

- ❑ Disk writes take millions of cycles
  - Block at once, not individual locations
  - Write through is impractical, so use write-back
    - Dirty bit in PTE set when page is written
- ❑ To reduce page fault rate, prefer least-recently used (LRU) replacement
  - Reference bit in PTE set to 1 on access to page
    - Periodically cleared to 0 by OS
  - Page with reference bit = 0 has not been used recently



# TLB

## (Translation-Lookaside Buffer)

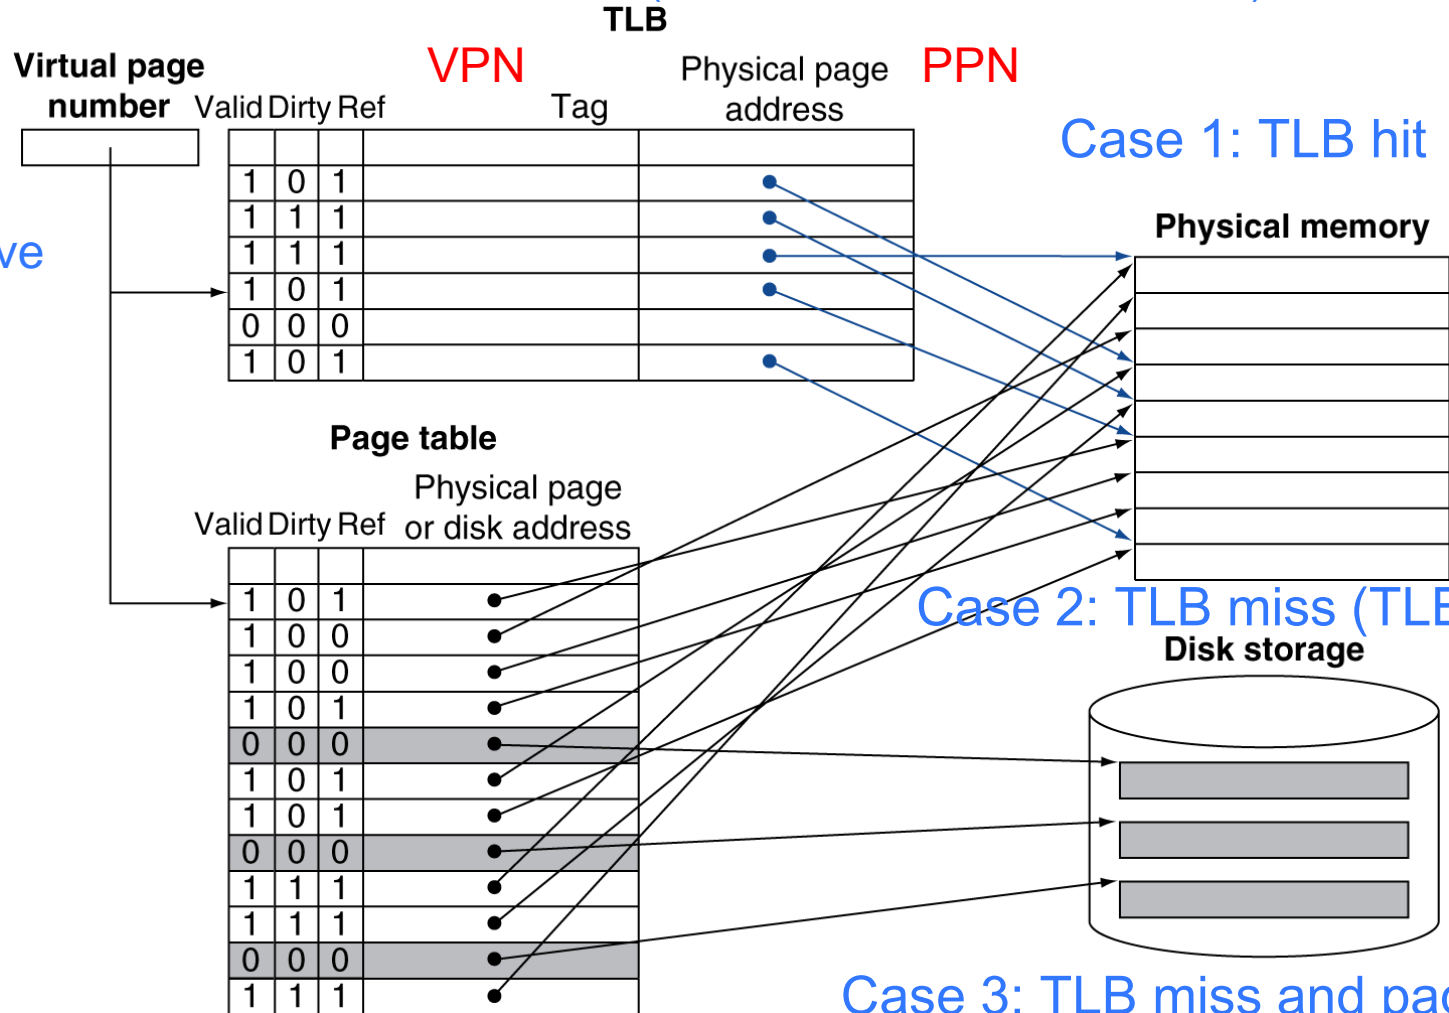
(OS Topic)

# Fast Translation Using a TLB

- ❑ Address translation requires extra memory references
  - One to access the PTE (page table entry)
  - Then the actual memory access
- ❑ Huge locality in accessing PTEs (VPN-PPN pairs)
  - Use the idea of cache again: store recently used PTEs (VPN-PPN pairs) in a fast cache
    - Called a Translation Look-aside Buffer (TLB)
    - Typical: 16-512 PTEs, 0.5-1 cycle for hit, 10-100 cycles for miss, 0.01%-1% miss rate

# Fast Translation Using a TLB

(Translation-lookaside buffer)



Fully  
associative  
mapping

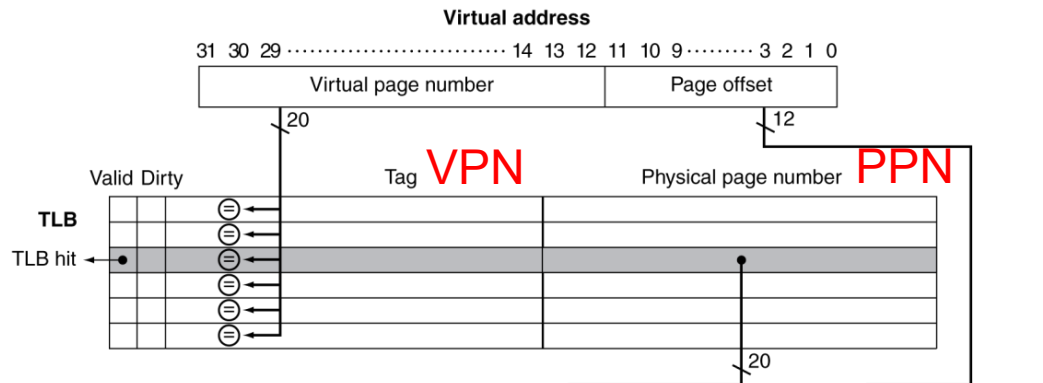
# TLB Misses (부연)

- ❑ If page is in memory
  - Load the PTE from memory and retry
  - Could be handled in hardware
  - Or in software
    - Raise a special exception, with optimized handler
- ❑ If page is not in memory (page fault)
  - OS handles fetching the page and updating the page table
  - Then restart the faulting instruction

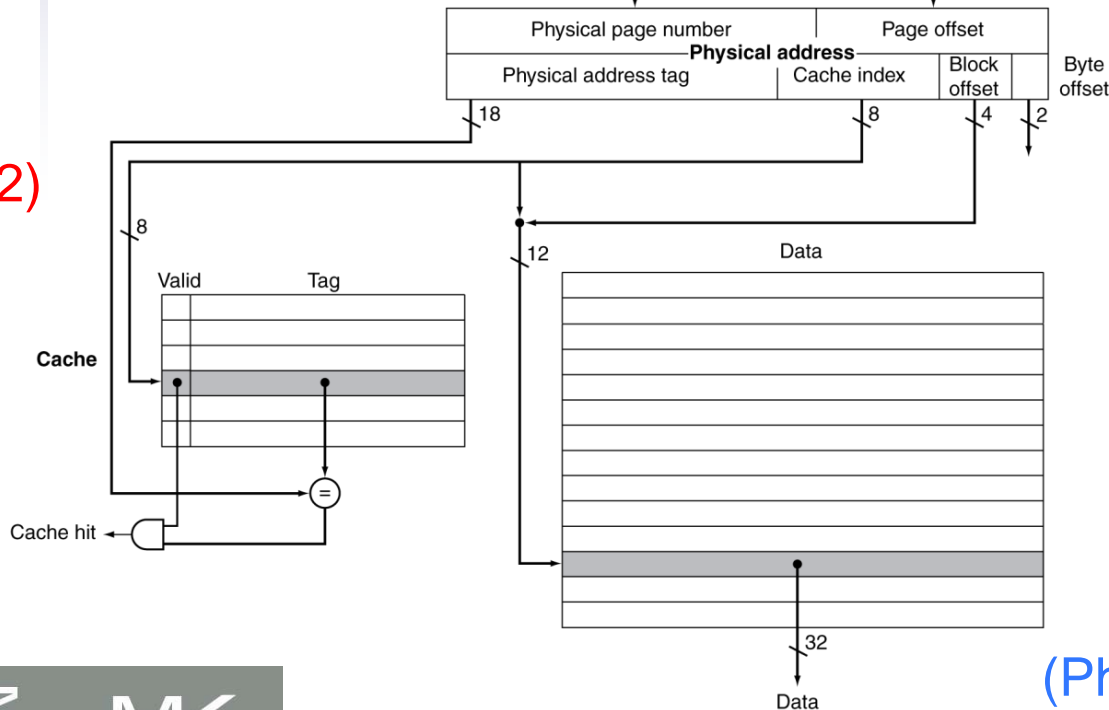
# Cache and Virtual Memory: Interactions

# TLB and Cache Interaction

(1)



(2)



❑ If cache tag uses physical address

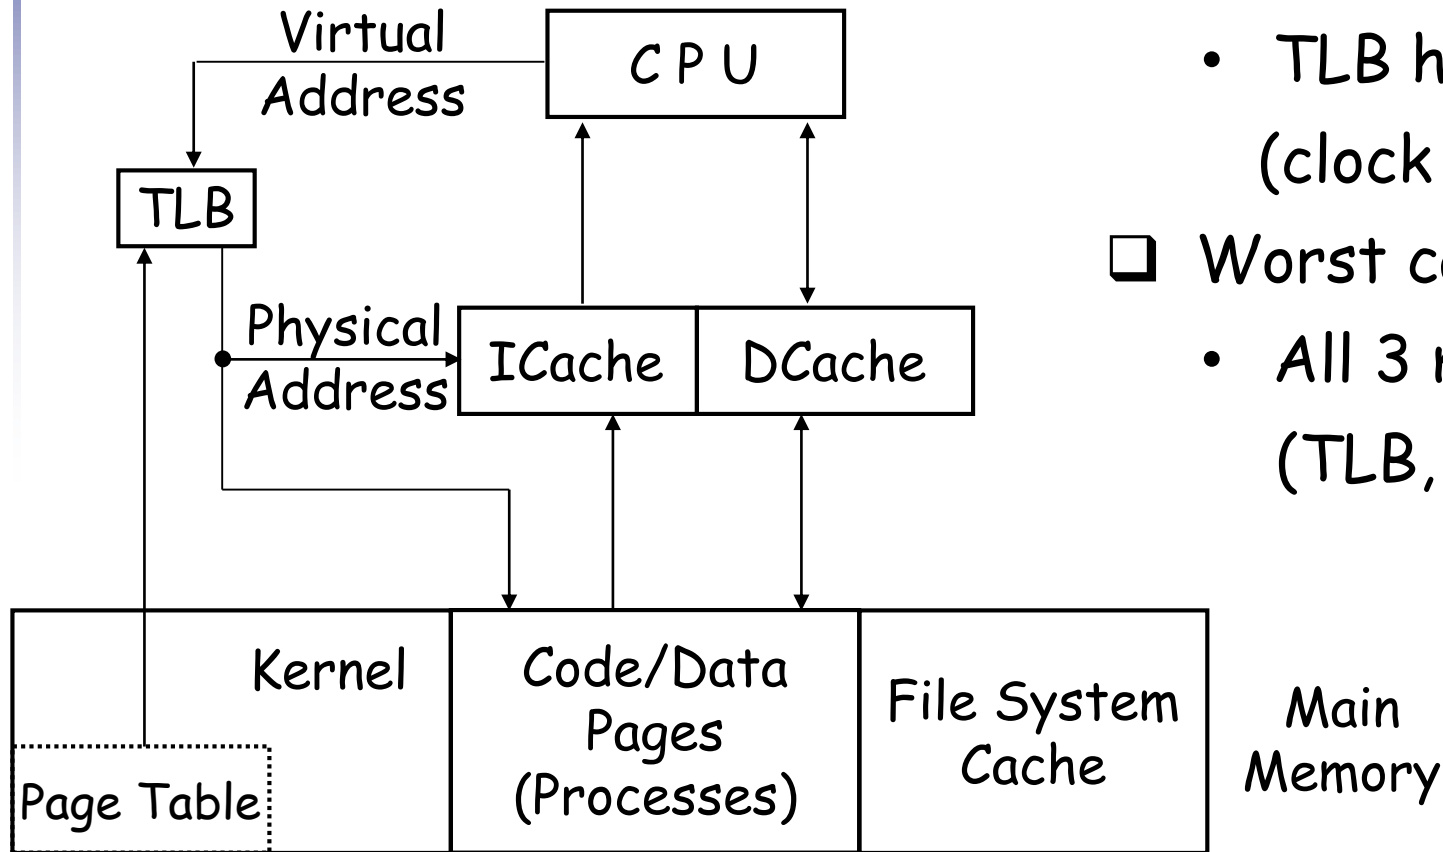
- Need to translate before cache lookup

❑ Alternative: use virtual address tag

- Complications due to aliasing
- Different virtual addresses for shared physical address

(Physical cache vs. virtual cache)

# Full Picture: Cache and Virtual Memory




- ❑ Best case
  - TLB hit, cache hit (clock cycle time)
- ❑ Worst case
  - All 3 misses (TLB, PT, cache)

# TLB, Cache, Main Memory (부연)

- ❑ TLB hit and cache hit: best case (clock cycle time)
- ❑ On TLB miss
  - Page table hit
    - Address translation using page table in main
    - Update TLB with new VPN-PPN pair
  - Page table miss (page fault)
    - OS handles fetching the page, updates the page table and TLB (process switch may occur)
    - Then restarts the faulting instruction
- ❑ On cache miss, check with main memory



# General-Purpose Computers



# Typical Values As of 2008 (참고)

Feature	Typical values for L1 caches	Typical values for L2 caches	Typical values for paged memory	Typical values for a TLB
Total size in blocks	250–2000	2500–25,000	16,000–250,000	40–1024
Total size in kilobytes	16–64	125–2000	1,000,000–1,000,000,000	0.25–16
Block size in bytes	16–64	64–128	4000–64,000	4–32
Miss penalty in clocks	10–25	100–1000	10,000,000–100,000,000	10–1000
Miss rates (global for L2)	2%–5%	0.1%–2%	0.00001%–0.0001%	0.01%–2%

Cache Performance:

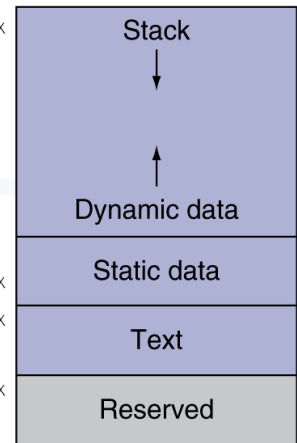
Three-C Model (optional)

# Sources of Misses

- ❑ Compulsory misses (aka cold start misses)
  - First access to a block
- ❑ Capacity misses
  - Due to finite cache size
  - A replaced block is later accessed again
- ❑ Conflict misses (aka collision misses)
  - In a non-fully associative cache
  - Due to competition for entries in a set
  - Would not occur in a fully associative cache of the same total size

# How to Compute Them

\$sp → 7fff fffc<sub>hex</sub>



\$gp → 1000 8000<sub>hex</sub>  
1000 0000<sub>hex</sub>

pc → 0040 0000<sub>hex</sub>  
0

## ❏ Compulsory misses

- Accessed regions (address trace) ÷ block size

## ❏ Cache simulation with given cache size and fully associative mapping

- Will get compulsory and capacity misses

## ❏ Cache simulation with given cache size and given mapping

- Will get compulsory and capacity and conflict misses

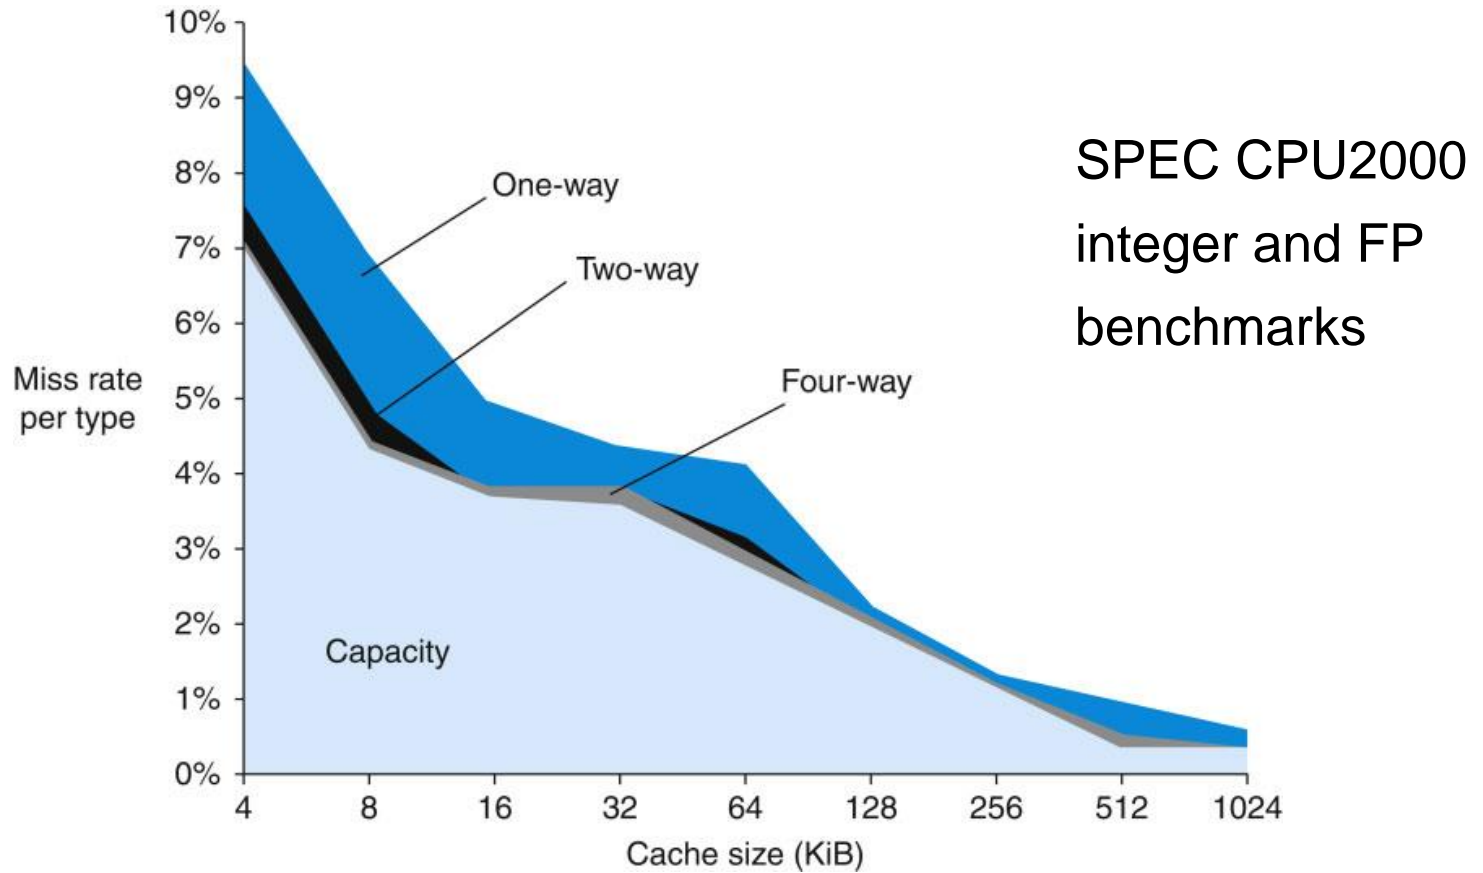
## ❖ Cache configuration: cache size, mapping, block size

# Cache Design Trade-offs

Design change	Effect on miss rate	Negative performance effect
Increase cache size	Decrease capacity misses	May increase access time
Increase associativity	Decrease conflict misses	May increase access time
Increase block size	Decrease compulsory misses	Increases miss penalty. For very large block size, may increase miss rate due to pollution.

# Three-C Model

- ❑ Compulsory misses: almost negligible (0.006%)



# Concluding Remarks

- ❑ Fast memories are small, large memories are slow
  - We really want fast, large memories ☹️
  - Caching gives this illusion 😊
- ❑ Principle of locality
  - Programs use a small part of their memory space frequently
- ❑ Memory hierarchy
  - L1 cache ↔ L2 cache ↔ ... ↔ DRAM memory ↔ disk
- ❑ Memory system design is critical for high performance



# Homework #14 (see Class Homepage)

1) Write a report summarizing the materials discussed in Topics 5-1 and 5-2 (이번 주 수업 내용)

\*\* 문장으로 써도 좋고 파워포인트 형태의 개조식 정리도 좋음

2) Solve Chapter 5 exercises 5.2.3, 5.3, 5.6, 5.7.3

□ Due: see Blackboard

- Submit electronically to Blackboard

# Big Picture

- ❑ Issue 1: Fundamental concepts and principles
  - What is computer, CSE, computer architecture?
- ❑ Issue 2: ISA (HW-SW interface) design
  - Ch. 1: computer performance
  - Ch. 2: language of computer; ISA
  - Ch. 3: data representation and ALU
- ❑ Issue 3: implementation of ISA (internal design)
  - Ch. 4: processor (data path, control, pipelining)
  - Ch. 5: memory system (cache memory)
- ❑ Ch. 6: short introduction to parallel processors