
Unix System Overview

System Programming


2019 여름 계절학기

한양대학교 공과대학 컴퓨터소프트웨어학부
홍석준


Programming

- ☐ **Unix System Overview**
- ☐ **Unix Standardization and Implementations**
- ☐ **File I/O**
- ☐ **Files and Directories**
- ☐ **Standard I/O Library**
- ☐ **System Data Files and Information**
- ☐ **Process Environment**
- ☐ **Process Control**
- ☐ **Signals**
- ☐ **Threads**
- ☐ **Advanced I/O**
- ☐ **Interprocess communication(IPC)**
- ☐ **Network IPC: Sockets**

History of UNIX



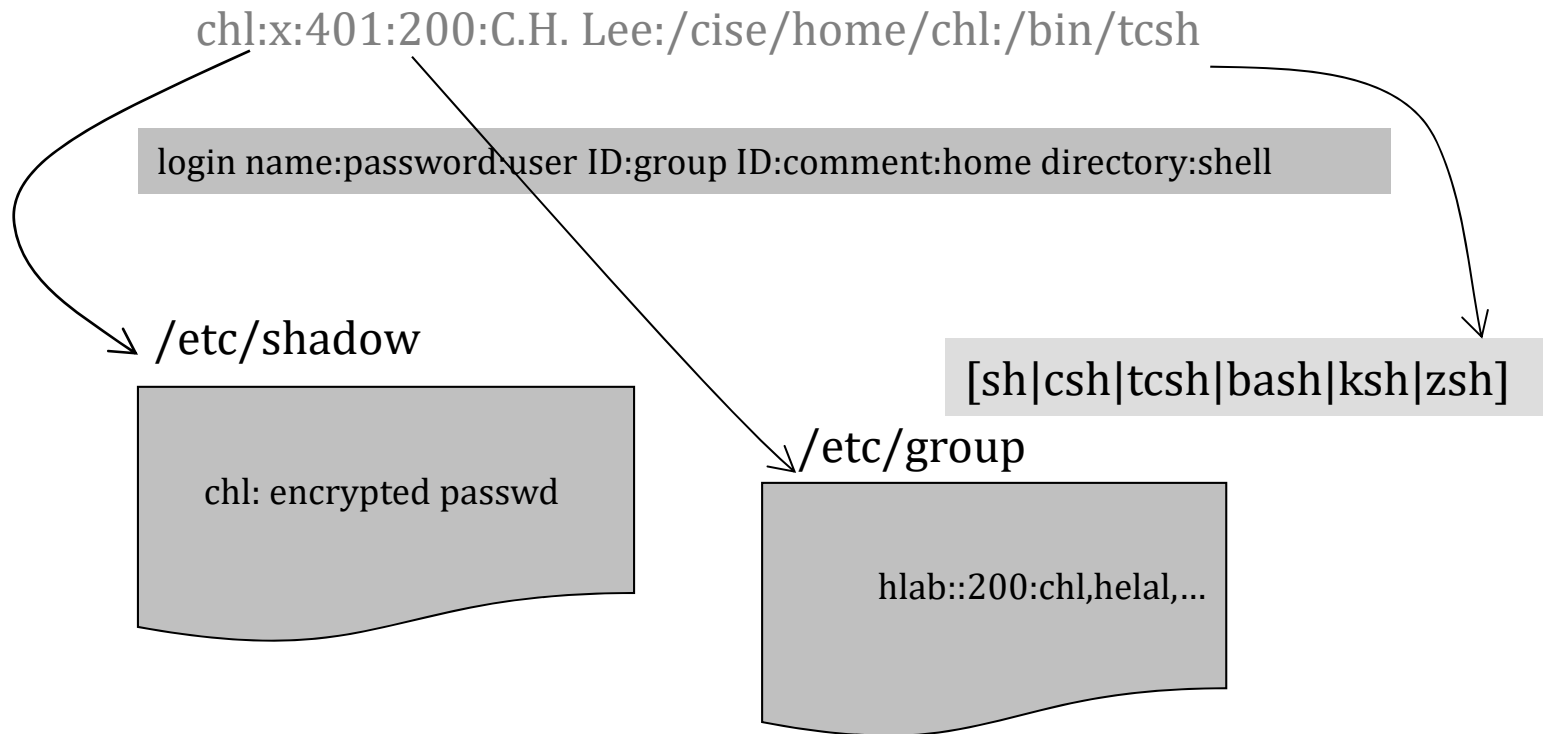
UNIX Architecture



□ Programming interface to the *kernel*

- System call interfaces and functions in the standard C library

❑ `/etc/passwd` composed of seven colon-separated fields:



□ Shell

- a command-line interpreter that reads user input and executes commands.
- Bourne Shell : UNIX version 7~
- Bourne-again shell : The GNU shell provided with all Linux systems

Name	Path	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10
Bourne shell	/bin/sh	•	•	copy of bash	•
Bourne-again shell	/bin/bash	optional	•	•	•
C shell	/bin/csh	link to tcsh	optional	link to tcsh	•
Korn shell	/bin/ksh	optional	optional	•	•
TENEX C shell	/bin/tcsh	•	optional	•	•

Figure 1.2 Common shells used on UNIX systems

❑ UNIX file system: a hierarchical arrangement of directories and files

- File
- A directory is a **file** that contains directory entries.

❑ Filename

- At least 255 character filenames (Most commercial UNIX)
- Special names: “/”, “.”, and “..”

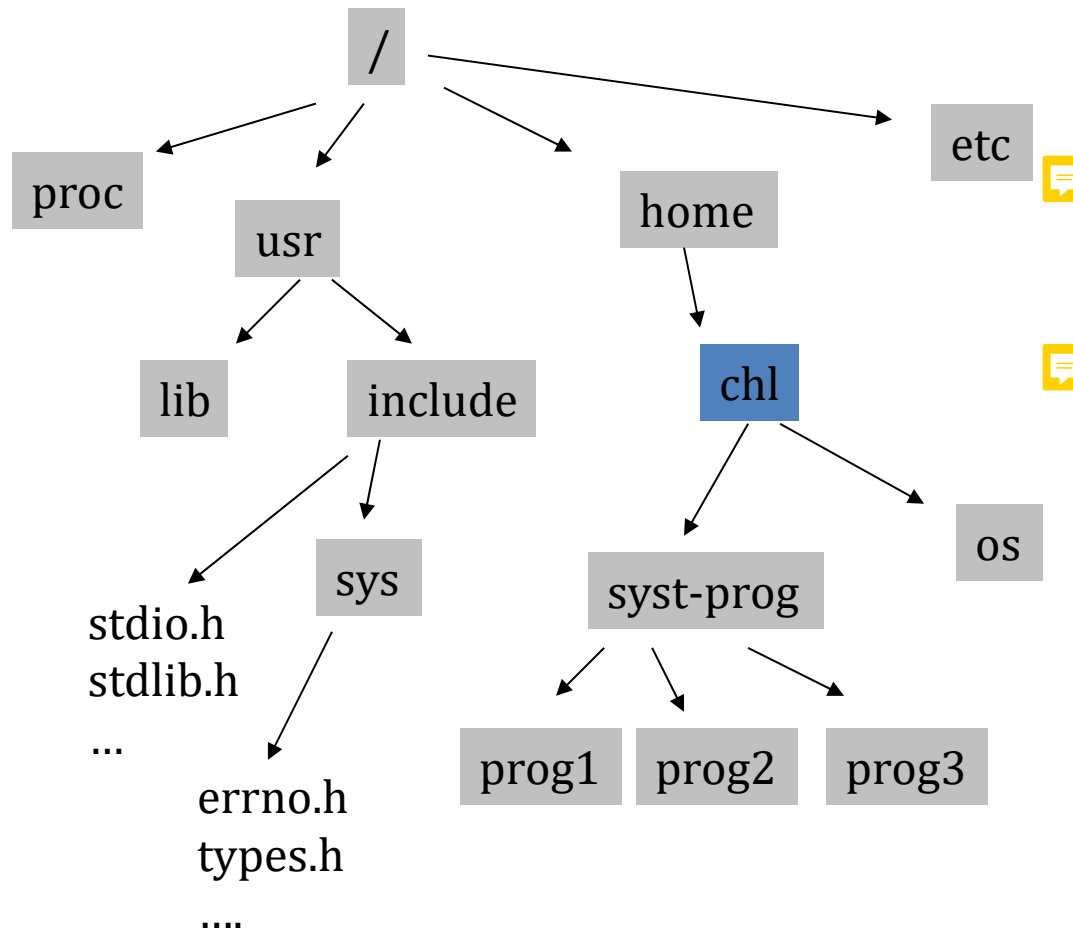
❑ Pathname

- A sequence of zero or more file names, separated by slash, and optionally starting with a slash
- Absolute pathname vs. relative pathname
 - A pathname that begins with a slash is called an absolute pathname
 - otherwise, it's called a relative pathname. Relative pathnames refer to files relative to the current directory

❑ Current Working Directory

- “.”
- pwd

UNIX File System



\$ pwd
/home/chl

\$ ls syst-prog
prog1 prog2 prog3

\$ ls ../.././chl/syst-prog
prog1 prog2 prog3

\$ ls /home/chl/syst-prog
prog1 prog2 prog3

❑ **File descriptors: small non-negative integers to identify the files**

❑ **Standard input, output, and error**

- stdin, stdout, and stderr
- Redirections: “<”, “>”, and “2>” in the case of csh

❑ **Unbuffered I/O vs. Standard I/O**

- (open, read, write, ...) vs. (fopen, fgetc, printf, ...)
- [Program 1.4](#) & [Program 1.5](#)

Program 1.3

```
#include "apue.h"
#include <dirent.h>

int main(int argc, char *argv[])
{
    DIR      *dp;
    struct    dirent  *dirp;

    if(argc != 2)
        err_quit("usage: ls directory_name");

    if((dp = opendir(argv[1])) == NULL)
        err_sys("can't open %s", argv[1]);

    while((dirp = readdir(dp)) != NULL)
        printf("%s\n", dirp->d_name);

    closedir(dp);
    exit(0);
}
```

Program 1.4

```
#include "apue.h"
#define BUFSIZE 4096
int main(void)
{
    int  n;
    char buf[BUFSIZE];
    while ((n = read(STDIN_FILENO, buf, BUFSIZE)) > 0)
        if (write(STDOUT_FILENO, buf, n) != n)
            err_sys("write error");

    if (n < 0)
        err_sys("read error");

    exit(0);
}
```

Program 1.5

```
#include "apue.h"
int main(void)
{
    int    c;

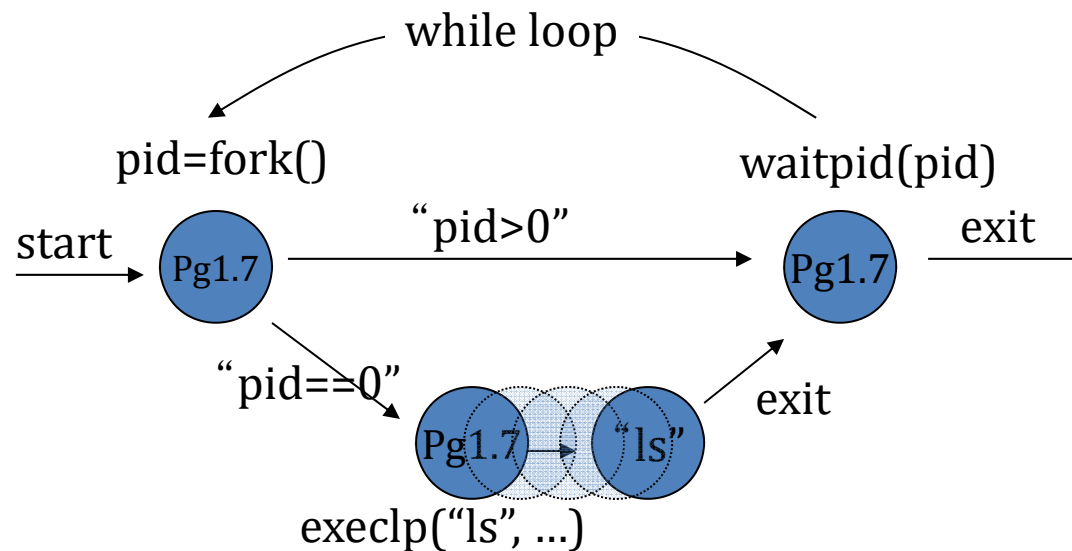
    while ((c = getc(stdin)) != EOF)
        if (putc(c, stdout) == EOF)
            err_sys("output error");

    if (ferror(stdin))
        err_sys("input error");

    exit(0);
}
```

I. Programs and Processes

- ❑ Program: an **executable file** residing on disk
- ❑ Process: an executing instance of a program
- ❑ Process ID: a nonnegative integer
- ❑ Process control: *fork*, *exec*, and *waitpid* in [Program 1.7](#)



Program 1.7

```
#include "apue.h"
#include <sys/wait.h>
int main(void)
{
    char  buf[MAXLINE]; /* from apue.h */
    pid_t pid;
    int   status;

    printf("%% ");
    while (fgets(buf, MAXLINE, stdin)
           != NULL) {
        if (buf[strlen(buf) - 1] == '\n')
            buf[strlen(buf) - 1] = 0;

        if ((pid = fork()) < 0) {
            err_sys("fork error");
        } else if (pid == 0) { /* child */
            execlp(buf, buf, (char *)0);
            err_ret("couldn't execute: %s", buf);
            exit(127);
        }

        /* parent */
        if ((pid = waitpid(pid, &status, 0)) < 0)
            err_sys("waitpid error");
        printf("%% ");
    }
    exit(0);
}
```

Threads

- ❑ **A thread of control**
- ❑ **All the thread within a process share**
 - Address space
 - File descriptors
 - Stacks
 - process-related attributes
- ❑ **Accesses to shared data should be synchronized**

Error Handling

- ❑ When an error occurs in one of the UNIX System functions, a negative value (or a null pointer) is returned is often returned.
- ❑ In case of an error, the `errno` is set to a value to give additional info.
- ❑ `/usr/include/(sys/)errno.h`

```
extern int errno;
```

```
#define EAGAIN  11 /* Resource temporarily unavailable */  
#define ENOMEM  12 /* Not enough core */  
#define EACCES  13 /* Permission denied */
```

– Never cleared if no error occurs

- ❑ `/usr/include/string.h`

```
char *strerror(int errnum);  
ex) strerror(EACCES)
```

- ❑ `/usr/include/stdio.h`

```
void perror(const char *msg);  
ex) perror(argv[0])
```

Signals

- ❑ A technique to notify a process that some condition has occurred.
- ❑ On receipt of a signal, a process can
 - ignore the signal (not recommended.)
 - let the default action occur.
 - call your own handler (catch the signal.)
- ❑ To generate a signal, [Ctrl-C], [Ctrl-backslash], or kill().
- ❑ [Program 1.10](#)

Program 1.10

```
#include "apue.h"
#include <sys/wait.h>
static void sig_int(int); /* our signal-catching function */
int main(void)
{
    char buf[MAXLINE]; /* from apue.h */
    pid_t pid;
    int status;
    if (signal(SIGINT, sig_int) == SIG_ERR)
        err_sys("signal error");
    printf("%% "); /* print prompt (printf requires %% to print %) */
    while (fgets(buf, MAXLINE, stdin) != NULL) {
        if (buf[strlen(buf) - 1] == '\n')
            buf[strlen(buf) - 1] = 0; /* replace newline with null */
    }
}
```

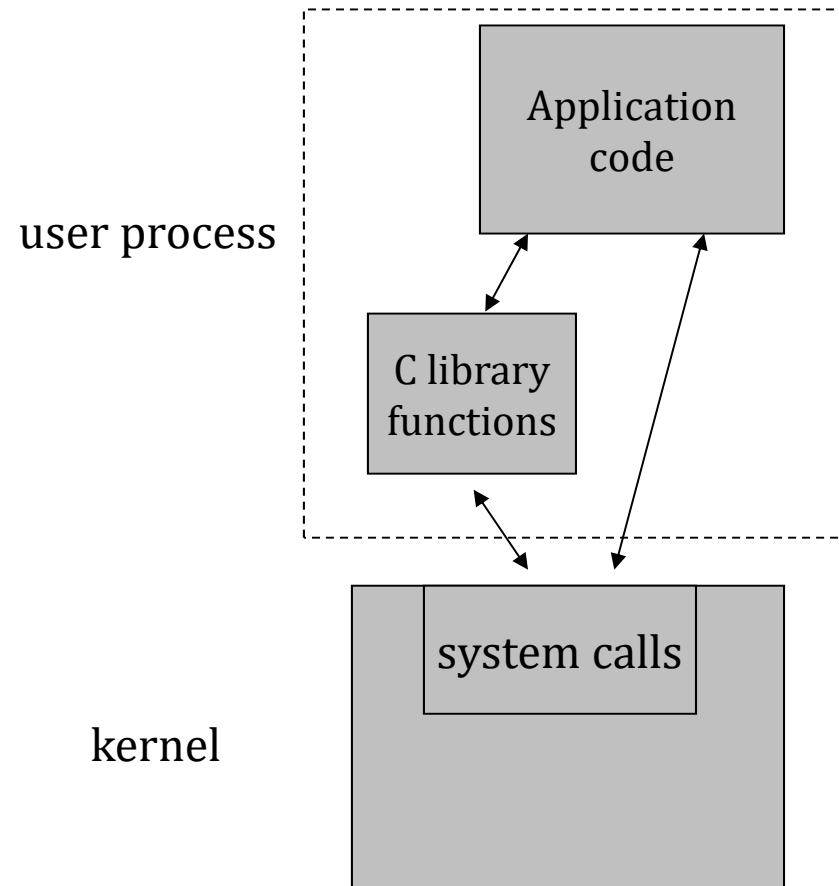
Program 1.10 (cont'd)

```
if ((pid = fork()) < 0) {
    err_sys("fork error");
} else if (pid == 0) { /* child */
    execlp(buf, buf, (char *)0);
    err_ret("couldn't execute: %s", buf);
    exit(127);
}
/* parent */
if ((pid = waitpid(pid, &status, 0)) < 0)
    err_sys("waitpid error");
printf("%%\n");
}
exit(0);
}
void sig_int(int signo)
{
    printf("interrupt\n%%\n");
}
```

UNIX Time Values

- ❑ Calendar time vs. Process time
- ❑ Calendar time in the number of seconds since 00:00:00 Jan. 1, 1970,
UTC (`time_t`)
- ❑ Process time in clock ticks (`clock_t`)
- ❑ Clock time (*wall clock time*), user CPU time, and system CPU time
 - `$ time ls`

Sys Calls and Library Functions



- 50 system calls for Unix Version 7, 110 for 4.4BSD, 120 for SVR4, 240-260 for Linux, and 320 for FreeBSD.
- System call – manual section 2
- Library – manual section 3C, 3m, ...
- More elaborate functionality
 - Atop of `sbrk()`, `malloc()` enables better memory allocation management.
 - Atop of `time()`, `gmtime()` provides broken-down time.
 - Atop of `read()`, `getc()` supports buffered I/O.

Unix man pages

□ Sections

- 1 – commands, e.g. `ls(1)`
- 2 – system calls and error numbers, e.g. `read(2)`
- 3 – functions and libraries such as 3C, 3M, etc
- 4 – file formats, e.g. `passwd(4)`
- 5 – miscellany, e.g. `environ(5)`
- 6 – games and demos
- 7 – special files, e.g. `hme(7D)`
- 8
- 9 – device driver interface

Unix man pages

\$ man man

\$ man -s 1 time

\$ man -s 2 time

\$ man -s 1 intro

\$ man -s 3 intro

Thank you for your attention !!

Q and A