

---

# File and Directories

---

System Programming

2019 여름 계절학기

한양대학교 공과대학 컴퓨터소프트웨어학부  
홍석준

# I. Contents

---

## ☐ File attributes handling

- `stat`

## ☐ Unix file system structure and symbolic links

## ☐ Directory operations

## I.stat, fstat, and lstat

```
#include <sys/stat.h>
```

```
int stat(const char *pathname, struct stat *buf);
```

```
int fstat(int fildes, struct stat *buf);
```

```
int lstat(const char *pathname, struct stat *buf);
```

- ❑ **stat/fstat** returns a structure of information about the named file.
- ❑ **lstat** returns information about the symbolic link, not the file referenced by the symbolic link.

## I.stat, fstat, and lstat

```
struct stat {
    mode_t    st_mode; /* file type & mode (permission) */
    ino_t     st_ino;  /* i-node number (serial number) */
    dev_t     st_dev;  /* device number (file system) */
    dev_t     st_rdev; /* device number for special files */
    nlink_t    st_nlink; /* number of links */
    uid_t     st_uid;  /* user ID of owner */
    gid_t     st_gid;  /* group ID of owner */
    off_t     st_size; /* size in bytes, for regular files */
    time_t    st_atime; /* time of last access */
    time_t    st_mtime; /* time of last modification */
    time_t    st_ctime; /* time of last file status change */
    long      st_blksize; /* best I/O block size */
    long      st_blocks; /* no. of disk blocks allocated */
};
```

## I. File Types

### ❑ Encoded in the `st_mode` member of the `stat` structure

- Regular file
- Directory file – pairs of (file name, pointer to information on the file)
- Character special file, e.g. tty
- Block special file, e.g. disk devices
- FIFO – named pipe
- Socket – used for network communication between processes
- Symbolic link – A type of file pointing to another file

### ❑ [Figure 4.3](#)

# Figure 4.3

```
#include "apue.h"
int main(int argc, char *argv[])
{ int          i;
  struct stat   buf;
  char          *ptr;

  for (i = 1; i < argc; i++) {
    printf("%s: ", argv[i]);
    if (stat(argv[i], &buf) < 0) {
      ret("lstat error");
      continue;
    }
    if (S_ISREG(buf.st_mode))
      ptr = "regular";
    else if (S_ISDIR(buf.st_mode))
      ptr = "directory";
    else if (S_ISCHR(buf.st_mode))
      ptr = "character special";
    else if (S_ISBLK(buf.st_mode))
      ptr = "block special";
    else if (S_ISFIFO(buf.st_mode))
      ptr = "fifo";
    else if (S_ISLNK(buf.st_mode))
      ptr = "symbolic link";
    else if (S_ISSOCK(buf.st_mode))
      ptr = "socket";
    else
      ptr = "*** unknown mode ***";
    printf("%s\n", ptr);
  }
  exit(0);
}
```

# I.Set-User-ID and Set-Group-ID

## ❑ User IDs and group IDs associated with each process

- real user ID, real group ID (`st_uid` & `st_gid`)
- effective user ID, effective group ID, supplementary group IDs
- saved set-user-ID, saved set-group-ID

## ❑ *set-user-ID* bit and *set-group-ID* bit in `st_mode`

- ☞ – *When this file is executed, set the effective user/group ID of the process to be the owner/group of the file.*
- As an example, `passwd(1)` is a set-user-ID program.

# I. File Access Permissions

## ❑ The nine file access permission bits from `<sys/stat.h>`

st_mode mask	Meaning
S_IRUSR	User-read
S_IWUSR	User-write
S_IXUSR	User-execute
S_IRGRP	Group-read
S_IWGRP	Group-write
S_IXGRP	Group-execute
S_IROTH	Other-read
S_IWOTH	Other-write
S_IXOTH	Other-execute

❑ `$ ls -l foo bar`

`-rwxr-xr-x` 1 stevens 0 Nov 16 16:23 bar

`-rw-r--r--` 1 stevens 0 Nov 16 16:23 foo

❑ `chmod(1), e.g. chmod g+w bar`



# I. File Access Permissions



## □ Directory

- R – to obtain a list of all the file names in the dir.
- W – to create/delete a file in the directory, both X and W are necessary.
- X - to pass through the directory comprising a pathname (e.g., execute permission in /, /usr, and /usr/include to open /usr/include/stdio.h), also called “**search bit**”.

## □ File

- R – O\_RDONLY and O\_RDWR for the open function
- W – O\_WRONLY, O\_RDWR, O\_TRUNC
- X – exec functions

□ **Q: To create a new file in a directory, which permission(s) do we need to have?**

□ **Q: To delete an existing file in a dir, which permission(s)?**

## I. File Access Permissions

- ☐ user ID and group ID of the file
  - ☐ effective user ID, effective group ID, and supplementary group IDs of the process
1. ☐ if the effective user ID is 0(super user), ...
  2. ☐ if the effective user ID equals the user ID, ...
  3. if the effective group ID or one of supplementary group IDs equals the group ID, ...
  4. if the appropriate other access permission bit is set, ...

☐ Q: what if permissions of 'a.out' are as follows?:

— -rw-rwx-r-x a.out

# I. Ownership of New Files and Directories

- ❑ The user ID of a new file is set to the effective user ID of the process.
- ❑ POSIX.1 options for the group ID of a new file
  - The effective group ID of the process
  - The group ID of the containing directory
    - FreeBSD, Mac OS X 10.3
    - set-group-ID bit (Linux, Solaris)

가해리 X

## I.access Function

```
#include <unistd.h>
```

```
int access(const char *pathname, int mode);
```

❑ mode: R\_OK, W\_OK, X\_OK, and F\_OK

❑ Accessibility test based on the real user ID/group ID to verify that the real user can access the given file.

– not the user ID of a process?

❑ [Figure 4.8](#)

# Figure 4.8



```
#include "apue.h"
#include <fcntl.h>
int main(int argc, char *argv[])
{
    if (argc != 2)
        err_quit("usage: a.out <pathname>");
    if (access(argv[1], R_OK) < 0)
        err_ret("access error for %s", argv[1]);
    else
        printf("read access OK\n");

    if (open(argv[1], O_RDONLY) < 0)
        err_ret("open error for %s", argv[1]);
    else
        printf("open for reading OK\n");

    exit(0);
}
```

# I.access Function

```
$ ls -l a.out
-rwxrwxr-w 1 sar          15945  Nov 30 12:10 a.out
$ ./a.out a.out
read access OK
open for reading OK
$ ls -l /etc/shadow
-r----- 1 root          1315   Jul 17 2002 /etc/shadow
$ ./a.out /etc/shadow
access error for /etc/shadow: Permission denied
open error for /etc/shadow: Permission denied
$ su
Passwd:
$ chown root a.out
$ chmod u+s a.out
$ ls -l a.out
-rwsrwxr-x 1 root          15945  Nov 30 12:10 a.out
$ exit
$ ./a.out /etc/shadow
access error for /etc/shadow: Permission denied
open for reading OK
```

*become superuser*

*enter superuser password*

*change file's user ID to root and turn on set-user-ID bit*

*check owner and SUID bit*

*go back to normal user*



```
#include <sys/stat.h>
```

```
mode_t umask(mode_t cmask);
```

- ❑ It sets the file mode creation mask for the process and returns the previous value.
- ❑ Any bits that are *on* in the file mode creation mask are turned *off* in the file's *mode*.
- ❑ [Figure 4.9](#)

```
$ umask
```

```
002
```

```
$ a.out
```

```
$ ls -l foo bar
```

```
-rw----- 1 sar 0 Nov 16 16:23 bar
```

```
-rw-rw-rw- 1 sar 0 Nov 16 16:23 foo
```

```
$ umask
```

```
002
```

*first print the current file mode creation mask*

*see if the file mode creation mask changed*

# Figure 4.9

```
#include "apue.h"
```

```
#include <fcntl.h>
```



```
#define RWRWRW (S_IRUSR|S_IWUSR|S_IRGRP|S_IWGRP|S_IROTH|S_IWOTH)
```

```
int main(void)
```

```
{
```

```
    umask(0);
```

```
    if (creat("foo", RWRWRW) < 0)
```

```
        err_sys("creat error for foo");
```

```
    umask(S_IRGRP | S_IWGRP | S_IROTH | S_IWOTH);
```

```
    if (creat("bar", RWRWRW) < 0)
```

```
        err_sys("creat error for bar");
```

```
    exit(0);
```

```
}
```

Handwritten red annotations: A large oval encircles the definition of RWRWRW. Below the definition, the permissions are broken down into bits: 1 1 0 for S\_IRGRP, 1 0 for S\_IWGRP, 1 1 0 for S\_IROTH, and 1 0 for S\_IWOTH. A red 'X' is drawn below the umask line, and a red arrow points from the oval to the umask line.



# I. chmod and fchmod Functions

```
#include <sys/stat.h>
int chmod(const char *pathname, mode_t mode);
int fchmod(int filedes, mode_t mode);
```

□ [Figure 4.12](#)


```
$ ls -l foo bar
-rw----- 1 sar ... bar
-rw-rw-rw- 1 sar ... foo
$ a.out
$ ls -l foo bar
???
```

unmask

mode	Description
S_ISUID S_ISGID S_ISVTX	set-user-ID on execution set-group-ID on execution saved-text (sticky bit)
S_IRWXU S_IRUSR S_IWUSR S_IXUSR	read, write, and execute by user (owner) read by user (owner) write by user (owner) execute by user (owner)
S_IRWXG S_IRGRP S_IWGRP S_IXGRP	read, write, and execute by group read by group write by group execute by group
S_IRWXO S_IROTH S_IWOTH S_IXOTH	read, write, and execute by other (world) read by other (world) write by other (world) execute by other (world)

# Figure 4.12

```
#include "apue.h"
int main(void)
{
    struct stat statbuf; /* turn on set-group-ID and turn off group-execute */

    if (stat("foo", &statbuf) < 0) 
        err_sys("stat error for foo");
    if (chmod("foo", (statbuf.st_mode & ~S_IXGRP) | S_ISGID) < 0)
        err_sys("chmod error for foo");

    if (chmod("bar", S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH) < 0)
        err_sys("chmod error for bar");

    exit(0);
}
```

# I.chown, fchown, and lchown

```
#include <unistd.h>
```

```
int chown(const char *pathname, uid_t owner, gid_t group);
```

```
int fchown(int fildes, uid_t owner, gid_t group);
```

```
int lchown(const char *pathname, uid_t owner, gid_t group);
```



- ❑ To change the user/group ID of a file.
- ❑ lchown changes the owners of the **symbolic link itself**.

# I.chown, fchown, and lchown

❑ If `_POSIX_CHOWN_RESTRICTED` is in effect,

- Only a superuser process can change the user ID of the file;
- A nonsuperuser process can change the group ID of the file if
  - the process owns the file (the *owner* equals to the user ID of the file), and
  - the *group* equals either the effective group ID of the process or one of the process's supplementary group IDs.
- You can't change the user ID of other users' files
- You can change the group ID of files that you own, but only to groups that you belong to.

# I. File Size

## ❑ **st\_size** of the **stat** structure: file size in bytes

- Regular file
- Directory
  - a multiple of a number such as 16 or 512
- Symbolic link
  - the actual number of bytes in the filename
  - For example,

```
lrwxrwxrwx 1 root    7 Sep 25 07:14 lib -> usr/lib
```

## ❑ **st\_blksize**: the preferred block size for I/O

## ❑ **st\_blocks**: the actual number of 512-byte blocks that are allocated

# I. File Size

## ❑ Holes in a File

```
$ ls -l core
```

```
-rw-r--r--  1 sar  8483248  Nov 18 12:18 core
```

```
$ du -s core
```

```
272 core
```

## ❑ **du reports the amount of disk space used by the file (272 512-byte blocks = 139,264 bytes)**

- BSD: 1024 byte blocks
- Solaris: 512 byte blocks

## I. File Truncation

---

```
#include <unistd.h>
```

```
int truncate(const char *pathname, off_t length);
```

```
int ftruncate(int filedes, off_t length);
```

□ They truncate an existing file to *length* bytes.

---

*Thank you for your attention !!*

---

Q and A