
File and Directories

System Programming

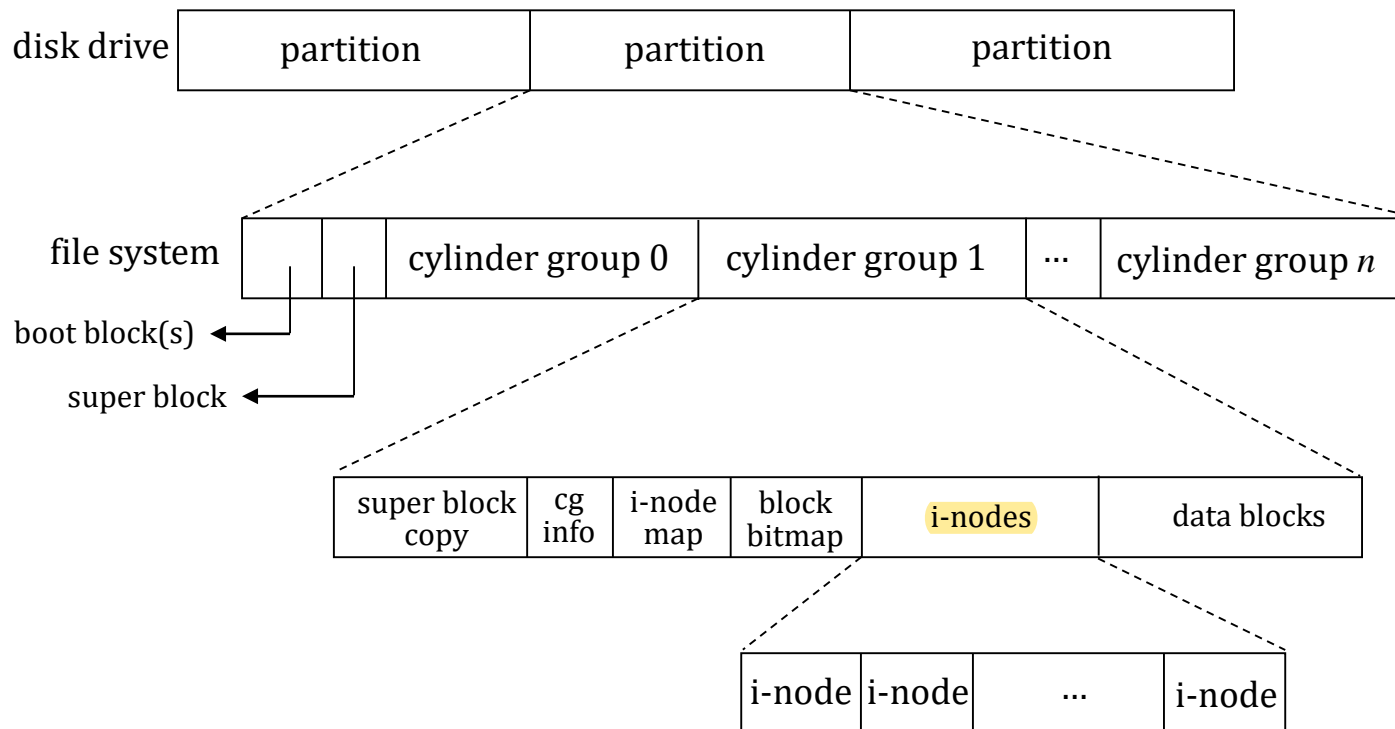
2019 여름 계절학기

한양대학교 공과대학 컴퓨터소프트웨어학부
홍석준

I. File Systems

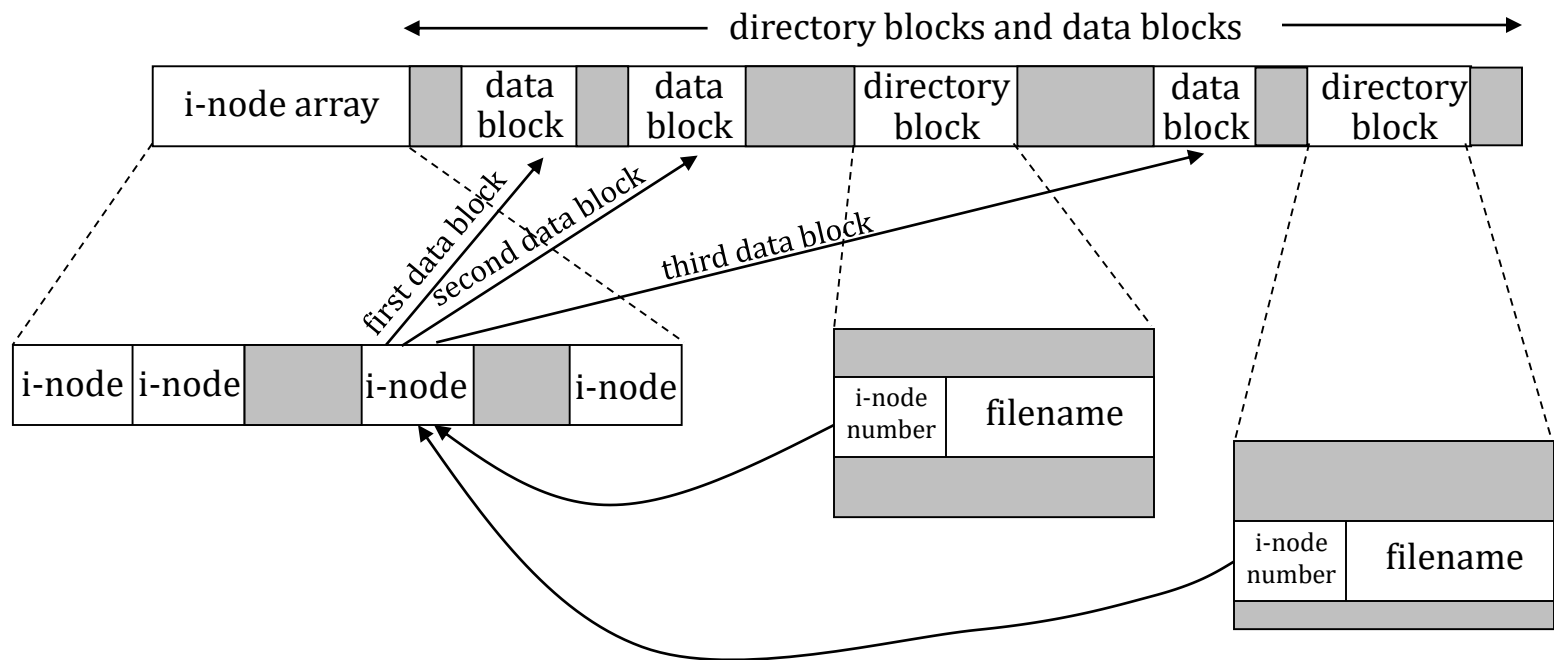
❑ Various implementations of the UNIX file system

❑ UFS




I. File Systems

- ❑ i-node contains info about the file, including file type, access permission, ref-count, size, ptrs to data blocks, and so on.
- ❑ Only two items (filename and i-node no.) are stored in the dir entry.



I. File Systems

- ❑ A link count in an i-node = the number of directory entries that point to the i-node
 - ❑ `st_nlink` in the `stat` structure
 - ❑ **Hard links vs. soft links**
 - Symbolic links (soft links)
 - The actual content of the file (the data blocks) contains the filename that the symbolic link points to.
- `lrwxrwxrwx 1 root 7 Sep 25 07:14 lib->usr/lib`
- ❑ No directory entry pointing to an i-node in a different file system.



I.link, unlink, remove, and rename Functions

```
#include <unistd.h>
```

```
int link (const char *existingpath, const char *newpath);
```

- ❑ Creates a new dir entry that references the existing path (, which increments the link count.)
- ❑ Both pathnames must be on the same file system (although POSIX.1 supports linking across file systems.)
- ❑ Only a superuser can create a link to a directory.

```
#include <unistd.h>
```

```
int unlink (const char *pathname) ;
```

- ❑ Removes the dir entry and decrements the link count (the file is deleted, when it reaches 0).
- ❑ If a symbolic link, unlink references the symbolic link itself.

I.link, unlink, remove, and rename Functions

```
#include <stdio.h>
```

```
int remove(const char *pathname) ;
```

❑ For a file, identical to unlink and, for a directory, to
rmdir

```
#include <stdio.h>
```

```
int rename(const char *oldname, const char  
*newname);
```

I.link, unlink, remove, and rename Functions

❑ [Program 4.16](#)

```
$ ls -l tempfile
-rw-r----- 1 sar      413265408   Jan 21   07:14  tempfile
$ df /home
Filesystem 1K-blocks    Used Available Use% Mounted on
/dev/hda4   11021440 1956332   9056108   18% /home
$ ./a.out &
1364
$ file unlinked
ls -l tempfile
ls: tempfile: No such file or directory
$ df /home
Filesystem 1K-blocks    Used Available Use% Mounted on
/dev/hda4   11021440 1956332   9056108   18% /home
$ done
df /home
Filesystem 1K-blocks    Used Available Use% Mounted on
/dev/hda4   11021440 1552352   9469088   15% /home
```

Program 4.16

```
#include "apue.h"
#include <fcntl.h>
int main(void)
{
    if (open("tempfile", O_RDWR) < 0)
        err_sys("open error");
    if (unlink("tempfile") < 0)
        err_sys("unlink error");

    printf("file unlinked\n");
    sleep(15);
    printf("done\n");

    exit(0);
}
```


I. Symbolic Links

❑ To get around the limitation of hard links

- Linking across file systems
- A hard link to a directory (only by superuser)

```
$ mkdir foo
```

```
$ touch foo/a
```


```
$ ln -s ../foo foo/testdir
```

```
$ ls -l foo
```

```
total 0
```

```
-rw-rw-r--  1 sar  0 Dec 6 06:06 a
```

```
lrwxrwxrwx  1 sar  6 Dec 6 06:06 testdir->../foo
```



I. symlink and readlink Functions

```
#include <unistd.h>
```

```
int symlink(const char *actualpath, const char *sympath);
```

- ❑ A new dir entry, *sympath*, is created that points to *actual path*.

```
#include <unistd.h>
```

```
ssize_t readlink(const char *pathname, char *buf, size_t  
    bufsize) ;
```

- ❑ `open` follows a symbolic link, while `readlink` opens the link itself and reads the name in the link.
- ❑ Equivalent to the actions of `open`, `read`, and `close`.

I. File Times

Field	Description	Example	ls(1) option
st_atime	Last access time of file data	read	-u
st_mtime	Last modification time of file data	write	default
st_ctime	Last change time of i-node status	chmod, chown	-c

- ❑ **The modification time is when the file contents were last modified.**
- ❑ **The changed-status time indicates when the i-node was last modified, e.g., changing the file access permission, the user ID, the number of links, etc.**
- ❑ **The three times for a file/directory and its parent directory**
 - For example, creating a new file affects the containing dir, and it affects the i-node for the new file. (Figure 4.20)

I.utime Function

```
#include <sys/types.h>
```

```
#include <utime.h>
```

```
int utime(const char *pathname, const struct utimbuf *times);
```

```
struct utimbuf {
```

```
    time_t actime;        /* access time */
```

```
    time_t modtime;       /* modification time */
```

```
}
```

- ❑ The `utime` changes the access/modification time of a file.
- ❑ If `times` is `NULL`, set to current time.
 - Effective UID must equal the real ID of the file, or write permission for the file.
- ❑ Otherwise, set to values pointed by `times`.
 - Effective UID must equal the real ID of the file, or superuser privilege
- ❑ [Program 4.21](#)

Program 4.21

```
#include "apue.h"
#include <fcntl.h>
#include <utime.h>
int main(int argc, char *argv[])
{
    int i, fd;
    struct stat statbuf;
    struct utimbuf timebuf;

    for (i = 1; i < argc; i++) {
        if (stat(argv[i], &statbuf) < 0) { /* fetch current times */
            err_ret("%s: stat error", argv[i]);
            continue;
        }
        if ((fd = open(argv[i], O_RDWR | O_TRUNC)) < 0) { /* truncate */
            err_ret("%s: open error", argv[i]);
            continue;
        }
        close(fd);
        timebuf.actime = statbuf.st_atime;
        timebuf.modtime = statbuf.st_mtime;
        if (utime(argv[i], &timebuf) < 0) { /* reset times */
            err_ret("%s: utime error", argv[i]);
            continue;
        }
    }
    exit(0);
}
```

I. `mkdir` and `rmdir` Functions

```
#include <sys/stat.h>
```

```
int mkdir(const char *pathname, mode_t mode) ;
```

- ❑ The `mode` is modified by the `umask` of the process.
- ❑ The user ID and group ID of the new directory.

```
#include <unistd.h>
```

```
int rmdir(const char *pathname);
```

- ❑ If the link count of the dir becomes 0, and no other process has the dir open, then the space occupied by the dir is freed.

I. Reading Directories

```
#include <dirent.h>
```

```
DIR *opendir(const char *pathname);
```

```
struct dirent *readdir(DIR *dp);
```

```
void rewinddir(DIR *dp);
```

```
int closedir(DIR *dp);
```

```
long telldir(DIR *dp);
```

```
void seekdir(DIR *dp, long loc);
```

```
struct dirent {
```

```
    ino_t    d_ino;                /* i-node number */
```

```
    char     d_name[NAME_MAX + 1]; /* null-terminated fname */  
}
```

- ☐ Only the kernel can write to a directory.
- ☐ Write and execute permission to create/delete files
- ☐ [Program 4.22](#)

Program 4.22

```
#include "apue.h"
#include <dirent.h>
#include <limits.h> /* function type that is called for each filename */
typedef      int      Myfunc(const char *, const struct stat *, int);
static Myfunc      myfunc;
static int      myftw(char *, Myfunc *);
static int      dopath(Myfunc *);
static long      nreg, ndir, nblk, nchr, nfifo, nslink, nsock, ntot;

int main(int argc, char *argv[])
{
    int ret;
    if (argc != 2)
        err_quit("usage: ftw <starting-pathname>");
    ret = myftw(argv[1], myfunc); /* does it all */
    ntot = nreg + ndir + nblk + nchr + nfifo + nslink + nsock;
    if (ntot == 0)
        ntot = 1; /* avoid divide by 0; print 0 for all counts */
    printf("regular files = %7ld, %5.2f %%%\n", nreg, nreg*100.0/ntot);
    printf("directories   = %7ld, %5.2f %%%\n", ndir, ndir*100.0/ntot);
    printf("block special = %7ld, %5.2f %%%\n", nblk, nblk*100.0/ntot);
    printf("char special  = %7ld, %5.2f %%%\n", nchr, nchr*100.0/ntot);
    printf("FIFOs        = %7ld, %5.2f %%%\n", nfifo, nfifo*100.0/ntot);
    printf("symbolic links = %7ld, %5.2f %%%\n", nslink, nslink*100.0/ntot);
    printf("sockets      = %7ld, %5.2f %%%\n", nsock, nsock*100.0/ntot);
    exit(ret);
}
```


Program 4.22

```
/*
 * Descend through the hierarchy, starting at "pathname".
 * The caller's func() is called for every file. */
#define    FTW_F    1                /* file other than directory */
#define    FTW_D    2                /* directory */
#define    FTW_DNR   3                /* directory that can't be read */
#define    FTW_NS    4                /* file that we can't stat */
static char *fullpath;                /* contains full pathname for every file */
static size_t pathlen;
static int                                /* we return whatever func() returns */
myftw(char *pathname, Myfunc *func)
{
    fullpath = path_alloc(& pathlen); /* malloc's for PATH_MAX+1 bytes */
                                   /* ({Prog pathalloc}) */
    if(pathlen <= strlen(pathname))
    {
        pathlen = strlen(pathname) * 2;
        if((fullpath = realloc(fullpath, pathlen)) == NULL )
            err_sys("realloc failed");
    }
    strcpy(fullpath, pathname);
    return(dopath(func));
}
```

Program 4.22

```
/*
 * Descend through the hierarchy, starting at "fullpath". If "fullpath" is anything other than a directory, we lstat() it,
 * call func(), and return. For a directory, we call ourself recursively for each name in the directory. */
static int      /* we return whatever func() returns */
dopath(Myfunc* func) {
    struct stat statbuf;
    struct dirent *dirp; DIR *dp;
    int ret, n;
    if (lstat(fullpath, &statbuf) < 0) /* stat error */
        return(func(fullpath, &statbuf, FTW_NS));
    if (S_ISDIR(statbuf.st_mode) == 0) /* not a directory */
        return(func(fullpath, &statbuf, FTW_F)); /* It's a directory. First call func() for the directory,
                                                * then process each filename in the directory. */
    if ((ret = func(fullpath, &statbuf, FTW_D)) != 0) return(ret);

    n = strlen(fullpath);
    if (n + NAME_MAX + 2 > pathlen){
        pathlen *= 2;
        if((fullpath = realloc(fullpath, pathlen)) == NULL)
            err_sys("realloc failed");
    }
    fullpath[n++] = '/';
    fullpath[n] = 0;

    if ((dp = opendir(fullpath)) == NULL) /* can't read directory */
        return(func(fullpath, &statbuf, FTW_DNR));
    while ((dirp = readdir(dp)) != NULL) {
        if (strcmp(dirp->d_name, ".") == 0 || strcmp(dirp->d_name, "..") == 0) continue; /* ignore dot and dot-dot */
        strcpy(ptr, dirp->d_name); /* append name after slash */
        if ((ret = dopath(func)) != 0) /* recursive */
            break; /* time to leave */
    }
    fullpath[n-1] = 0;

    if (closedir(dp) < 0)
        err_ret("can't close directory %s", fullpath);
    return(ret);
}
```

Program 4.22

```
static int
myfunc(const char *pathname, const struct stat *statptr, int type) {
    switch (type) {
        case FTW_F:
            switch (statptr->st_mode & S_IFMT) {
                case S_IFREG: nreg++; break;
                case S_IFBLK: nblk++; break;
                case S_IFCHR: nchr++; break;
                case S_IFIFO:      nfifo++;    break;
                case S_IFLNK:      nlink++;    break;
                case S_IFSOCK: nsock++; break;
                case S_IFDIR: err_dump("for S_IFDIR for %s", pathname);
                               /* directories should have type = FTW_D */
            }
            break;
        case FTW_D: ndir++; break;
        case FTW_DNR: err_ret("can't read directory %s", pathname); break;
        case FTW_NS: err_ret("stat error for %s", pathname); break;
        default: err_dump("unknown type %d for pathname %s", type, pathname);
    }
    return(0);
}
```

I.chdir, fchdir, and getcwd

```
#include <unistd.h>
```

```
int chdir(const char *pathname) ;
```

```
int fchdir(int fildes) ;
```

```
char *getcwd(char *buf, size_t size) ;
```

I. Special Device Files

- ❑ Every file system is known by its major/minor device numbers stored in a `dev_t` object.
- ❑ `major` and `minor` macros to access major/minor numbers.
- ❑ The `st_dev` is the dev no. of the file system containing the file.
- ❑ The `st_rdev` contains the dev no. of the character/block special files.
- ❑ [Program 4.25](#)

Program 4.25

```
#include "apue.h"
#ifdef SOLARIS
#include <sys/mkdev.h>
#endif
int main(int argc, char *argv[]) {
    int i;
    struct stat buf;
    for (i = 1; i < argc; i++) {
        printf("%s: ", argv[i]);
        if (stat(argv[i], &buf) < 0) {
            err_ret("stat error");
            continue;
        }
        printf("dev = %d/%d", major(buf.st_dev), minor(buf.st_dev));
        if (S_ISCHR(buf.st_mode) || S_ISBLK(buf.st_mode)) {
            printf(" (%s) rdev = %d/%d",
                (S_ISCHR(buf.st_mode)) ? "character" : "block",
                major(buf.st_rdev), minor(buf.st_rdev));
        }
        printf("\n");
    }
    exit(0);
}
```

Thank you for your attention !!

Q and A