
Standard I/O Library

System Programming

2019 여름 계절학기

한양대학교 공과대학 컴퓨터소프트웨어학부
홍석준

I. Standard I/O Library

- ❑ The standard I/O library by Dennis Ritchie around 1975 based on the Portable I/O library by Mike Lesk**
- ❑ The ISO C standard**
- ❑ Additional interfaces by the Single UNIX Specification**
- ❑ Easy-to-use library for buffer allocation, I/O in optimal-sized chunks, etc.**

I. Streams and `FILE` Objects

- ❑ When we open or create a file with the standard I/O library, a stream is associated with the file.
- ❑ The standard I/O function `fopen` returns a pointer to a **FILE object**, i.e. a *file pointer*, containing all the information to manage the stream.
 - The file descriptor, a pointer to a buffer, the size of the buffer, # of characters in the buffer, an error flag, and the like.
- ❑ **Three predefined streams:** `stdin`, `stdout`, and `stderr`

I. Buffering

- ❑ **To use the minimum number of `read` and `write` calls.**
- ❑ **Three types: fully buffered, line buffered, and unbuffered**

- ❑ **Fully buffered**
 - Actual I/O takes place, when the buffer is filled.
 - The buffer is obtained by the first I/O function on a stream.
 - The buffer can be flushed (i.e., writing out its contents.)
 - Automatically when the buffer fills, or
 - `fflush` flushes a stream.

I. Buffering

❑ Line buffered

- Actual I/O is performed, when a new line char is encountered on input or output.
- Typically used on a stream to refer to a terminal.
- Two caveats
 - Actual I/O might take place if the buffer is filled before a newline.
 - All line-buffered output streams are flushed, if input is requested from an unbuffered or a line-buffered stream.

❑ Unbuffered

- No buffering
- The `stderr` is normally unbuffered.

I. Buffering

❑ ISO C requirements

- `stdin` and `stdout` are fully buffered, iff it is not an interactive device.
- `stderr` is never fully buffered.

❑ In most implementations,

- `stderr` is always unbuffered.
- All other streams are line buffered if referring to a terminal device; otherwise, they are fully buffered.

❑ These defaults can be changed by either of `setbuf` and `setvbuf`.

I. Buffering

```
#include <stdio.h>
void setbuf(FILE *fp, char *buf) ;
int setvbuf(FILE *fp, char *buf, int mode, size_t
    size) ;
```

❑ To enable/disable buffering, we set setbuf's *buf* to non-null/NULL.

❑ setvbuf

- *mode*: `_IOFBF`, `_IOLBF`, and `_IONBF`.
- Optional *buf* and *size* arguments

```
#include <stdio.h>
int fflush(FILE *fp) ;
```

❑ A stream or all output streams, if *fp* is NULL, are flushed.

I. Buffering

Function	<i>mode</i>	<i>buf</i>	Buffer and length	Type of buffering
setbuf		non-null	user <i>buf</i> of length <code>BUFSIZ</code>	fully buffered or line buffered
		NULL	(no buffer)	unbuffered
setvbuf	_IOFBF	non-null	user <i>buf</i> of length <i>size</i>	fully buffered
		NULL	system buffer of appropriate length	
	_IOLBF	non-null	user <i>buf</i> of length <i>size</i>	line buffered
		NULL	System buffer of appropriate length	
	_IONBF	(ignored)	(no buffer)	unbuffered

I. Opening a Stream

```
#include <stdio.h>
```

```
FILE *fopen(const char *pathname, const char *type);
```

```
FILE *freopen(const char *pathname, const char *type, FILE  
*fp);
```

```
FILE *fdopen(int filedes, const char *type);
```

- ❑ **freopen opens a file on a specified stream, closing the stream first, if it's already open. (typically used to open the file as one of stdin, stdout, and stderr.)**
- ❑ **fdopen is often used with pipes and network communication channels. (no fopen for these special files, so we have to call the device-specific function to obtain a file descriptor, and then fdopen)**

I. Opening a Stream

❑ **b** stands for a binary file, but no effect for Unix.

❑ **With fdopen,**

- No truncation for opening for write, since it has already been opened.
- Append mode can not create the file, since it already exists.

<i>type</i>	Description
r or rb	open for reading
w or wb	truncate to 0 length or create for writing
a or ab	append; open for writing at EOF, or create for writing
r+, r+b or rb+	open for reading and writing
w+, w+b or wb+	truncate to 0 length or create for reading and writing
a+, a+b or ab+	open or create for reading and writing at EOF

I. Opening a Stream

```
#include <stdio.h>
```

```
int fclose(FILE *fp);
```

- ❑ **Any buffered output data is flushed, and an automatically allocated buffer is released.**
- ❑ **When a process terminates, all I/O streams are flushed and closed.**

I. Reading and Writing a Stream

❑ Unformatted I/O

- Character-at-a-time I/O, e.g. `getc`
- Line-at-a-time I/O, e.g. `fgets` and `fputs`
- Direct I/O, e.g. `fread` and `fwrite`
 - a.k.a. binary I/O, object-at-a-time, record-oriented I/O, or structure-oriented I/O
 - Read or write some number of objects, often used for binary files

I. Reading and Writing a Stream

```
#include <stdio.h>
int getc(FILE *fp) ;
int fgetc(FILE *fp) ;
int getchar(void) ;
```

- ❑ **All three return the next char (as an unsigned char converted to an int) if OK, EOF on end of file or error.**
- ❑ **getchar is equivalent to getc(stdin)**
- ❑ **getc can be implemented as a macro, but fgetc is guaranteed to be a function.**

```
#include <stdio.h>
int ferror(FILE *fp) ;
int feof(FILE *fp) ;
void clearerr(FILE *fp) ; /* clear both flags */
```

- ❑ **The *getc* functions return the same value whether an error occurs or the end of file is reached. We call *ferror* or *feof* to distinguish between the two.**

I. Reading and Writing a Stream

```
#include <stdio.h>
int ungetc(int c, FILE *fp) ;
```

- ❑ Pushback of a char
- ❑ does not need to be the same char read.

```
❑ #include <stdio.h>
int putc(int c, FILE *fp) ;
int fputc(int c, FILE *fp) ;
int putchar(int c) ;
```

```
❑ putchar(c) = putc(c, stdout)
```

```
❑ putc can be implemented as a macro, but  
fputc must as a function.
```

I. Line-at-a-Time I/O

```
#include <stdio.h>
```

```
char *fgets(char *buf, int n, FILE *fp);
```

```
char *gets(char *buf);
```

- ❑ Both return `buf` if OK, NULL on end of file or error.
- ❑ `fgets` read up to `n-1` chars, including a newline, and the buffer is null-terminated.
- ❑ `gets` is a deprecated function subject to buffer overflow, if the line is longer than the buffer, and it does not store a new line.

I. Line-at-a-Time I/O

```
#include <stdio.h>
int fputs(const char *str, FILE *fp) ;
int puts(const char *str) ;
```

- ❑ **fputs write the null-terminated string to the stream.**
 - No need to be line-at-a-time output, since the string need not contain a newline at the last char (which is usually the case.)
- ❑ **puts writes the null-terminated string to the standard output, and then a newline.**

I. Standard I/O Efficiency

```
#include "apue.h"
#define BUFSIZE 4096
int main(void) {
    int n;
    char buf[BUFSIZE];
    while ((n = read(STDIN_FILENO, buf, BUFSIZE)) > 0)
        if (write(STDOUT_FILENO, buf, n) != n)
            err_sys("write error");
    if (n < 0)
        err_sys("read error");
    exit(0);
}
```

Fig 3.5 - read() and

```
#include "apue.h"
int main(void) {
    int c;
    while ((c = getc(stdin)) != EOF)
        if (putc(c, stdout) == EOF)
            err_sys("output error");
    if (ferror(stdin))
        err_sys("input error");
    exit(0);
}
```

Fig 5.4 - getc() and putc()

```
#include "apue.h"
int main(void) {
    char buf[MAXLINE];
    while (fgets(buf, MAXLINE, stdin) != NULL)
        if (fputs(buf, stdout) == EOF)
            err_sys("output error");
    if (ferror(stdin))
        err_sys("input error");
    exit(0); }

```

Fig 5.5 – fgets() and fputs()

I. Standard I/O Efficiency

- ❑ **Copying a 98.5 MB file from `stdin` to `stdout`**
- ❑ **User CPU time dominated by a loop which is executed**
 - 100 million times in the character-at-a-time versions.
 - 3,144,984 times in the line-at-a-time version.
 - 12,611 times in the `read` version (for a buffer size of 8192.)
- ❑ **System CPU time is the same as before, because the standard I/O routines performs buffering using the optimal I/O size.**

Function	User CPU (seconds)	System CPU (seconds)	Clock time (seconds)	Bytes of program text
best time from Figure 3.5	0.01	0.18	6.67	
<code>fgets</code> , <code>fputs</code>	2.59	0.19	7.15	139
<code>getc</code> , <code>putc</code>	10.84	0.27	12.07	120
<code>fgetc</code> , <code>fputc</code>	10.44	0.27	11.42	120
single byte time from Fig3.5	124.89	161.65	288.64	

I. Binary I/O

```
#include <stdio.h>
size_t fread(void *ptr, size_t size, size_t nobj, FILE *fp) ;
size_t fwrite(const void *ptr, size_t size, size_t nobj, FILE *fp);
```

□ Read or write objects

– Example 1

```
float data[10]
if (fwrite(&data[2], sizeof(float), 4, fp) != 4)
    err_sys("fwrite error");
```

– Example 2

```
struct {
    short count;
    long total;
    char name[NAMESIZE];
} item;
if (fwrite(&item, sizeof(item), 1, fp) != 1)
    err_sys("fwrite error");
```

I. Binary I/O

- ❑ **fread returns the number of objects less than `nobj` on the end of file or error**
- ❑ **If `fwrite` returns the value less than the `nobj`, an error has occurred.**
- ❑ **Portability problem with binary I/O**
 - The offset of a member within a structure may differ between compilers and systems (i.e. different alignments)
 - Byte ordering for multibyte integers and floating-point values

I. Positioning a Stream

```
#include <stdio.h>
long ftell(FILE *fp);
int fseek(FILE *fp, long offset, int whence);
void rewind(FILE *fp);
```

- ❑ **whence: SEEK_SET, SEEK_CUR, SEEK_END** that are the same as `lseek`
- ❑ **A stream is set to the beginning of the file with `rewind`.**

```
#include <stdio.h>
off_t ftello(FILE *fp);
int fseeko(FILE *fp, off_t offset, int whence);
```

- ❑ **Introduced in the Single UNIX Specification**
- ❑ **The type of the `offset` is `off_t` instead of `long`.**

I. Positioning a Stream

```
#include <stdio.h>
```

```
int fgetpos(FILE *fp, fpos_t *pos)
```

```
int fsetpos(FILE *fp, const fpos_t *pos);
```

- ❑ **New with the ISO C standard**

- ❑ **fgetpos stores the file's current position indicator in *pos*, which can be used in a later call to fsetpos to reposition the stream.**

I. Formatted I/O

```
#include <stdio.h>
int printf(const char *format, ...);
int fprintf(FILE *fp, const char *format, ...);
int sprintf(char *buf, const char *format, ...);
int snprintf(char *buf, size_t n, const char *format, ...);

#include <stdarg.h>
#include <stdio.h>
int vprintf(const char *format, va_list arg);
int vfprintf(FILE *fp, const char *format, va_list arg);
int vsprintf(char *buf, const char *format, va_list arg);
int vsnprintf(char *buf, size_t n, const char *format, va_list arg);

#include <stdio.h>
int scanf(const char *format, ...);
int fscanf(FILE *fp, const char *format, ...);
int sscanf(const char *buf, const char *format, ...);

#include <stdarg.h>
#include <stdio.h>
int vscanf(const char *format, va_list arg);
int vfscanf(FILE *fp, const char *format, va_list arg);
int vsscanf(const char *buf, const char *format, va_list arg);
```

I. Implementation Details

```
#include <stdio.h>
```

```
int fileno(FILE *fp);
```

- ❑ **It returns the file descriptor associated with the stream (if you want to call `dup`, `fcntl`, etc.)**

- ❑ **Print buffering for various standard I/O streams (Figure 5.11)**
 - `stdin` and `stdout`: line buffered
 - `stderr`: unbuffered
 - Regular files: fully buffered

I. Print buffering for various standard I/O streams

Fig 5.11

```
#include "apue.h"
void pr_stdio(const char *, FILE *);
int main(void) {
    FILE *fp;
    fputs("enter any character\n", stdout);
    if (getchar() == EOF) err_sys("getchar
error");
    fputs("one line to standard error\n", stderr);
    pr_stdio("stdin", stdin);
    pr_stdio("stdout", stdout);
    pr_stdio("stderr", stderr);

    if ((fp = fopen("/etc/motd", "r")) == NULL)
        err_sys("fopen error");
    if (getc(fp) == EOF) err_sys("getc error");

    pr_stdio("/etc/motd", fp);
    exit(0);
}
```

```
void pr_stdio(const char *name, FILE *fp)
{
    printf("stream = %s, ", name);
    /* The following is nonportable. */
    if (fp->_IO_file_flags & _IO_UNBUFFERED)
        printf("unbuffered");
    else if (fp->_IO_file_flags & _IO_LINE_BUF)
        printf("line buffered");
    else /* if neither of above */
        printf("fully buffered");
    printf(", buffer size = %d\n",
           fp->_IO_buf_end - fp->_IO_buf_base);
}
```

I. Print buffering for various standard I/O streams

Fig 5.11

```
$ ./a.out stdin, stdout, and stderr connected to terminal
enter any character
we type a newline
one line to standard error
stream = stdin, line buffered, buffer size = 1024
stream = stdout, line buffered, buffer size = 1024
stream = stderr, unbuffered, buffer size = 1
stream = /etc/motd, fully buffered, buffer size = 4096
$ ./a.out < /etc/termcap > std.out 2> std.err
$ cat std.err
one line to standard error
$ cat std.out
enter any character
stream = stdin, fully buffered, buffer size = 4096
stream = stdout, fully buffered, buffer size = 4096
stream = stderr, unbuffered, buffer size = 1
stream = /etc/motd, fully buffered, buffer size = 4096
```

I. Temporary Files

```
#include <stdio.h>
char *tmpnam(char *ptr) ;
FILE *tmpfile(void) ;
```

- ❑ **tmpnam generates a unique file name using the path prefix P_tmpdir.**
 - If *ptr* is `NULL`, the generated pathname is stored in an internal static area.
 - else an array of at least `L_tmpnam` is assumed.
- ❑ **tmpfile creates a temporary file (type `wb+`) that is automatically removed when it is closed or on program termination.**
 - Equivalent to (`tmpnam`, then create the file, and immediately `unlink` it)

I. Temporary Files

Figure 5.12 – Demonstrate `tmpnam` and `tmpfile` functions

```
#include "apue.h"
int main(void) {
    char name[L_tmpnam], line[MAXLINE];
    FILE *fp;
    printf("%s\n", tmpnam(NULL)); /* first temp name */
    tmpnam(name); /* second temp name */
    printf("%s\n", name);

    if ((fp = tmpfile()) == NULL) /* create temp file */
        err_sys("tmpfile error");
    fputs("one line of output\n", fp); /* write to temp file */
    rewind(fp); /* then read it back */
    if (fgets(line, sizeof(line), fp) == NULL)
        err_sys("fgets error");
    fputs(line, stdout); /* print the line we wrote */
    exit(0);
}
```

```
$ ./a.out
/tmp/fileC1Icwc
/tmp/filemSkHSe
one line of output
```

Thank you for your attention !!

Q and A