
Network IPC: Sockets

System Programming

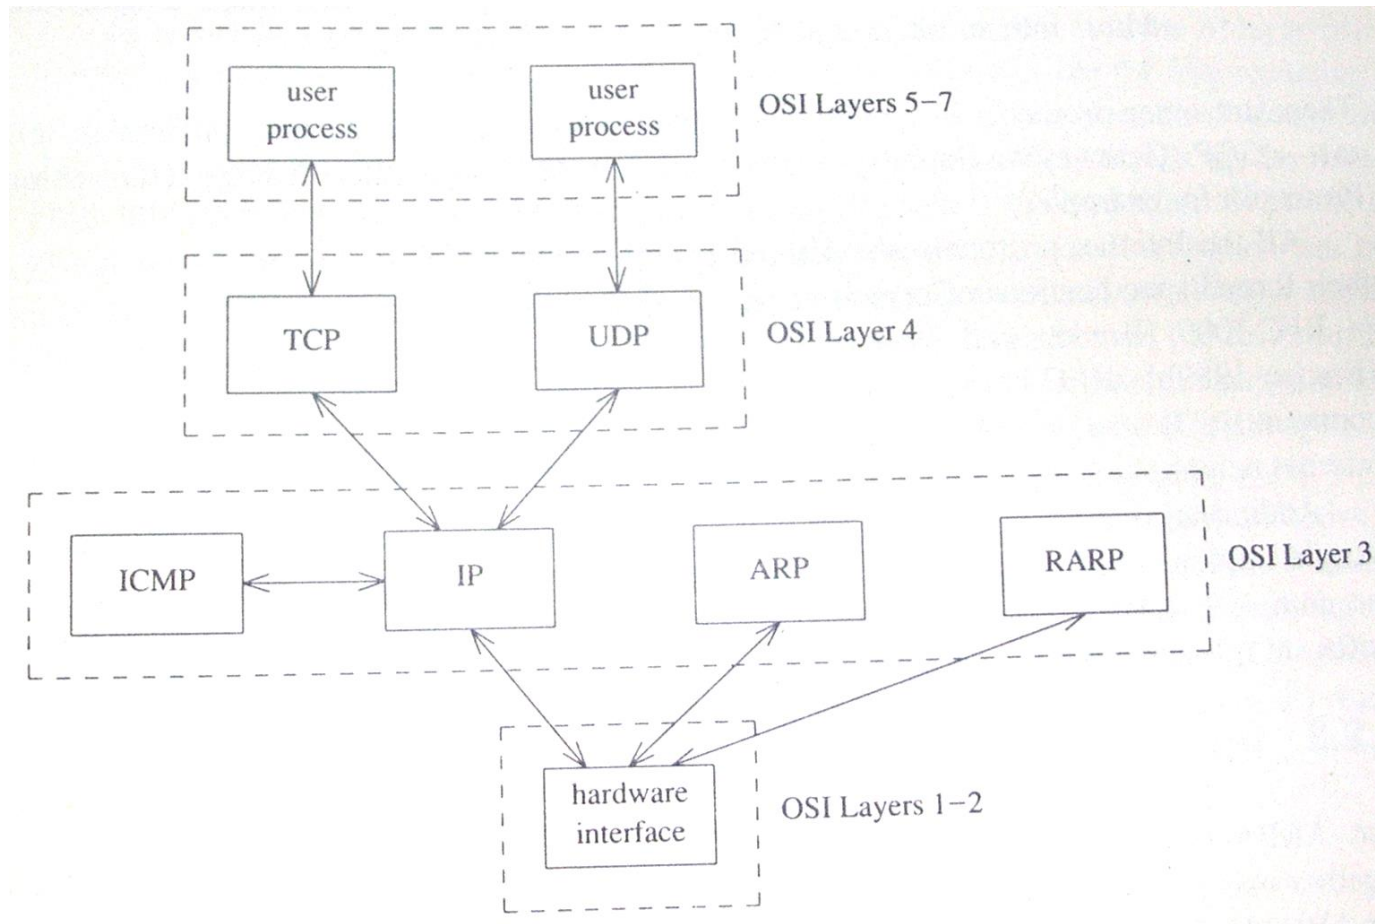
2019 여름 계절학기

한양대학교 공과대학 컴퓨터소프트웨어학부
홍석준

Introduction

- **Socket network IPC interface to support process communication on the same machine or on different machines**
 - Classical IPCs : pipes, FIFOs, message queues, semaphores, and shared memory on the same computer
- **The POSIX.1 socket API based on the 4.4BSD socket interface**
 - Originally introduced in 4.2BSD in the early 1980s.

The Internet Protocol Suite



Connection-Oriented vs. Connectionless

❑ Connectionless protocol (Datagram interface)

- A datagram is a self-contained message.
- Sending a datagram is analogous to mailing someone a letter.
- No logical connection between peers
 - Subject to out-of-order delivery and loss of packet
- UDP (User Datagram Protocol)

❑ Connection-oriented protocol (Stream interface)

- Like making a phone call.
- A point-to-point virtual connection between both ends of the call
- TCP (Transmission Control Protocol)

Socket Descriptors

- ❑ A socket is an abstraction of a communication endpoint.
- ❑ Socket descriptors to access sockets
 - Implemented as file descriptors

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol);
```

- ❑ Domain: AF_INET, AF_INET6, AF_UNIX, AF_UNSPEC
- ❑ Types: SOCK_DGRAM, SOCK_RAW, SOCK_SEQPACKET, SOCK_STREAM
- ❑ Protocol: zero value to indicate the default protocol for the given domain and socket type
 - TCP for AF_INET and SOCK_STREAM
 - UDP for AF_INET and SOCK_DGRAM

Socket Types

Type	Description
SOCK_DGRAM	fixed-length, connectionless, unreliable messages
SOCK_RAW	datagram interface to IP (optional in POSIX.1)
SOCK_SEQPACKET	fixed-length, sequenced, reliable, connection-oriented messages
SOCK_STREAM	sequenced, reliable, bidirectional, connection-oriented byte streams

Address Formats

- ❑ An address identifies a socket endpoint. Its format is specific to the particular domain, being cast to a generic `sockaddr` address structure to be passed to socket functions.

```
struct sockaddr {  
    sa_family_t    sa_family; /* address family */  
    char           sa_data[]; /* variable-length address */  
    ...  
};
```

- ❑ Internet address defined in `<netinet/in.h>`

```
struct in_addr {  
    in_addr_t    s_addr; /* IPv4 address */  
};  
  
struct sockaddr_in {  
    sa_family_t    sin_family; /* address family */  
    in_port_t      sin_port; /* port number */  
    struct in_addr sin_addr; /* IPv4 address */  
};
```

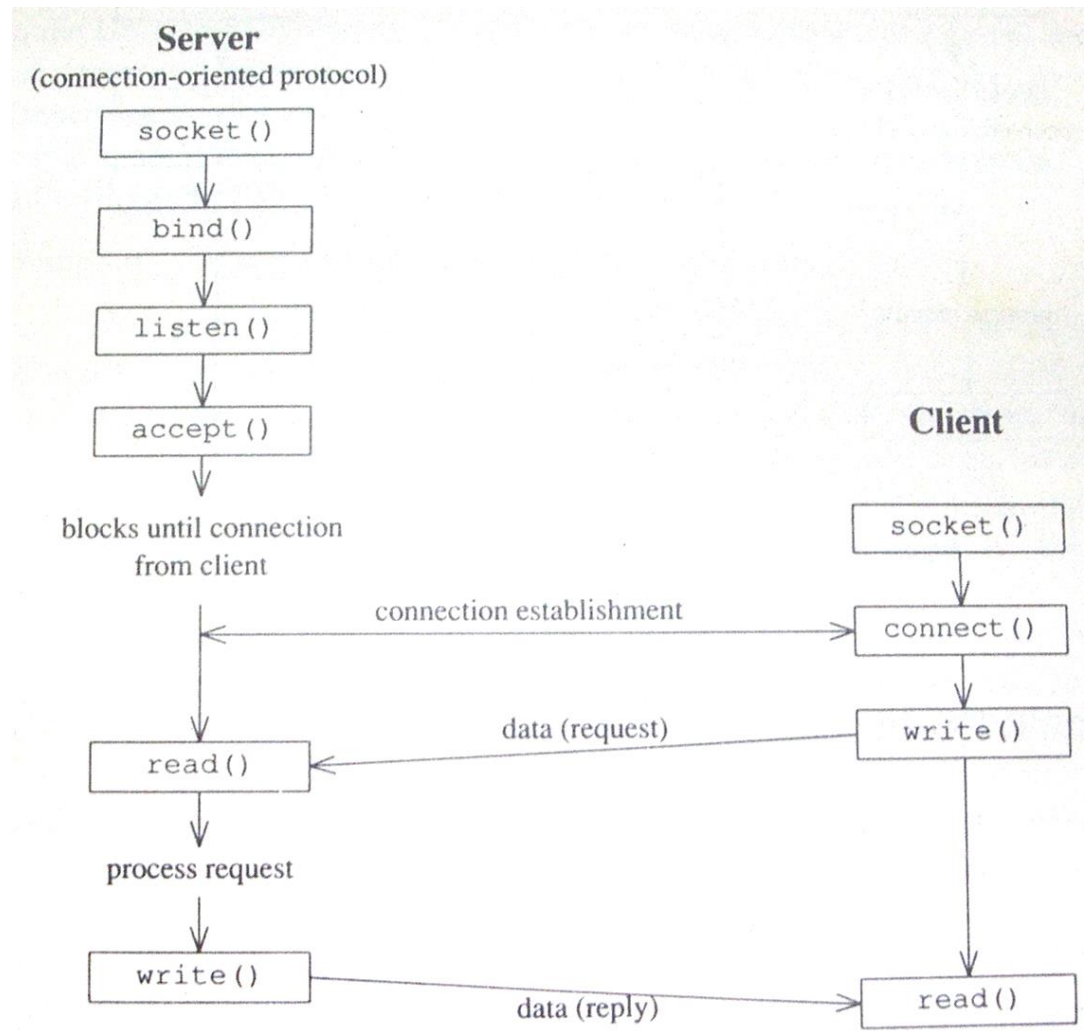
```
#include <arpa/inet.h>
```

```
const char *inet_ntop(int domain, const void *addr,  
    char *str, socklen_t size);
```

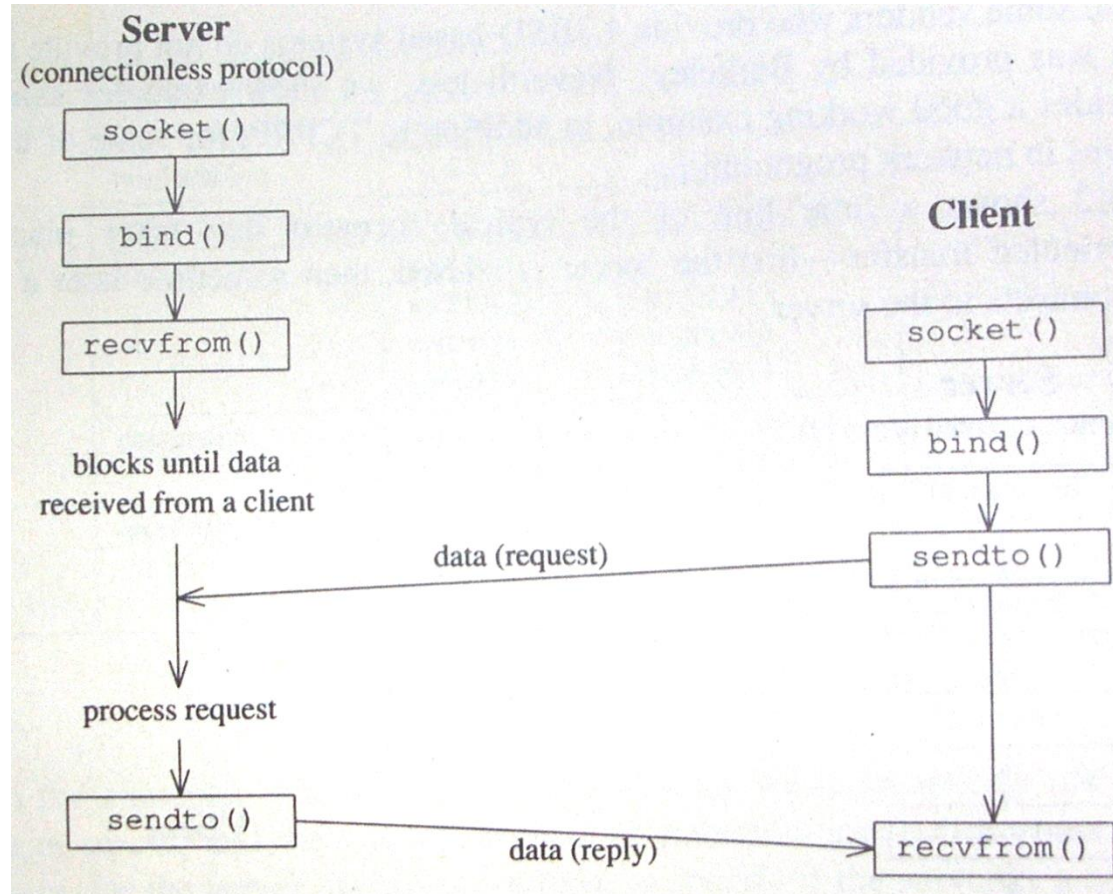
```
int inet_pton(int domain, const char *str, void *addr);
```

- ❑ Conversion between the binary address format and the dotted-decimal notation
- ❑ Only AF_INET and AF_INET6 domains supported

Socket System Calls for Connection-Oriented Protocol



Socket System Calls for Connectionless Protocol



Associating Addresses with Sockets

```
#include <sys/socket.h>
```

```
int bind(int sockfd, const struct sockaddr *addr, socklen_t  
        len) ;
```

- ❑ To associate an address with a socket, to discover the address bound to a socket, and to find out the peer's address
- ❑ For a client's socket, we can let the system choose a default address for us, whereas a well-known address should be associated with the server's socket.
- ❑ The system will choose an address and bind it to our socket, if we call `connect` or `listen` without first binding an address to the socket.

```
int getsockname(int sockfd, struct sockaddr *addr, socklen_t  
                *alenp) ;
```

```
int getpeername(int sockfd, struct sockaddr *addr, socklen_t  
                *alenp) ;
```

returned
value

Connection Establishment

```
#include <sys/socket.h>
```

```
int connect(int sockfd, const struct sockaddr *addr,  
            socklen_t len);
```

- ❑ It creates a connection to the specified server, *addr*. If *sockfd* is not bound to an address, a default address will be bound to the caller.

- ❑ [Figure 16.9](#)

Connection Establishment

```
#include <sys/socket.h>
int listen(int sockfd, int backlog);
```

- ❑ The *backlog* argument provides a hint to the system of the number of outstanding connect requests that it should enqueue on behalf of the process.

Connection Establishment

```
#include <sys/socket.h>
```

```
int accept(int sockfd, struct sockaddr *addr, socklen_t  
    *len) ;
```

- ❑ The file descriptor returned is a socket descriptor that is connected to the client that called `connect`. This new socket descriptor has the same socket type and address family as *sockfd*.
- ❑ The original socket passed to `accept` is not associated with the connection, but instead remains available to receive additional connect requests.
- ❑ On return, `accept` will fill in the client's address in *addr*.
- ❑ [Figure 16.10](#)

- ❑ `read/write` can be used to communicate with a socket.

```
#include <sys/socket.h>
```

```
ssize_t send(int sockfd, const void *buf, size_t nbytes, int  
flags) ;
```

```
#include <sys/socket.h>
```

```
ssize_t sendto(int sockfd, const void *buf, size_t nbytes, int  
flags, const struct sockaddr *destaddr, socklen_t destlen) ;
```

- ❑ With a connection-oriented socket, the destination address is ignored, as the destination is implied by the connection. (With a connectionless socket, we can't use `send` unless the destination address is first set by calling `connect`.)

```
#include <sys/socket.h>
```

```
ssize_t recv(int sockfd, void *buf, size_t nbytes, int  
    flags) ;
```

```
#include <sys/socket.h>
```

```
ssize_t recvfrom(int sockfd, void *buf, size_t len, int  
    flags, struct sockaddr *addr, socklen_t *addrlen) ;
```

Thank you for your attention !!

Q and A